

Rethinking Statistics

EES 5891-03

Bayesian Statistical Methods

Jonathan Gilligan

Class #2: Tuesday, August 30 2022

Golems and Rethinking

Statistical Tools as Golems

- Statistical tools will do what you tell them to do ...
 - but if you're not careful, what you tell them to do may not be what you want them to do
- The goal of this book is to help you:
 - Learn to use statistical golems wisely
 - Learn to choose the right golem for the job
 - Learn to engineer your own golems if the ready-to-use golems aren't right for your job.




The Golem of Prague (Photo: Prague Post)

Hypothesis Testing

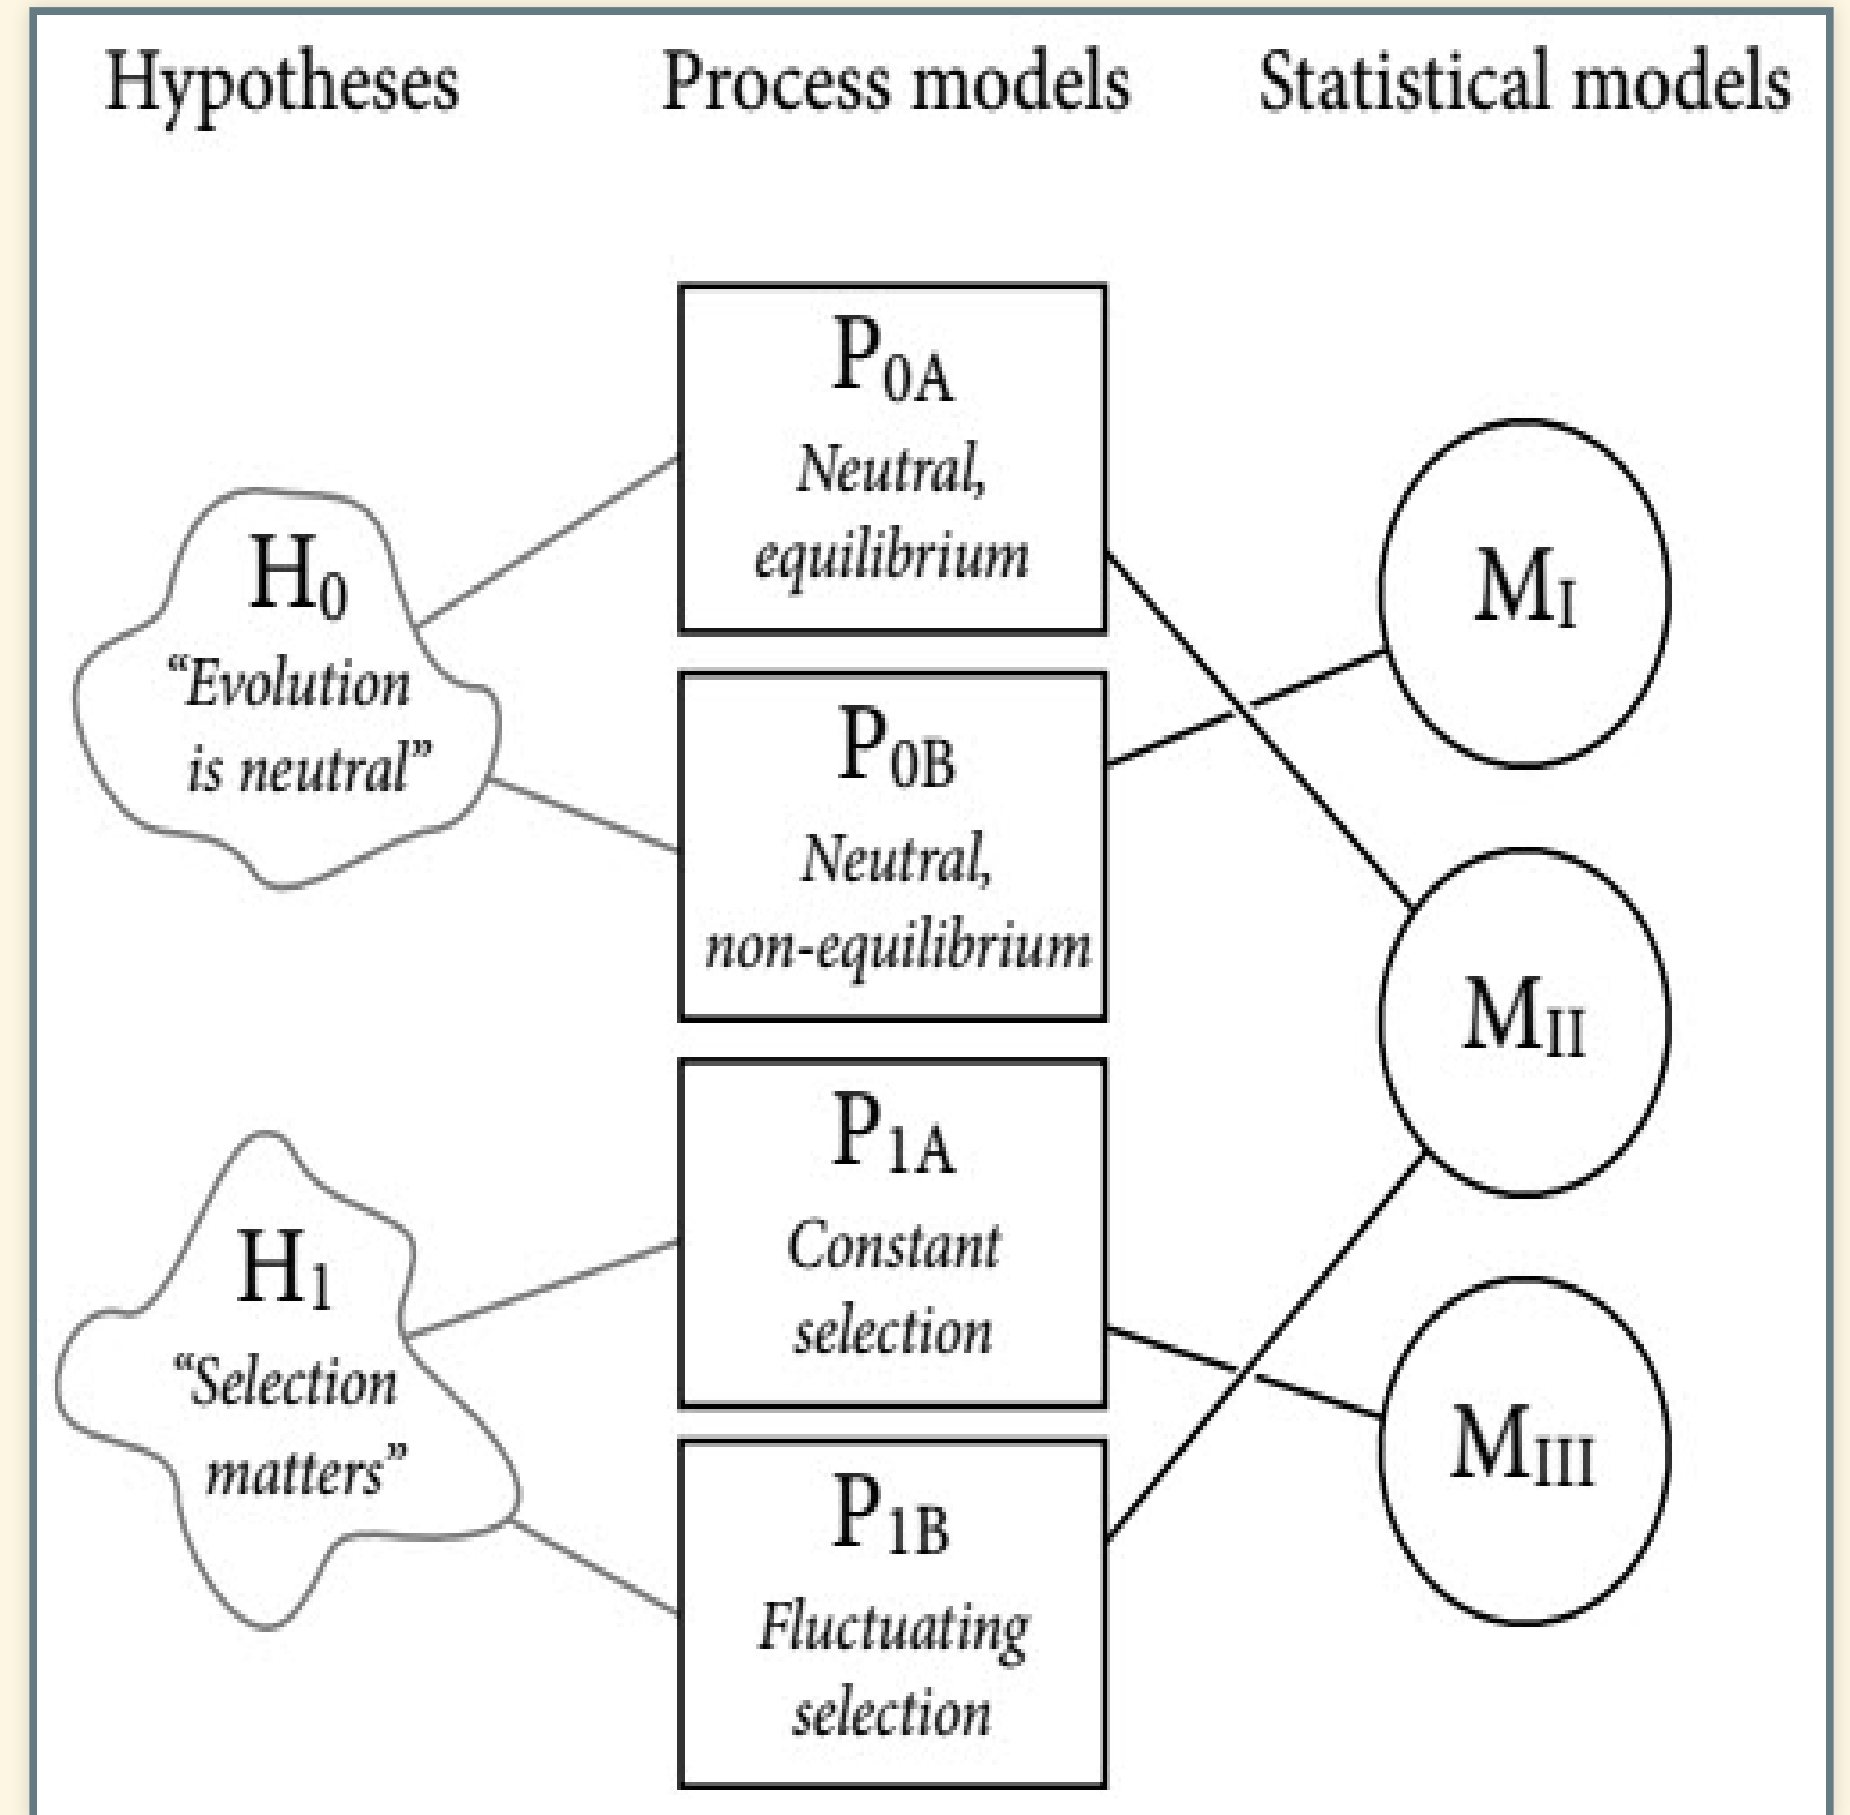
- Karl Popper (1902–1994)
 - Science can never prove that a hypothesis is true
 - But it can prove that an incorrect hypothesis is false
 - The more false hypotheses we rule out, we narrow down the list of potentially true hypotheses.

*When you have eliminated the impossible,
whatever remains, however improbable,
must be the truth*
— Arthur Conan Doyle/Sherlock Holmes



Problems with falsification

- The predictions of a hypothesis may not be as clear as many people assume.
- Depending on what other assumptions you make, two different hypotheses may predict the same kind of data.
 - If your data looks like M_I , it rules out (falsifies) H_1
 - But if your data looks like M_{II} , it doesn't rule out either hypothesis.
- A given hypothesis may predict many different possible kinds of data, depending on what other assumptions you make.
 - If your data doesn't look like M_I , that doesn't imply that it's less likely H_1 is true.



Null-Hypothesis Significance Testing

- Most statistical tests aim to rule out a *null hypothesis*, not to falsify the actual research hypothesis.
- Often, there's not one unique alternative hypothesis to the null hypothesis, so even if we reject the null hypothesis, there are many other possibilities.
- Bayesian methods give us better, more powerful golems to answer the questions we're really interested in.
 - But they're still golems and we have to be thoughtful and careful about how to use them.

Discussion

Bayes's Theorem

Bayes's Theorem

- The core of this part of the course is Bayes's theorem.
- Notation:
 - Conditional probability: $P(a|b)$ means *the probability of a, given b*.
- Bayes's theorem:

$$P(H|D) = \frac{P(D|H) \times P(H)}{P(D)},$$

where

- $P(H|D)$ is the *posterior*: The probability that H is true, given that you observed D .
- $P(D|H)$ is the *likelihood*: The probability that you would observe D , if H is true.
- $P(H)$ is the *prior* probability of H , based on what you knew before observing D
- $P(D)$ is the *evidence*: The probability that you would observe D , regardless whether H is true.
 - If H is binary (true or false), then

$$P(D) = P(D|H) \times P(H) + P(D|\text{not}H) \times (1 - P(H))$$

Bayes's Theorem (cont)

- We can apply Bayes's theorem to a numbers too. For a variable x that we want to predict:

$$P(x|D) = \frac{P(D|x) \times P(x)}{P(D)}$$

- In this case,


$$P(D) = \int_{-\infty}^{\infty} P(D|x) \times P(x)dx,$$

where the *prior* $P(x)$ is a probability density function.

Sampling

Sampling

- You have a globe and want to figure out what fraction of the earth's surface is water.
- Toss the globe in the air, catch it, and note whether your index finger is on water or land: outcomes are W and L .
- At every toss, use Bayes's theorem to update your estimate of the fraction that is water.




First toss

- Before you toss the globe, pick a prior probability distribution for the fraction that's water.
- Suppose we don't know anything.
 - Pick $p \sim \text{Uniform}(0, 1)$, a uniform prior:
- Toss the globe and your finger lands on water.
- Update the probability:

$$P(p|W) = \frac{P(W|p)P(p)}{p(W)},$$

where p is the probability of water, and W is measuring water.

Prior:




The calculation:

$$P(p|W) = \frac{P(W|p)P(p)}{p(W)},$$

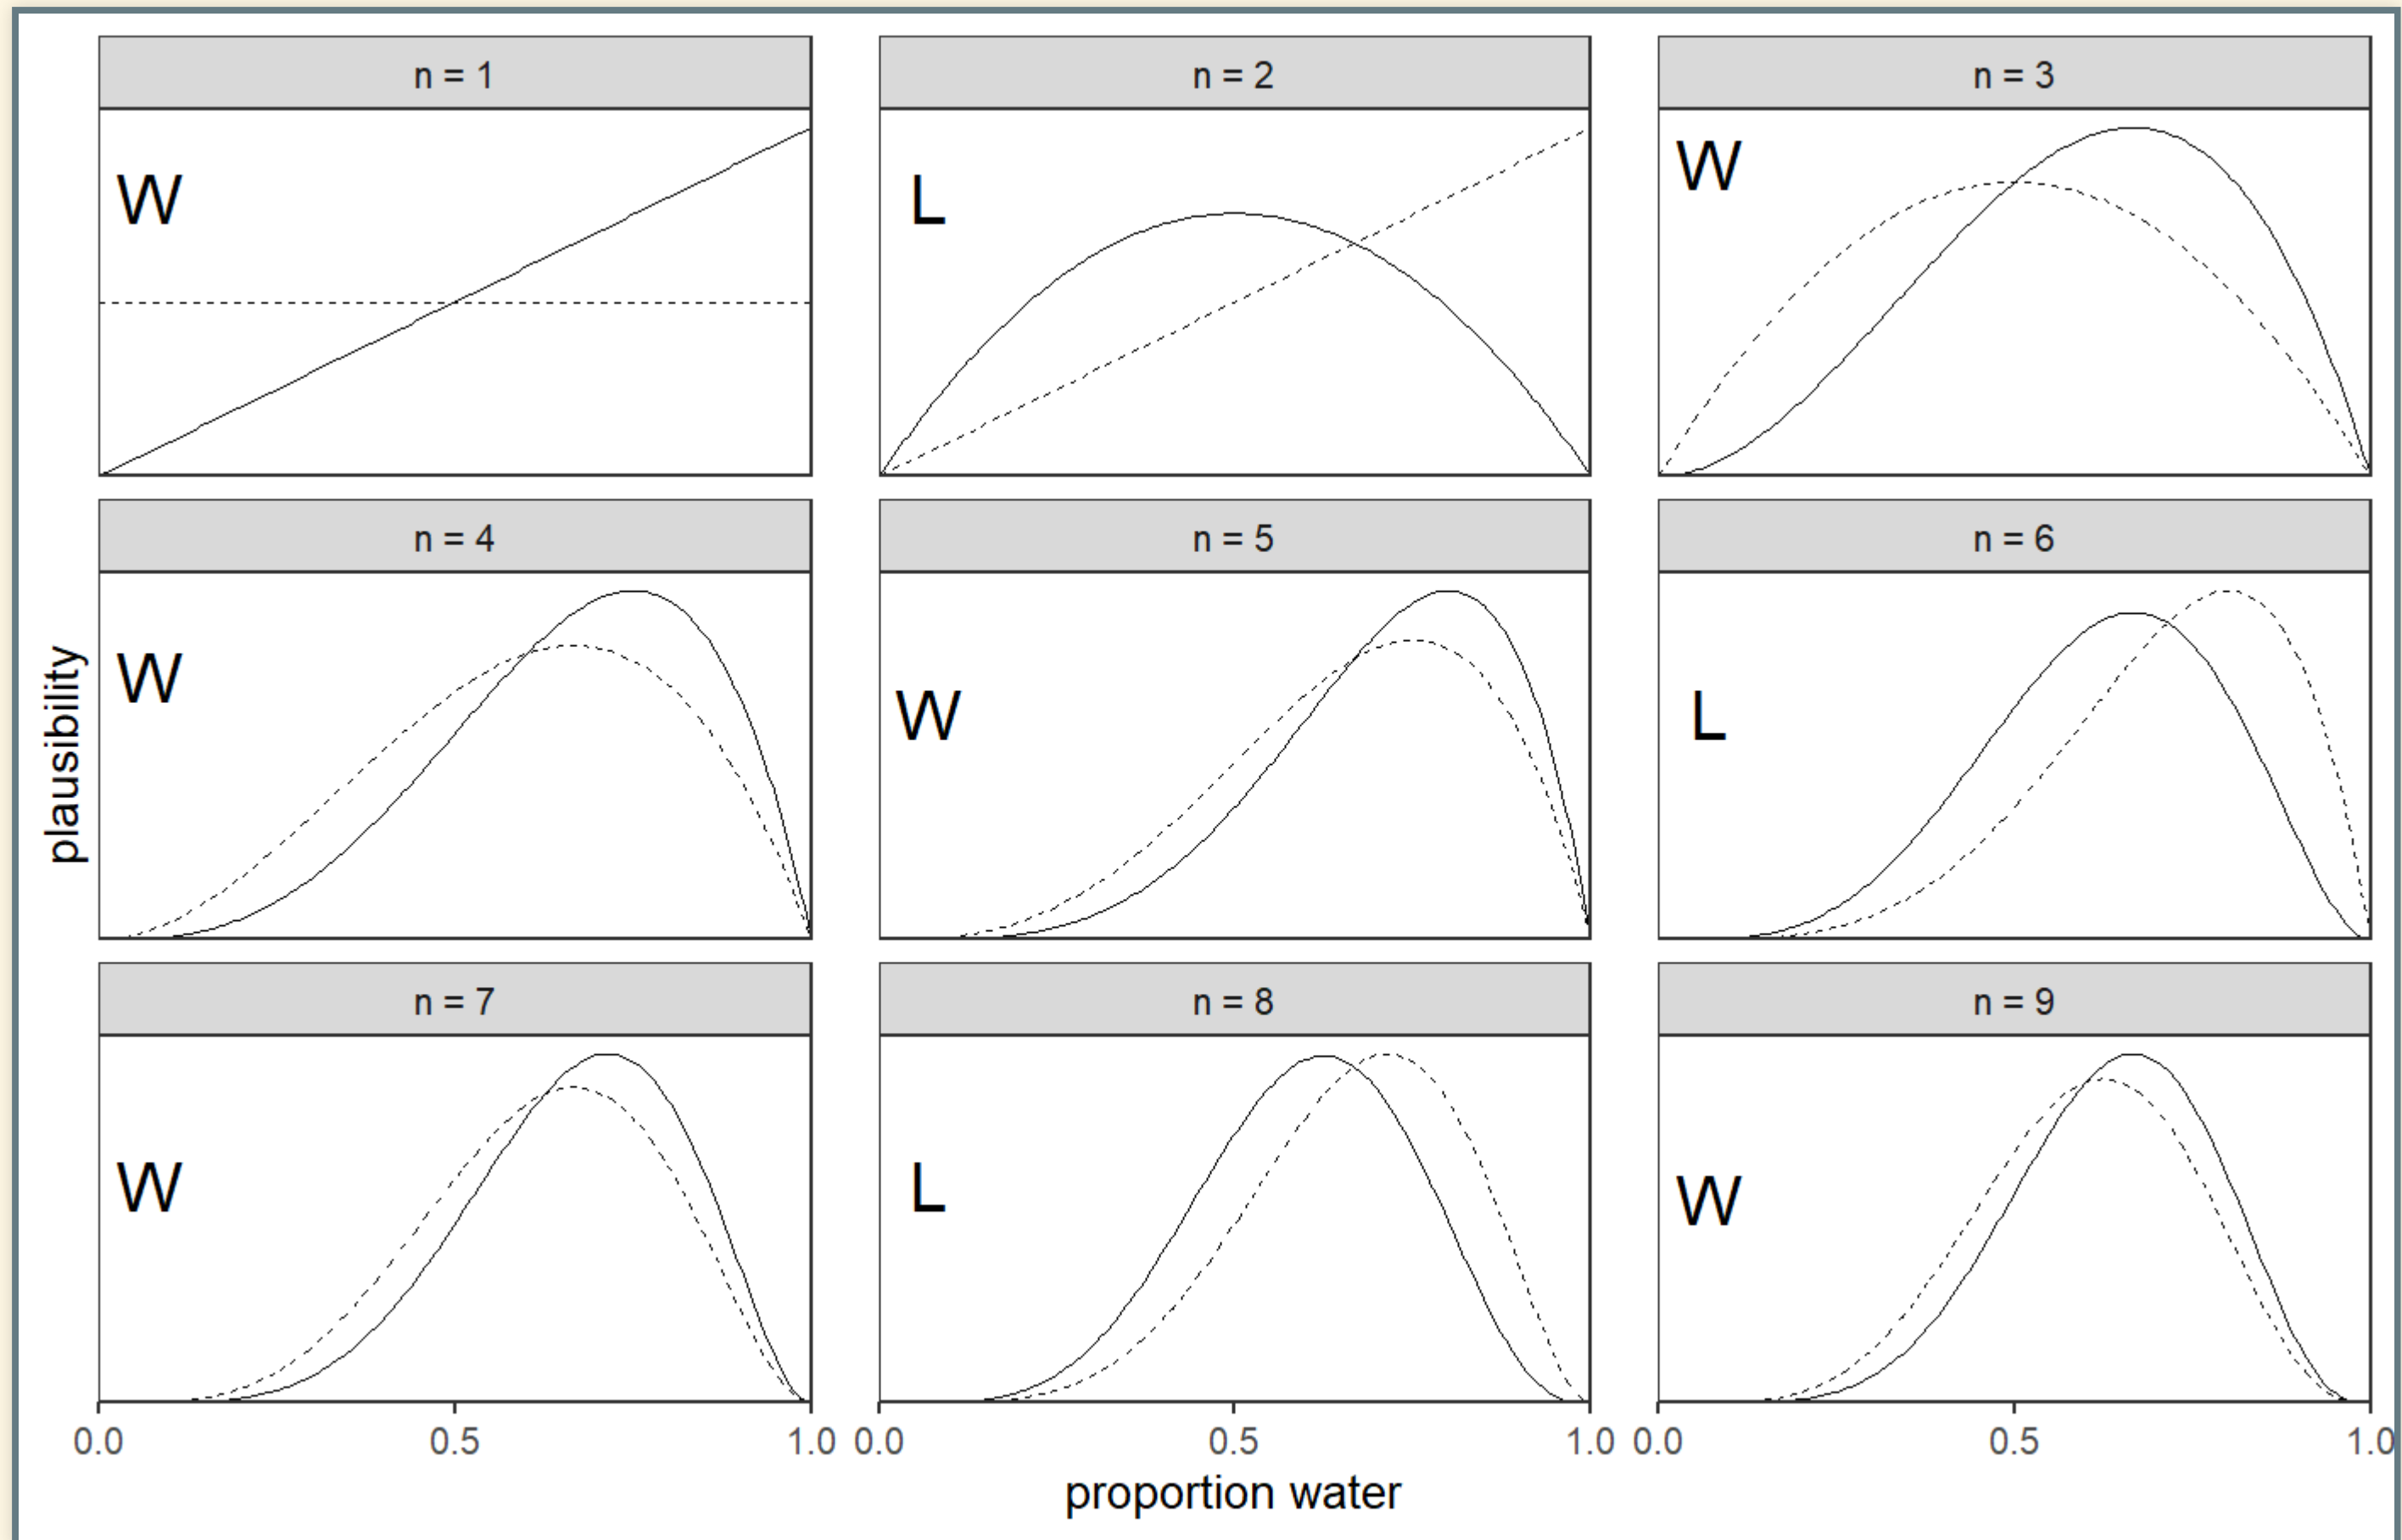
- $P(W|p) = p$
- $P(p) = 1$ (it's a uniform distribution)

$$\begin{aligned} P(W) &= \int_0^1 P(W|p)P(p)dp = \int_0^1 p \times 1dp \\ &= (p^2/2)\Big|_0^1 = 1/2 \end{aligned}$$

- so the posterior $P(p|W) = 2p$
- Use this posterior as the prior for the next toss...



Subsequent tosses

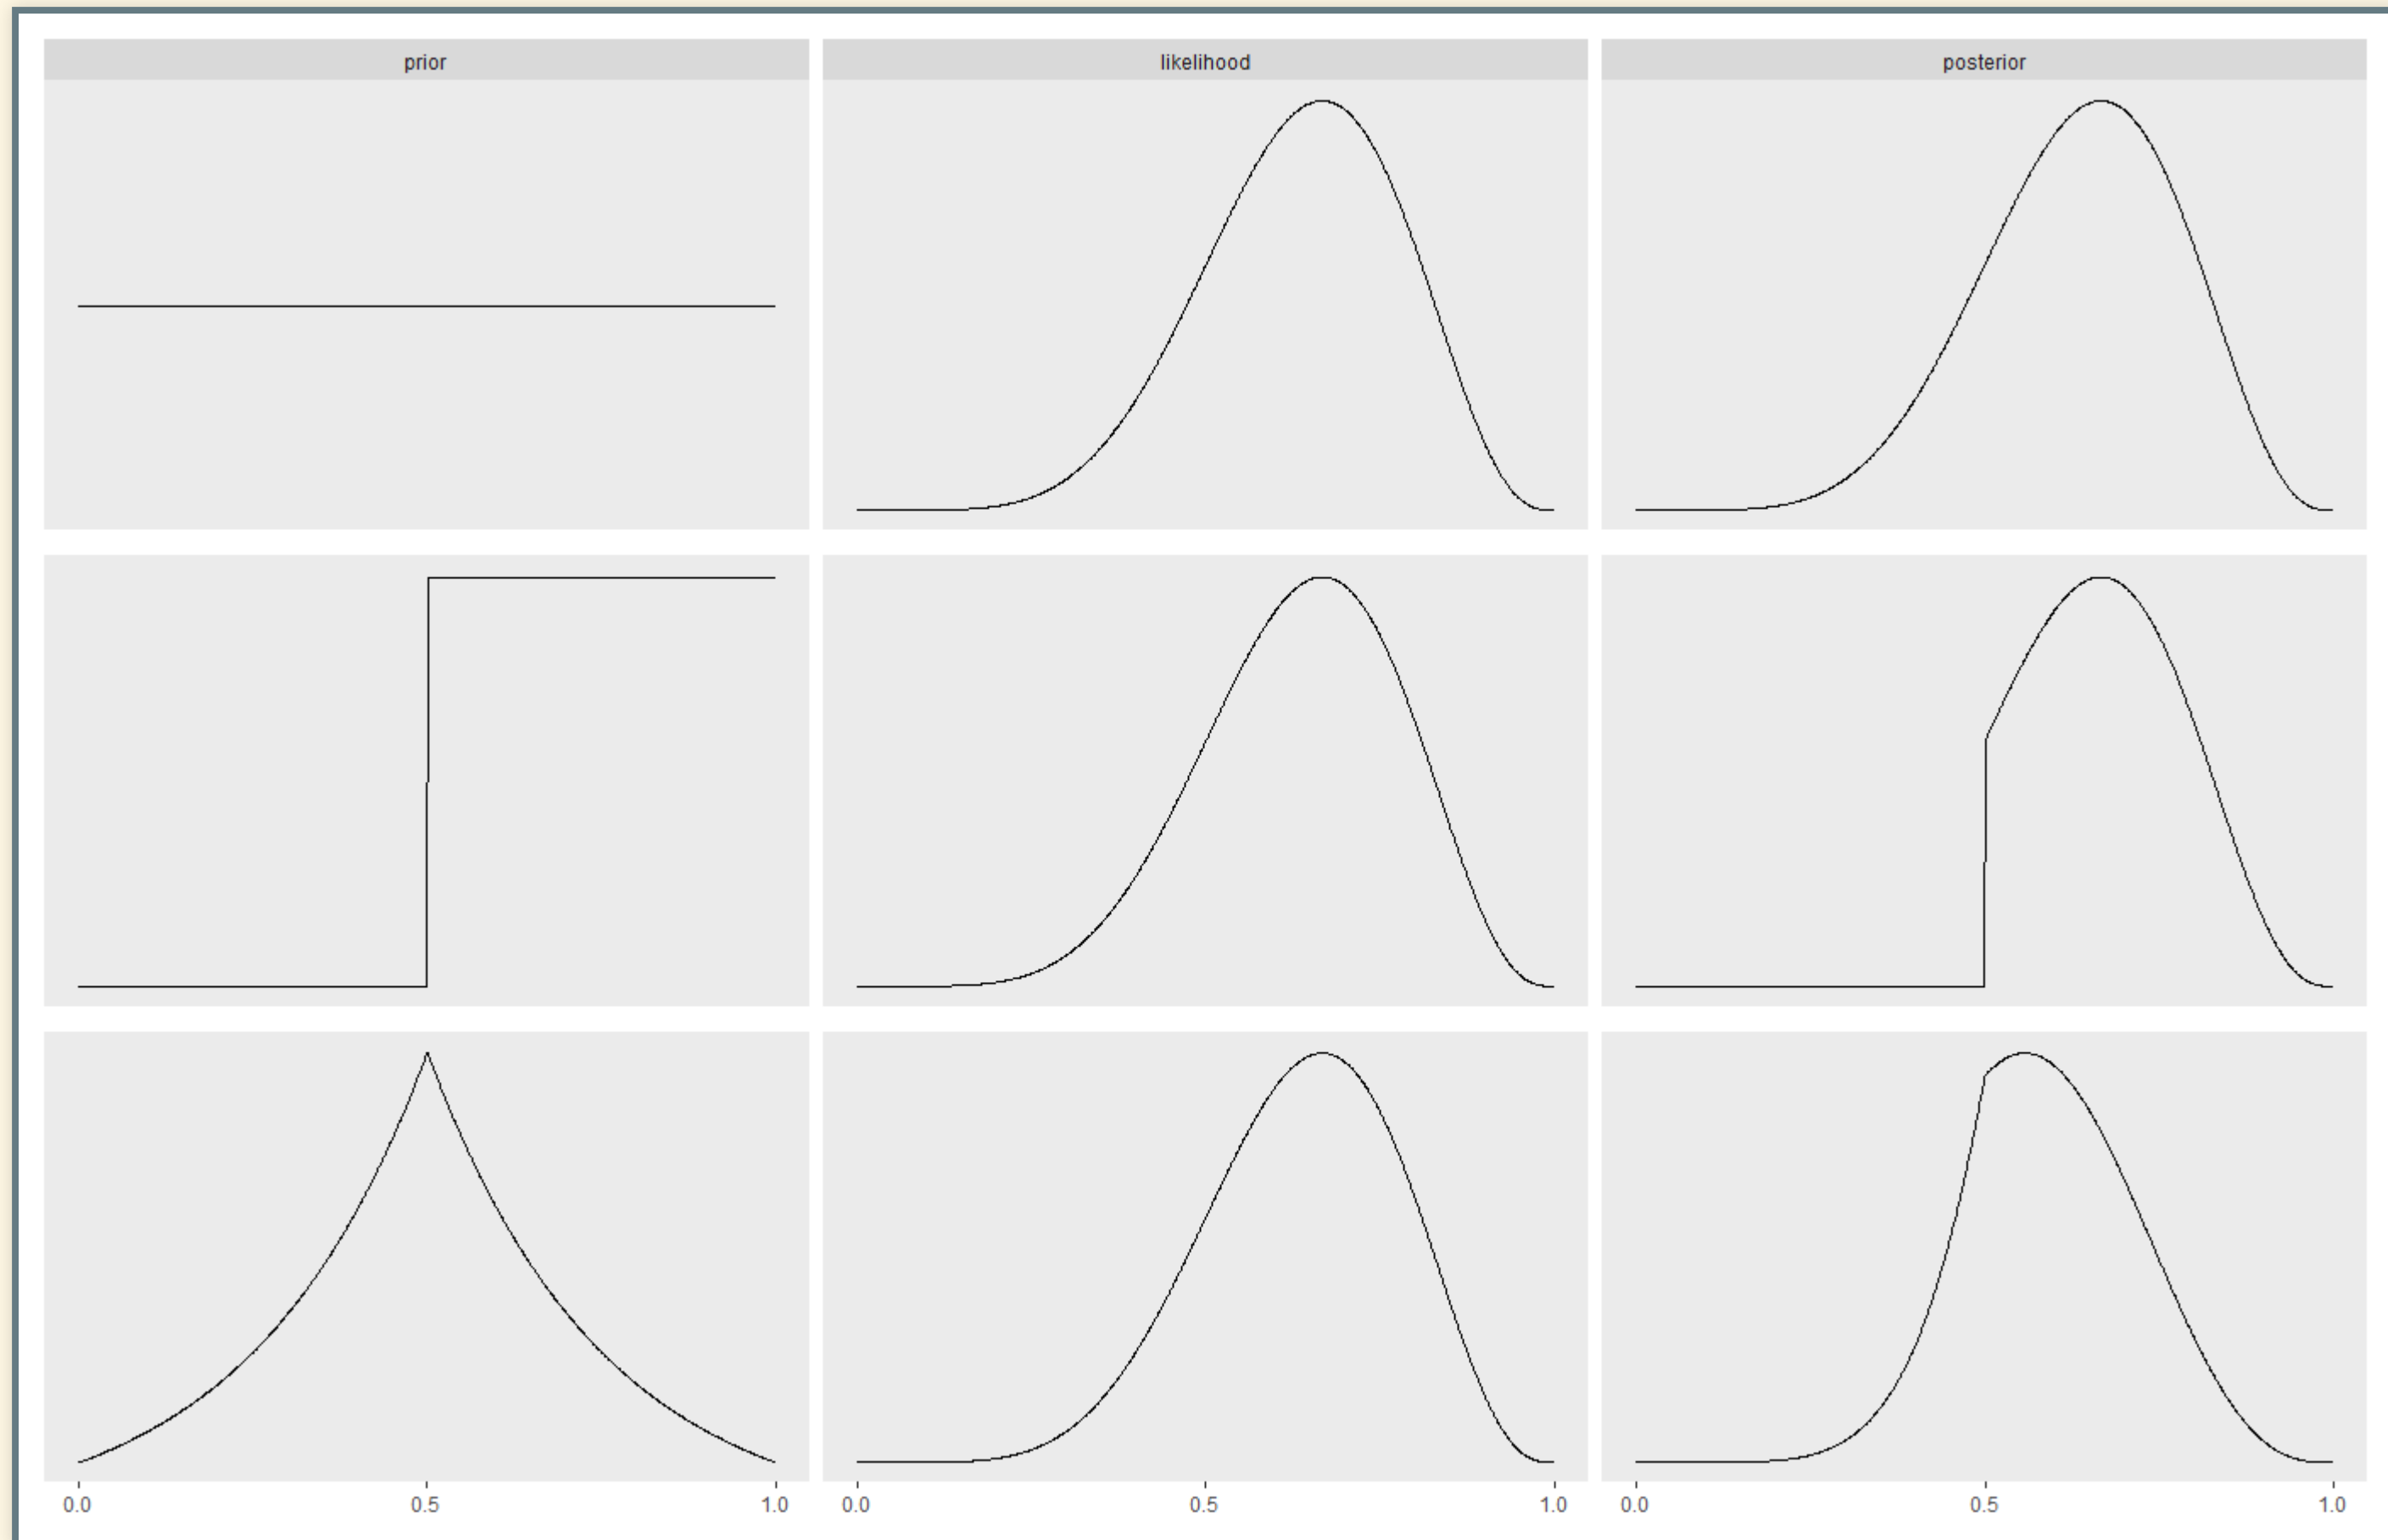


Developing a Model

Developing a Model

- Observations and parameters are drawn from probability distributions:
 - Likelihood: $W \sim \text{Binomial}(N, p)$, where N is the total number of tosses.
 - Prior $p \sim \text{Uniform}(0, 1)$
 - “ \sim ” means a random variable drawn from a probability distribution.
- We use the likelihood and the prior to calculate the posterior.
- We can't easily do this with analytical math using pencil and paper.
 - Computational methods:
 - Grid approximation
 - Quadratic approximation
 - Monte Carlo sampling

Examples




Grid Approximation

1. Define a grid:
 - specify a number of points to sample your function at.
 - Take evenly spaced values for each parameter (e.g., the proportion of water).
 - This example uses one parameter and a one-dimensional grid for simplicity.
 - For models with more than one parameter, the grid has 2, 3, or more dimensions—one per parameter.
2. Calculate the value of the *prior* at each grid point
3. Calculate the *likelihood* at each grid point
4. Compute an *unstandardized posterior* by multiplying the *prior* and *posterior* at each grid point.
5. Finally, standardize the *posterior* by dividing each value by the sum of all values in the *unstandardized posterior*.
- The more grid points you use, the more accurate your estimate will be, but the more computer power you'll need.
 - For one parameter, a 1000 point grid is simple.
 - For 2 parameters, a 1000 point grid for each of them means 1 million points.
 - For 3 parameters, it means 1 billion points.
 - For 30 parameters, your grid would have more points than there are atoms in the universe.

Quadratic Approximation

- Focus our attention near the part of the distribution that has the highest probability density.
- This region looks a lot like a Gaussian (normal) distribution.
 1. Find the posterior mode (the highest value).
 - Hill-climbing algorithms
 2. Estimate the curvature of the posterior near the mode
 - Approximate probability density as a Gaussian
 - Approximate the logarithm of the probability density as a quadratic function.
 3. We can calculate the integral of a Gaussian easily.



Programming Models in R

Grid approximation

- Sample functions on a regular grid and approximate integrals by the sum of the samples.

```
# define number of points in the grid
grid_points = 200


# define grid
p_grid <- seq(from = 0, to = 1,
              length.out = grid_points)

# define prior
prior <- rep(1, grid_points)

# compute likelihood at each value in grid
likelihood <- dbinom(6, size = 9, prob = p_grid)

# compute product of likelihood and prior
unstd.posterior <- likelihood * prior

# standardize the posterior, so it sums to 1
posterior <- unstd.posterior / sum(unstd.posterior)
```



Quadratic Approximation

```
library(rethinking)

globe_qa <- quap(
  alist(
    W ~ dbinom( W + L, p),  # binomial likelihood
    p ~ dunif(0, 1)         # uniform prior
  ),
  data = list(W = 6, L = 3)
)
```

Now display a summary:

```
precis(globe_qa)
```

```
##           mean           sd        5.5%        94.5%
## p 0.6666761 0.1571315 0.4155496 0.9178027
```