

# Rethinking Statistics

EES 5891-03

Bayesian Statistical Methods

Jonathan Gilligan

Class #2: Tuesday, August 30 2022

# Golems and Rethinking

# Statistical Tools as Golems

- Statistical tools will do what you tell them to do ...
  - but if you're not careful, what you tell them to do may not be what you want them to do
- The goal of this book is to help you:
  - Learn to use statistical golems wisely
  - Learn to choose the right golem for the job
  - Learn to engineer your own golems if the ready-to-use golems aren't right for your job.




The Golem of Prague (Photo: Prague Post)



# Hypothesis Testing

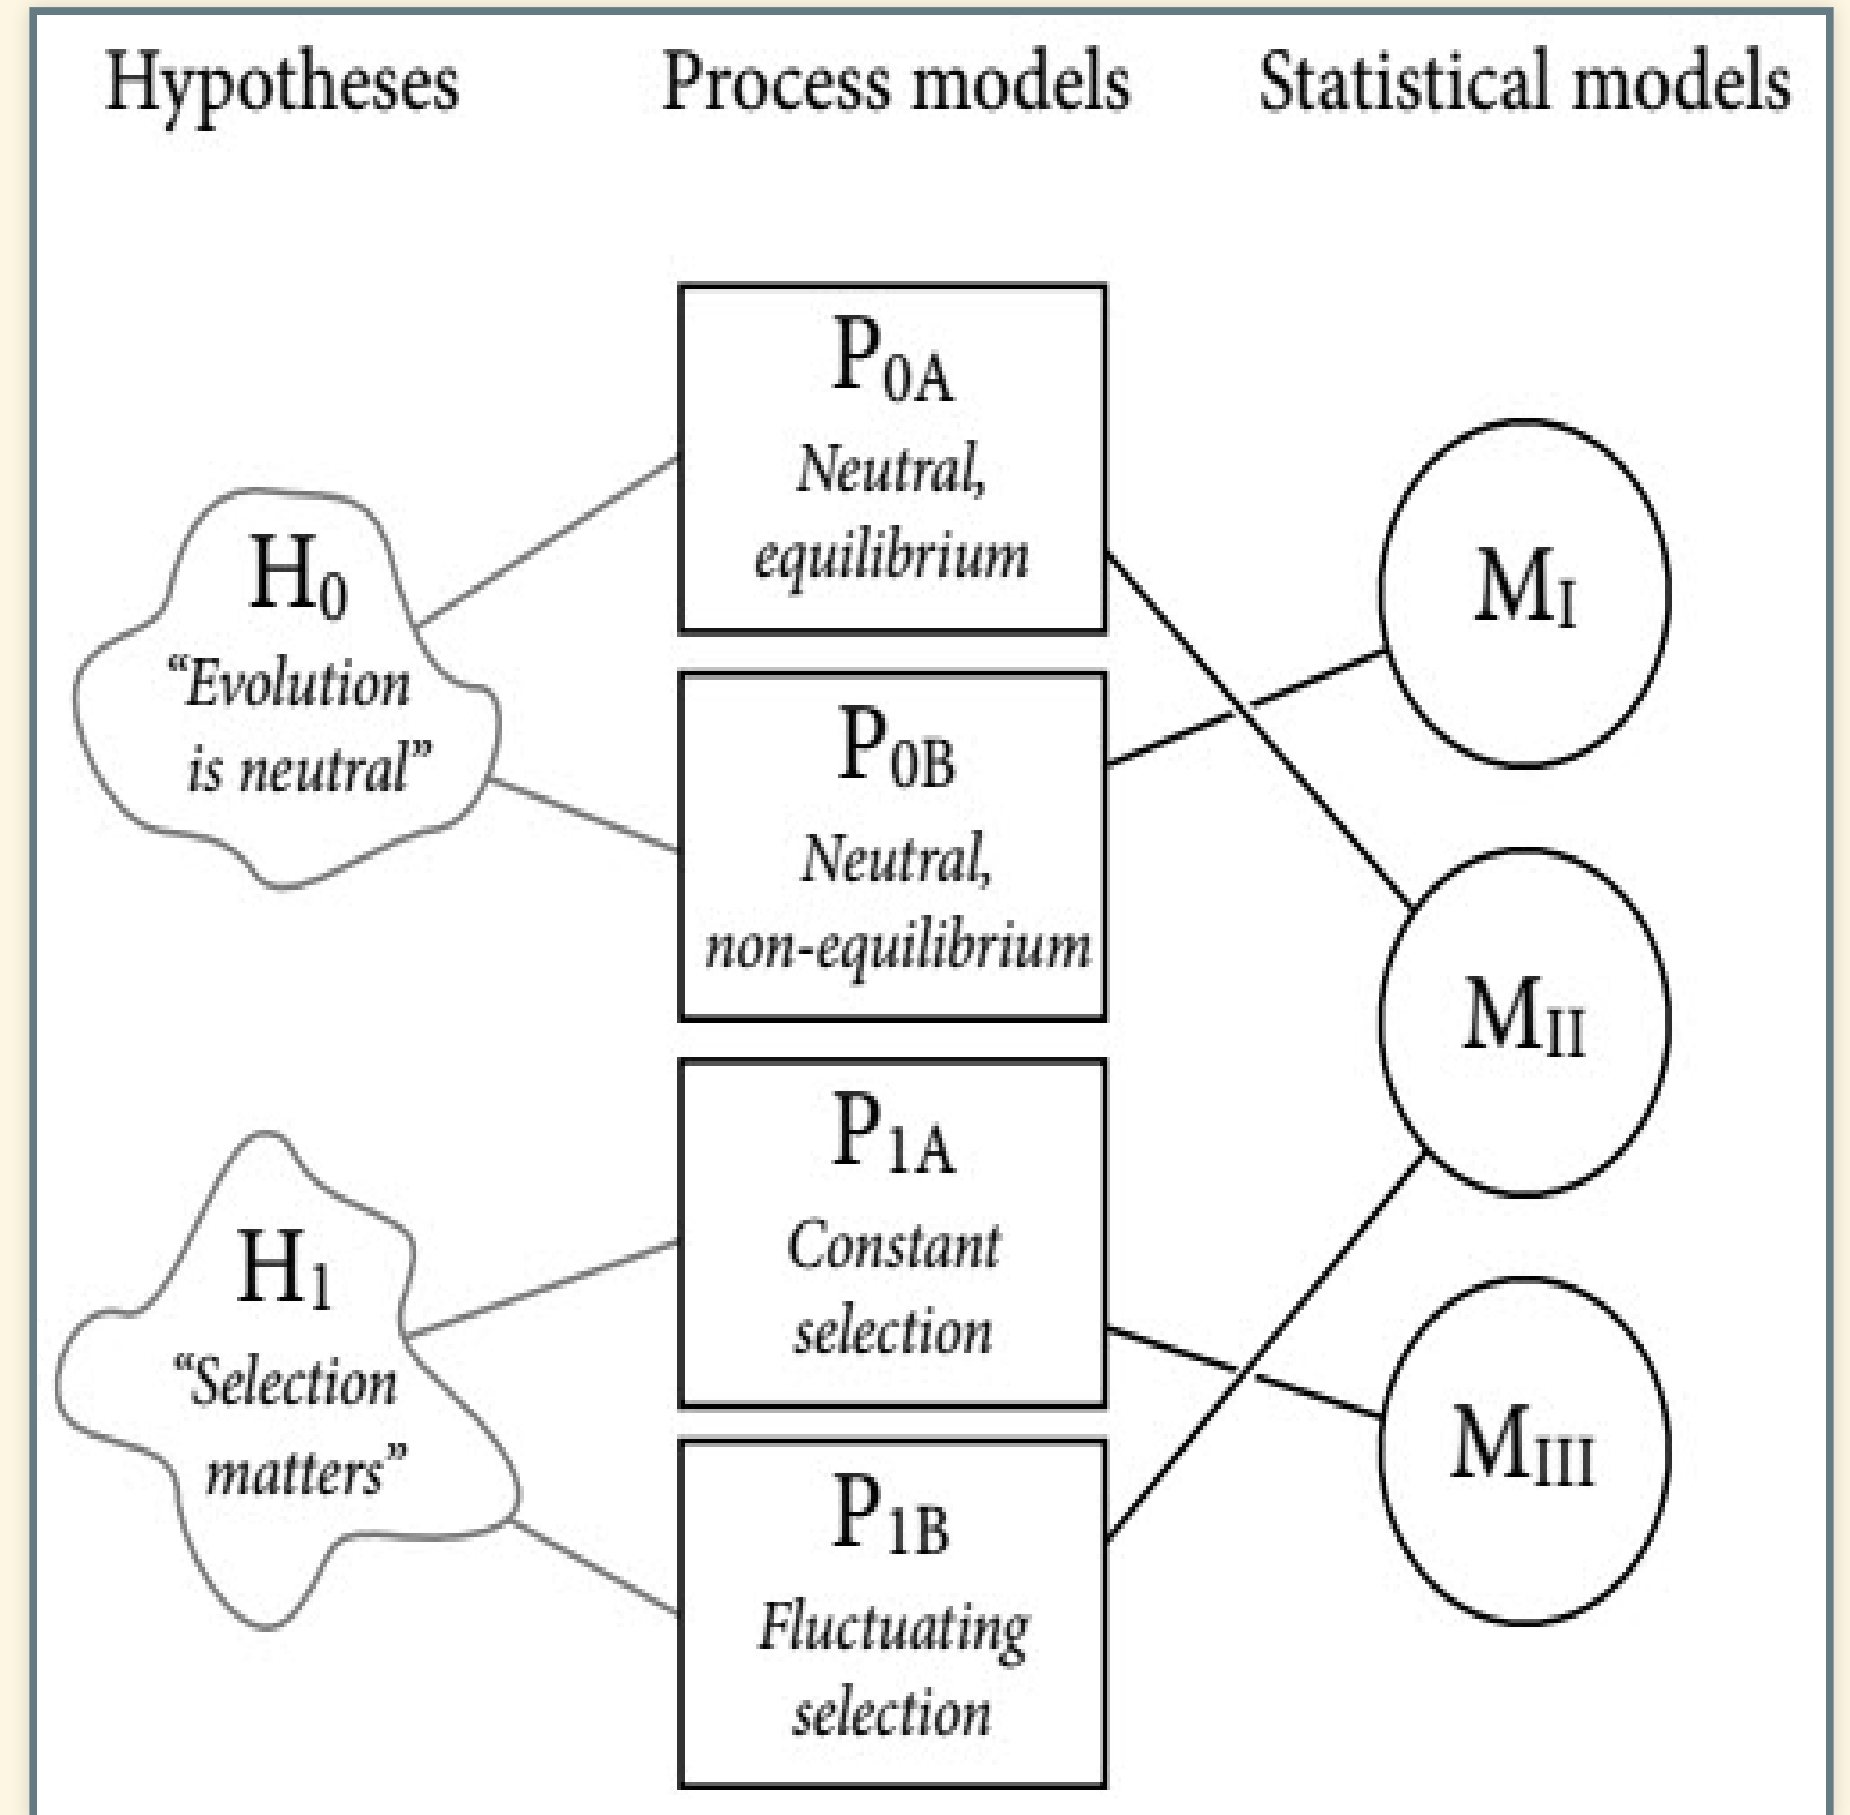
- Karl Popper (1902–1994)
  - Science can never prove that a hypothesis is true
  - But it can prove that an incorrect hypothesis is false
  - The more false hypotheses we rule out, we narrow down the list of potentially true hypotheses.

*When you have eliminated the impossible,  
whatever remains, however improbable,  
must be the truth*  
— Arthur Conan Doyle/Sherlock Holmes



# Problems with falsification

- The predictions of a hypothesis may not be as clear as many people assume.
- Depending on what other assumptions you make, two different hypotheses may predict the same kind of data.
  - If your data looks like  $M_I$ , it rules out (falsifies)  $H_1$
  - But if your data looks like  $M_{II}$ , it doesn't rule out either hypothesis.
- A given hypothesis may predict many different possible kinds of data, depending on what other assumptions you make.
  - If your data doesn't look like  $M_I$ , that doesn't imply that it's less likely  $H_1$  is true.



# Null-Hypothesis Significance Testing

- Most statistical tests aim to rule out a *null hypothesis*, not to falsify the actual research hypothesis.
- Often, there's not one unique alternative hypothesis to the null hypothesis, so even if we reject the null hypothesis, there are many other possibilities.
- Bayesian methods give us better, more powerful golems to answer the questions we're really interested in.
  - But they're still golems and we have to be thoughtful and careful about how to use them.

# Discussion

# Bayes's Theorem



# Bayes's Theorem

- The core of this part of the course is Bayes's theorem.
- Notation:
  - Conditional probability:  $P(a | b)$  means *the probability of a, given b*.
- Bayes's theorem:  $P(H | D) = \frac{P(D | H) \times P(H)}{P(D)}$ , where
  - $P(H | D)$  is the *posterior*: The probability that  $H$  is true, given that you observed  $D$ .
  - $P(D | H)$  is the *likelihood*: The probability that you would observe  $D$ , if  $H$  is true.
  - $P(H)$  is the *prior* probability of  $H$ , based on what you knew before observing  $D$
  - $P(D)$  is the *evidence*: The probability that you would observe  $D$ , regardless whether  $H$  is true.
    - If  $H$  is binary (true or false), then  $P(D) = P(D | H) \times P(H) + P(D | \text{not } H) \times (1 - P(H))$


# Bayes's Theorem (cont)

- We can apply Bayes's theorem to a numbers too. For a variable  $x$  that we want to predict:  $\left[ P(x | D) = \frac{P(D | x) \times P(x)}{P(D)} \right]$
- In this case,  $\left[ P(D) = \int_{-\infty}^{\infty} P(D | x) \times P(x) \mathrm{d}x, \right]$  where the *prior*  $\left( P(x) \right)$  is a probability density function.

# Sampling

# Sampling


- You have a globe and want to figure out what fraction of the earth's surface is water.
- Toss the globe in the air, catch it, and note whether your index finger is on water or land: outcomes are  $W$  and  $L$ .
- At every toss, use Bayes's theorem to update your estimate of the fraction that is water.



# First toss

- Before you toss the globe, pick a prior probability distribution for the fraction that's water.
- Suppose we don't know anything.
  - Pick  $(p \sim \text{Uniform}(0,1))$ , a uniform prior:
- Toss the globe and your finger lands on water.
- Update the probability:  $[ P(p | W) = \frac{P(W | p) P(p)}{p(W)}, ]$  where  $p$  is the probability of water, and  $W$  is measuring water.

Prior:






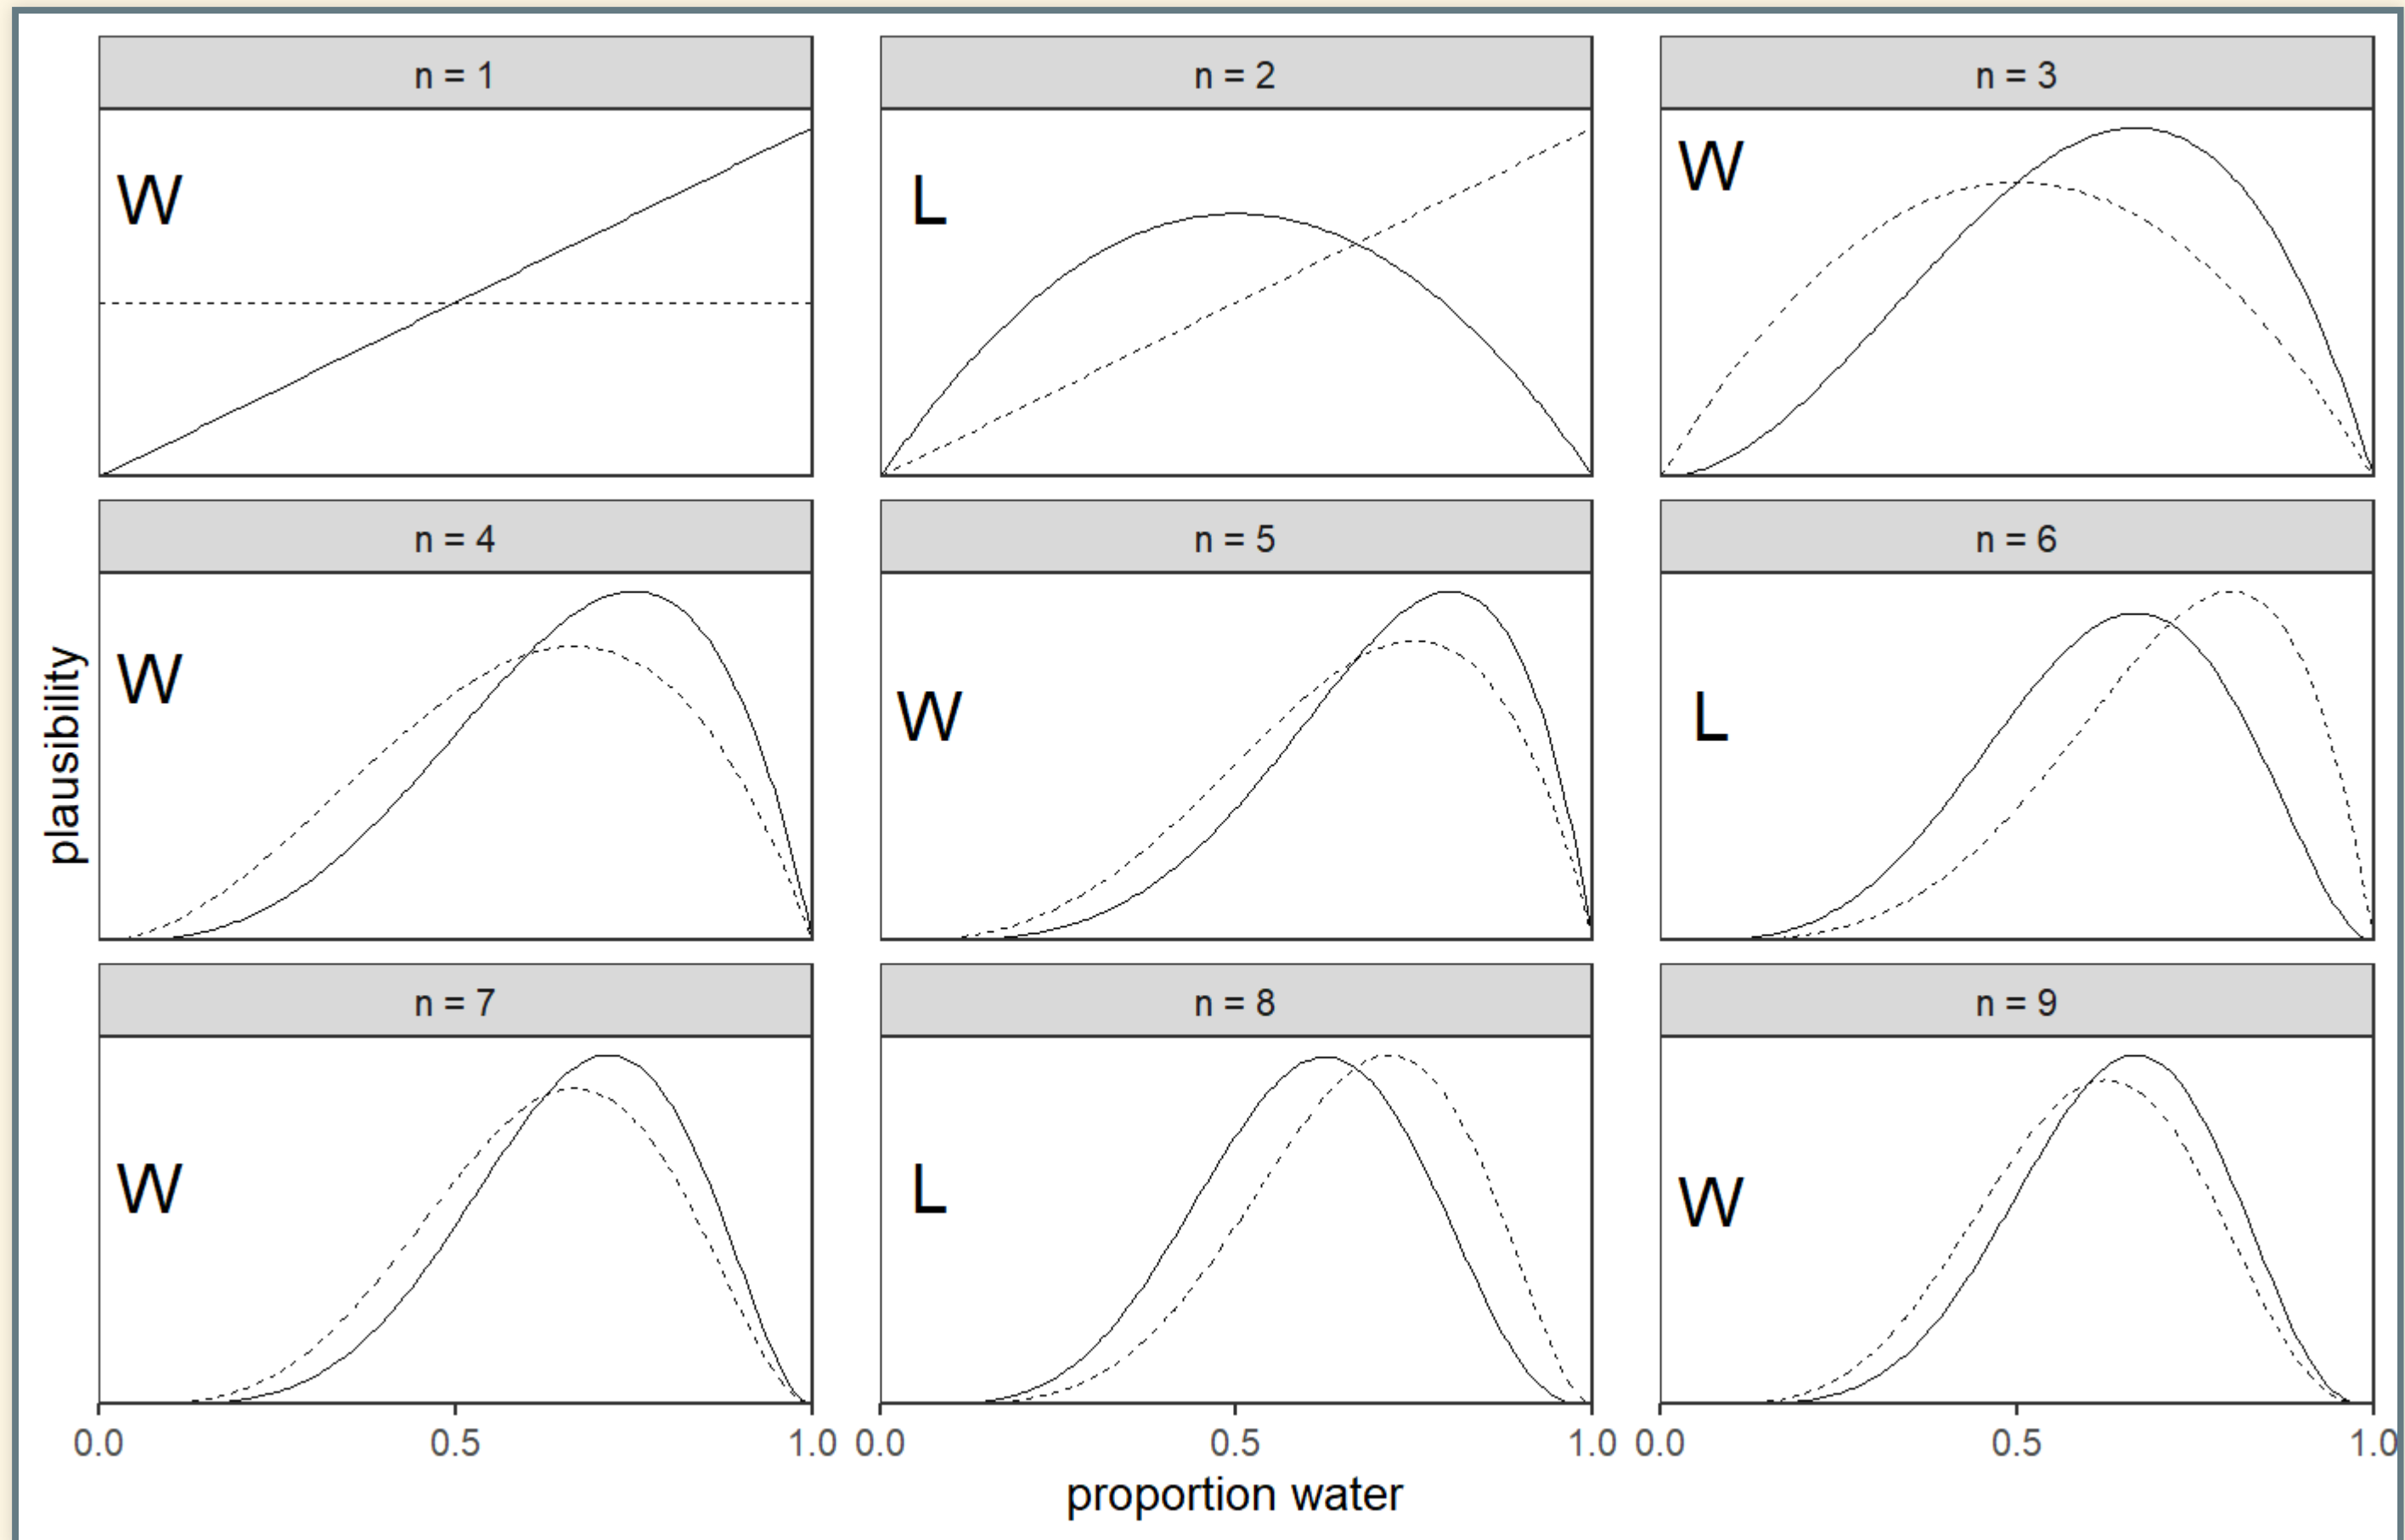
# The calculation:

$$P(p | W) = \frac{P(W | p) P(p)}{P(W)}$$

- $P(W | p) = p$
- $P(p) = 1$  (it's a uniform distribution)  
$$P(W) = \int_0^1 P(W | p) P(p) \mathrm{d}p = \int_0^1 p \times 1 \mathrm{d}p = \left. \frac{p^2}{2} \right|_0^1 = \frac{1}{2}$$
- so the posterior  $P(p | W) = 2p$
- Use this posterior as the prior for the next toss...



# Subsequent tosses

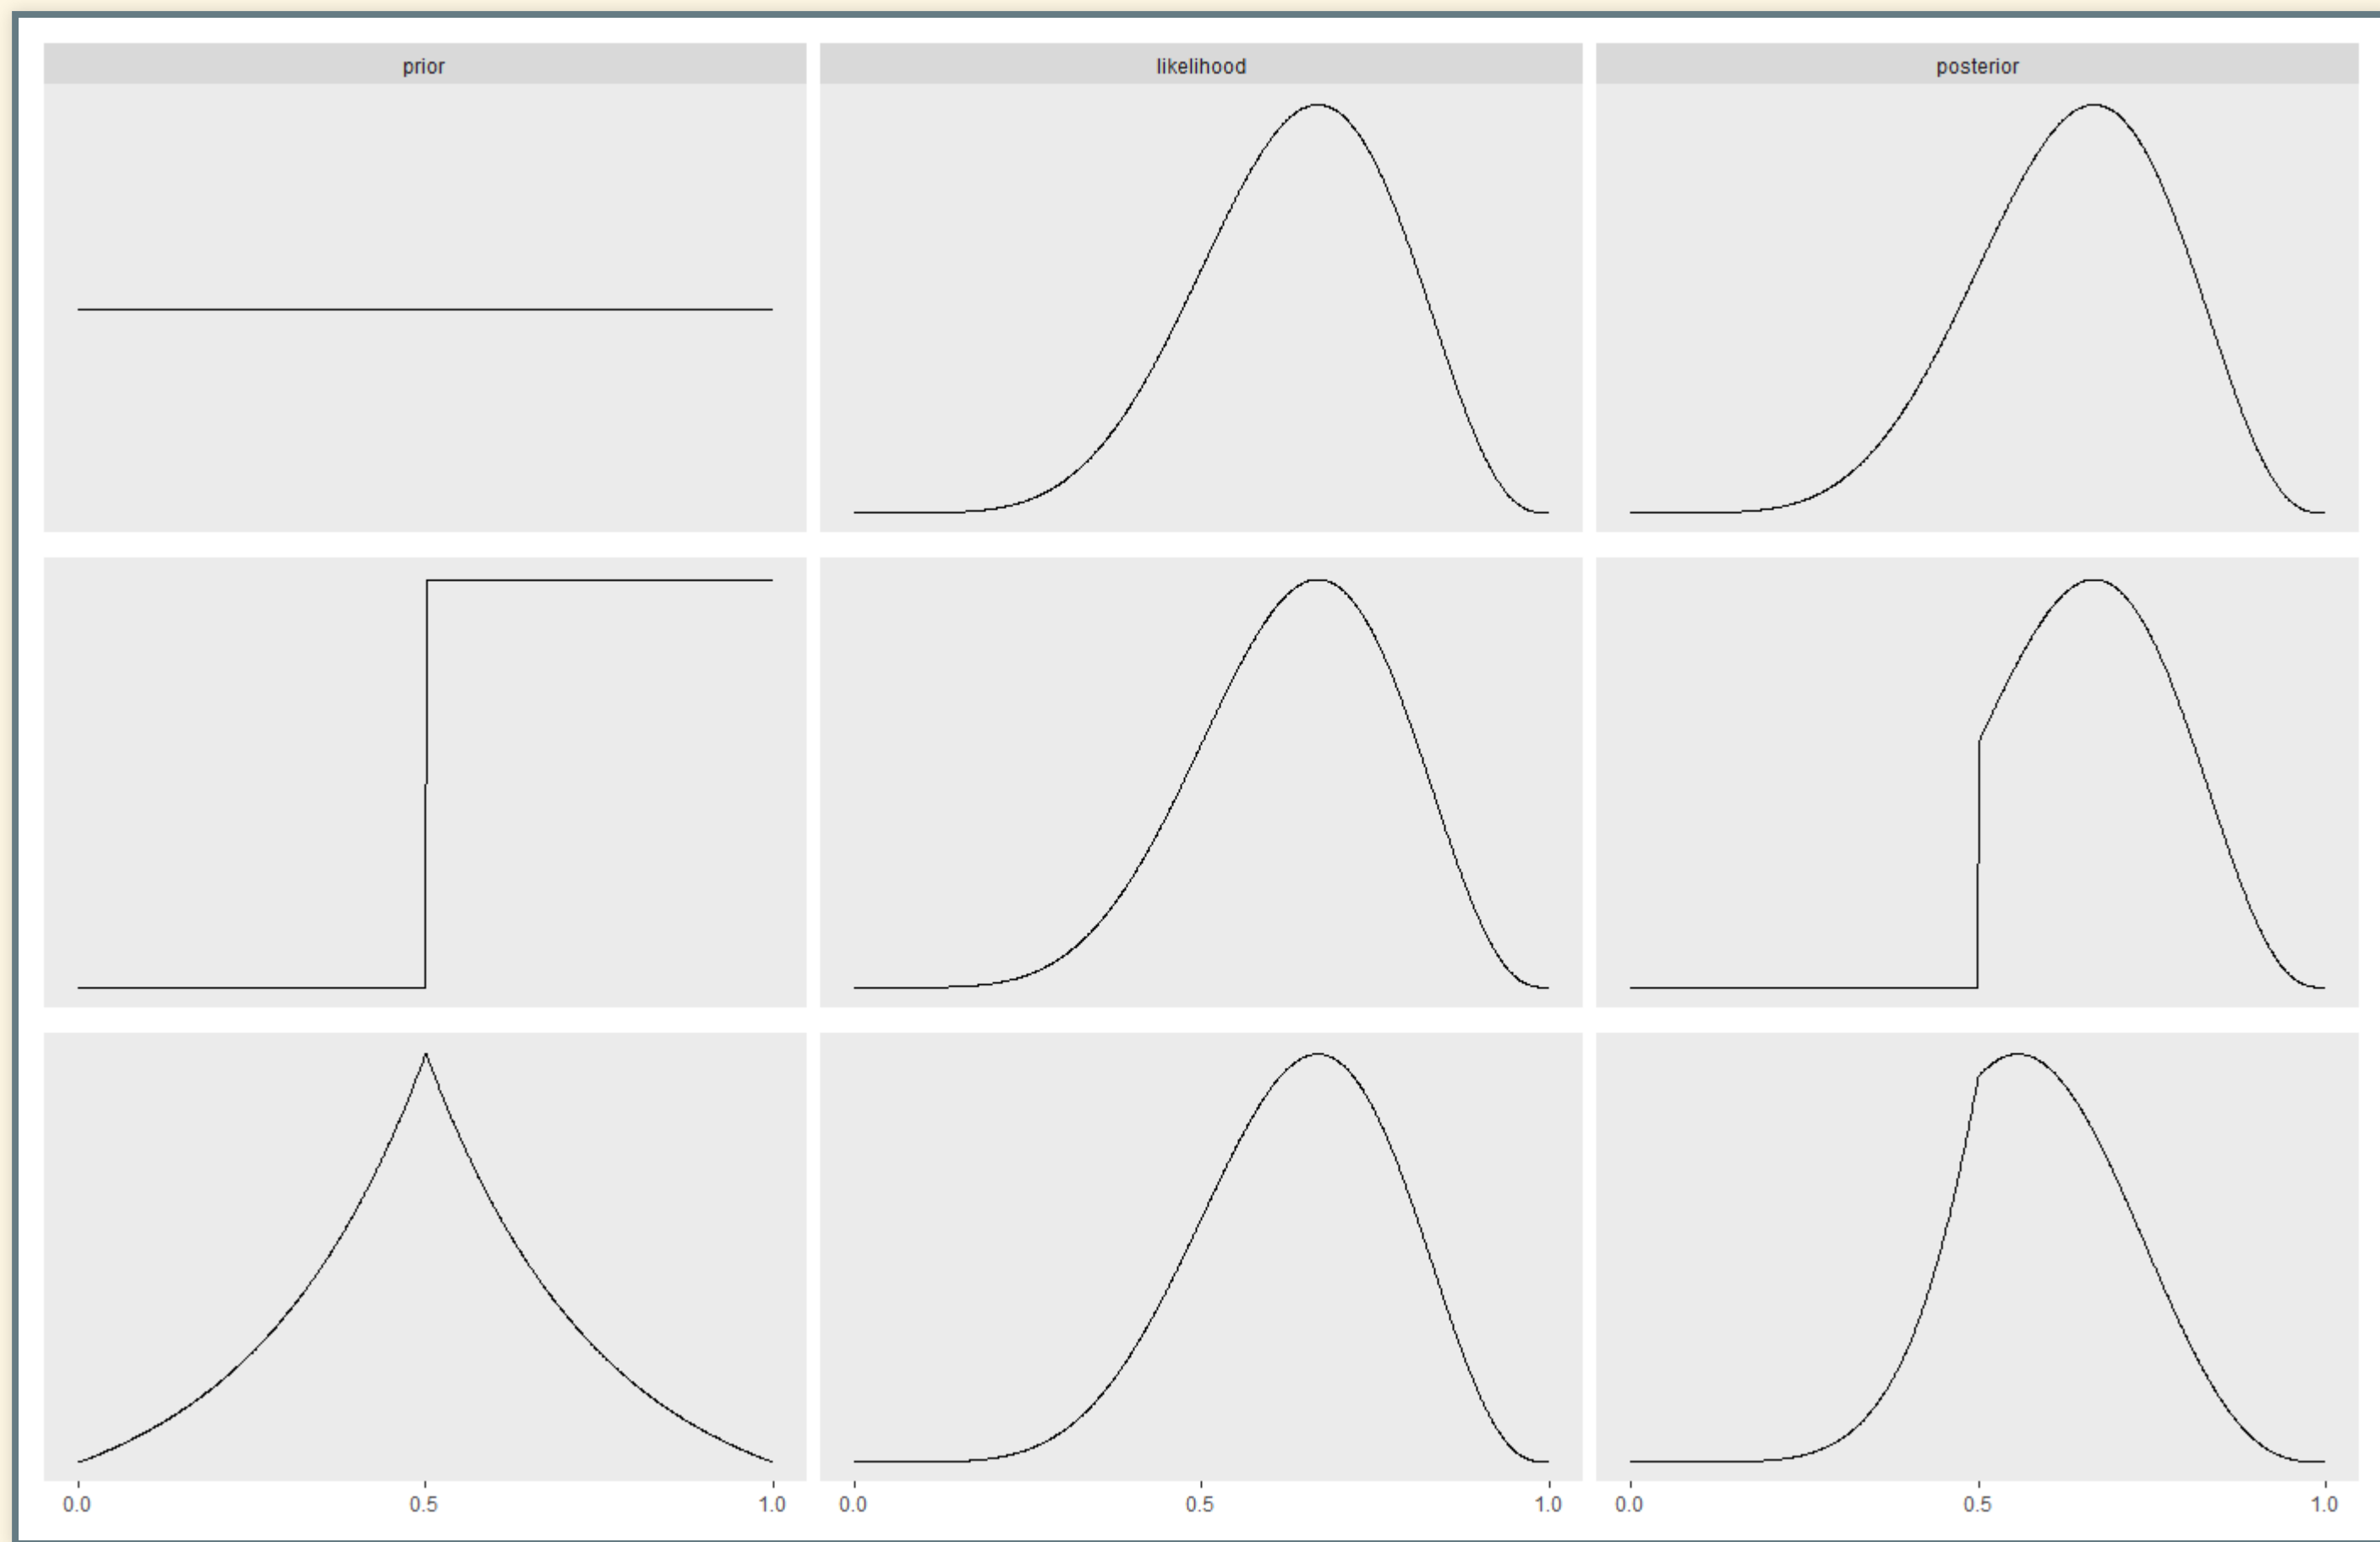


# Developing a Model

# Developing a Model

- Observations and parameters are drawn from probability distributions:
  - Likelihood:  $(W \sim \text{Binomial}(N, p))$ , where  $(N)$  is the total number of tosses.
  - Prior  $(p \sim \text{Uniform}(0,1))$
  - “ $(\sim)$ ” means a random variable drawn from a probability distribution.
- We use the likelihood and the prior to calculate the posterior.
- We can't easily do this with analytical math using pencil and paper.
  - Computational methods:
    - Grid approximation
    - Quadratic approximation
    - Monte Carlo sampling

# Examples






# Grid Approximation

1. Define a grid:
  - specify a number of points to sample your function at.
  - Take evenly spaced values for each parameter (e.g., the proportion of water).
    - This example uses one parameter and a one-dimensional grid for simplicity.
    - For models with more than one parameter, the grid has 2, 3, or more dimensions—one per parameter.
2. Calculate the value of the *prior* at each grid point
3. Calculate the *likelihood* at each grid point
4. Compute an *unstandardized posterior* by multiplying the *prior* and *likelihood* at each grid point.
5. Finally, standardize the *posterior* by dividing each value by the sum of all values in the *unstandardized posterior*.
- The more grid points you use, the more accurate your estimate will be, but the more computer power you'll need.
  - For one parameter, a 1000 point grid is simple.
  - For 2 parameters, a 1000 point grid for each of them means 1 million points.
  - For 3 parameters, it means 1 billion points.
  - For 30 parameters, your grid would have more points than there are atoms in the universe.

# Quadratic Approximation

- Focus our attention near the part of the distribution that has the highest probability density.
- This region looks a lot like a Gaussian (normal) distribution.
  1. Find the posterior mode (the highest value).
    - Hill-climbing algorithms
  2. Estimate the curvature of the posterior near the mode
    - Approximate probability density as a Gaussian
    - Approximate the logarithm of the probability density as a quadratic function.
  3. We can calculate the integral of a Gaussian easily.



# Programming Models in R

# Grid approximation

- Sample functions on a regular grid and approximate integrals by the sum of the samples.

```
# define number of points in the grid
grid_points = 200


# define grid
p_grid <- seq(from = 0, to = 1,
              length.out = grid_points)

# define prior
prior <- rep(1, grid_points)

# compute likelihood at each value in grid
likelihood <- dbinom(6, size = 9, prob = p_grid)

# compute product of likelihood and prior
unstd.posterior <- likelihood * prior

# standardize the posterior, so it sums to 1
posterior <- unstd.posterior / sum(unstd.posterior)
```



# Quadratic Approximation

```
library(rethinking)

globe_qa <- quap(
  alist(
    W ~ dbinom( W + L, p),  # binomial likelihood
    p ~ dunif(0, 1)         # uniform prior
  ),
  data = list(W = 6, L = 3)
)
```

Now display a summary:

```
precis(globe_qa)
```

```
##           mean           sd        5.5%        94.5%
## p 0.6666761 0.1571315 0.4155496 0.9178027
```