

Hamiltonian Monte Carlo with Stan and Ulam

EES 5891-03

Bayesian Statistical Methods

Jonathan Gilligan

Class #12: Thursday, October 06 2022

Hamiltonian Monte Carlo Sampling

Prepare data

- `rugged` data set:
 - Each row is a different country
 - `rugged`: Topographic roughness of the terrain
 - `rpcgdp_2000`: Per-capita GDP in 2000
- Use the logarithm of GDP.
- quasi-standardize the variables.
- Model
$$\begin{aligned} G &\sim \text{Normal}(\mu, \sigma) \\ \mu &= \alpha_{\text{CID}} + \beta_{\text{CID}} R \\ \alpha_{\text{CID}} &\sim \text{Normal}(1, 0.1) \\ \beta_{\text{CID}} &\sim \text{Normal}(0, 0.3) \\ \sigma &\sim \text{Exponential}(1) \end{aligned}$$

```
library(tidyverse)
library(rethinking)

data(rugged)
d <- rugged

dd <- d %>% filter(complete.cases(rgdppc_2000)) %>%
  mutate(
    log_gdp = log(rgdppc_2000),
    log_gdp_std = log_gdp / mean(log_gdp),
    rugged_std = rugged / max(rugged),
    cid = ifelse(cont_africa == 1, 1, 2)
  )
```

```
lst_rug <- alist(
  log_gdp_std ~ dnorm(mu, sigma),
  mu <- a[cid] + b[cid] * (rugged_std - 0.215),
  a[cid] ~ dnorm(1, 0.1),
  b[cid] ~ dnorm(0, 0.3),
  sigma ~ dexp(1)
)
```

quap vs. ulam

- **quap** calculates the posterior using a *quadratic approximation*
 - Focus on finding the peak of the posterior (maximum *a-posteriori*) for each parameter.
 - Near the peak, the posterior looks like a quadratic function (parabola).
 - Approximate the width
 - Works well with Gaussian (Normal) distributions, not so well with other functions.
- **ulam** calculates the posterior using a *Hamiltonian Monte Carlo* sampler.
 - Generates samples from the posterior.
 - Doesn't directly calculate the *maximum a-posteriori*.
- Use statistical functions to analyze the posterior (mean, median, standard deviation, etc.)
- Runs slower than **quap** but works with a greater variety of prior and likelihood functions.

ulam and stan

```
lst_rug <- alist(  
  log_gdp_std ~ dnorm(mu, sigma),  
  mu <- a[cid] + b[cid] * (rugged_std - 0.215),  
  a[cid] ~ dnorm(1, 0.1),  
  b[cid] ~ dnorm(0, 0.3),  
  sigma ~ dexp(1)  
)  
dat_slim <- list(log_gdp_std = dd$log_gdp_std,  
                rugged_std = dd$rugged_std,  
                cid = as.integer(dd$cid))  
mdl_rug_ulam <- ulam(lst_rug, data = dat_slim,  
                    chains = 4, cores = 4)
```

```
## Running MCMC with 4 parallel chains, with 1 thread(s)  
per chain...
```

```
##  
## Chain 1 Iteration:    1 / 1000 [  0%] (Warmup)  
## Chain 1 Iteration:  100 / 1000 [ 10%] (Warmup)  
## Chain 1 Iteration:  200 / 1000 [ 20%] (Warmup)  
## Chain 1 Iteration:  300 / 1000 [ 30%] (Warmup)  
## Chain 1 Iteration:  400 / 1000 [ 40%] (Warmup)  
## Chain 1 Iteration:  500 / 1000 [ 50%] (Warmup)  
## Chain 1 Iteration:  501 / 1000 [ 50%] (Sampling)  
## Chain 1 Iteration:  600 / 1000 [ 60%] (Sampling)  
## Chain 1 Iteration:  700 / 1000 [ 70%] (Sampling)  
## Chain 1 Iteration:  800 / 1000 [ 80%] (Sampling)  
## Chain 1 Iteration:  900 / 1000 [ 90%] (Sampling)  
## Chain 1 Iteration: 1000 / 1000 [100%] (Sampling)
```

Ulam creates this Stan model:

```
cat(mdl_rug_ulam@model)
```

```
## data{  
##     vector[170] log_gdp_std;  
##     vector[170] rugged_std;  
##     int cid[170];  
## }  
## parameters{  
##     vector[2] a;  
##     vector[2] b;  
##     real<lower=0> sigma;  
## }  
## model{  
##     vector[170] mu;  
##     sigma ~ exponential( 1 );  
##     b ~ normal( 0 , 0.3 );  
##     a ~ normal( 1 , 0.1 );  
##     for ( i in 1:170 ) {  
##         mu[i] = a[cid[i]] + b[cid[i]] * (rugged_std[i]  
- 0.215);  
##     }  
##     log_gdp_std ~ normal( mu , sigma );  
## }
```

Running `ulam`

- The model specification is pretty much the same as for `quap`.
- Data must **only** have the observed variables in the model
 - Do transforms (`log()`, `standardize()`) before calling `ulam`.
- When you call `ulam()` it.
 1. Generates `stan` code for the model.
 2. Calls the `stan` compiler to translate the `stan` model into C++.
 3. Compiles the C++ model.
 4. Runs the model with the data you provided.
 - Default: 1000 iterations per chain
 - Half are *warm-up* and half are samples from posterior
 - *warm-up* is where the model tunes itself
 5. Returns the results.

```
lst_rug <- alist(  
  log_gdp_std ~ dnorm(mu, sigma),  
  mu <- a[cid] + b[cid] * (rugged_std - 0.215),  
  a[cid] ~ dnorm(1, 0.1),  
  b[cid] ~ dnorm(0, 0.3),  
  sigma ~ dexp(1)  
)  
mdl_rug_quap <- quap(lst_rug, data = dat_slim)  
dat_slim <- list(log_gdp_std = dd$log_gdp_std,  
                rugged_std = dd$rugged_std,  
                cid = as.integer(dd$cid))  
mdl_rug_ulam <- ulam(lst_rug, data = dat_slim,  
                    chains = 4, cores = 4)
```

Running `ulam`

- `chains`: A *chain* is a sequence of samples.
 - *Markov chain*: Each sample depends on the previous sample.
 - Different chains start from different starting points.
 - *Mixing*: After a certain number of samples, chains that start at different starting points should become indistinguishable.
 - You should sample multiple chains so you can check whether the chains are well-mixed.
- `cores`: How many processor cores to use:
 - When `cores = 1`, the chains execute sequentially
 - When `cores = chains`, all chains execute in parallel (faster).
 - There's no extra benefit in having more `cores` than `chains`.

```
lst_rug <- alist(  
  log_gdp_std ~ dnorm(mu, sigma),  
  mu <- a[cid] + b[cid] * (rugged_std - 0.215),  
  a[cid] ~ dnorm(1, 0.1),  
  b[cid] ~ dnorm(0, 0.3),  
  sigma ~ dexp(1)  
)  
  
dat_slim <- list(log_gdp_std = dd$log_gdp_std,  
                rugged_std = dd$rugged_std,  
                cid = as.integer(dd$cid))  
  
mdl_rug_ulam <- ulam(lst_rug, data = dat_slim,  
                    chains = 4, cores = 4)
```

Interpreting Model Results

Model Output

```
lst_rug <- alist(
  log_gdp_std ~ dnorm(mu, sigma),
  mu <- a[cid] + b[cid] * (rugged_std - 0.215),
  a[cid] ~ dnorm(1, 0.1),
  b[cid] ~ dnorm(0, 0.3),
  sigma ~ dexp(1)
)
dat_slim <- list(log_gdp_std = dd$log_gdp_std,
                rugged_std = dd$rugged_std,
                cid = as.integer(dd$cid))
mdl_rug_ulam <- ulam(lst_rug, data = dat_slim,
                    chains = 4, cores = 4)
```

- For each parameter, **precis** shows:
 - Statistical properties: mean, standard deviation, 5.5% and 94.5% quantiles,
 - **n_eff**: The effective number of samples (smaller than the actual number because of autocorrelation)
 - **Rhat4**: The Gelman-Rubin convergence diagnostic:
 - variance between chains / variance within chain
 - 1 when chains have converged completely
 - ≤ 1.1 indicates good convergence (well-mixed)

```
show(mdl_rug_ulam)
```

```
## Hamiltonian Monte Carlo approximation
## 2000 samples from 4 chains
##
## Sampling durations (seconds):
##           warmup sample total
## chain:1      0.08    0.05    0.13
## chain:2      0.09    0.06    0.15
## chain:3      0.08    0.05    0.13
## chain:4      0.08    0.05    0.13
##
## Formula:
## log_gdp_std ~ dnorm(mu, sigma)
## mu <- a[cid] + b[cid] * (rugged_std - 0.215)
## a[cid] ~ dnorm(1, 0.1)
## b[cid] ~ dnorm(0, 0.3)
## sigma ~ dexp(1)
```

```
precis_show(precis(mdl_rug_ulam, depth = 2, digits = 2))
```

##		mean	sd	5.5%	94.5%	n_eff	Rhat4
##	a[1]	0.89	0.02	0.86	0.91	3016	1
##	a[2]	1.05	0.01	1.03	1.07	2414	1
##	b[1]	0.13	0.08	0.01	0.26	3307	1
##	b[2]	-0.14	0.06	-0.24	-0.05	3287	1
##	sigma	0.11	0.01	0.10	0.12	2306	1

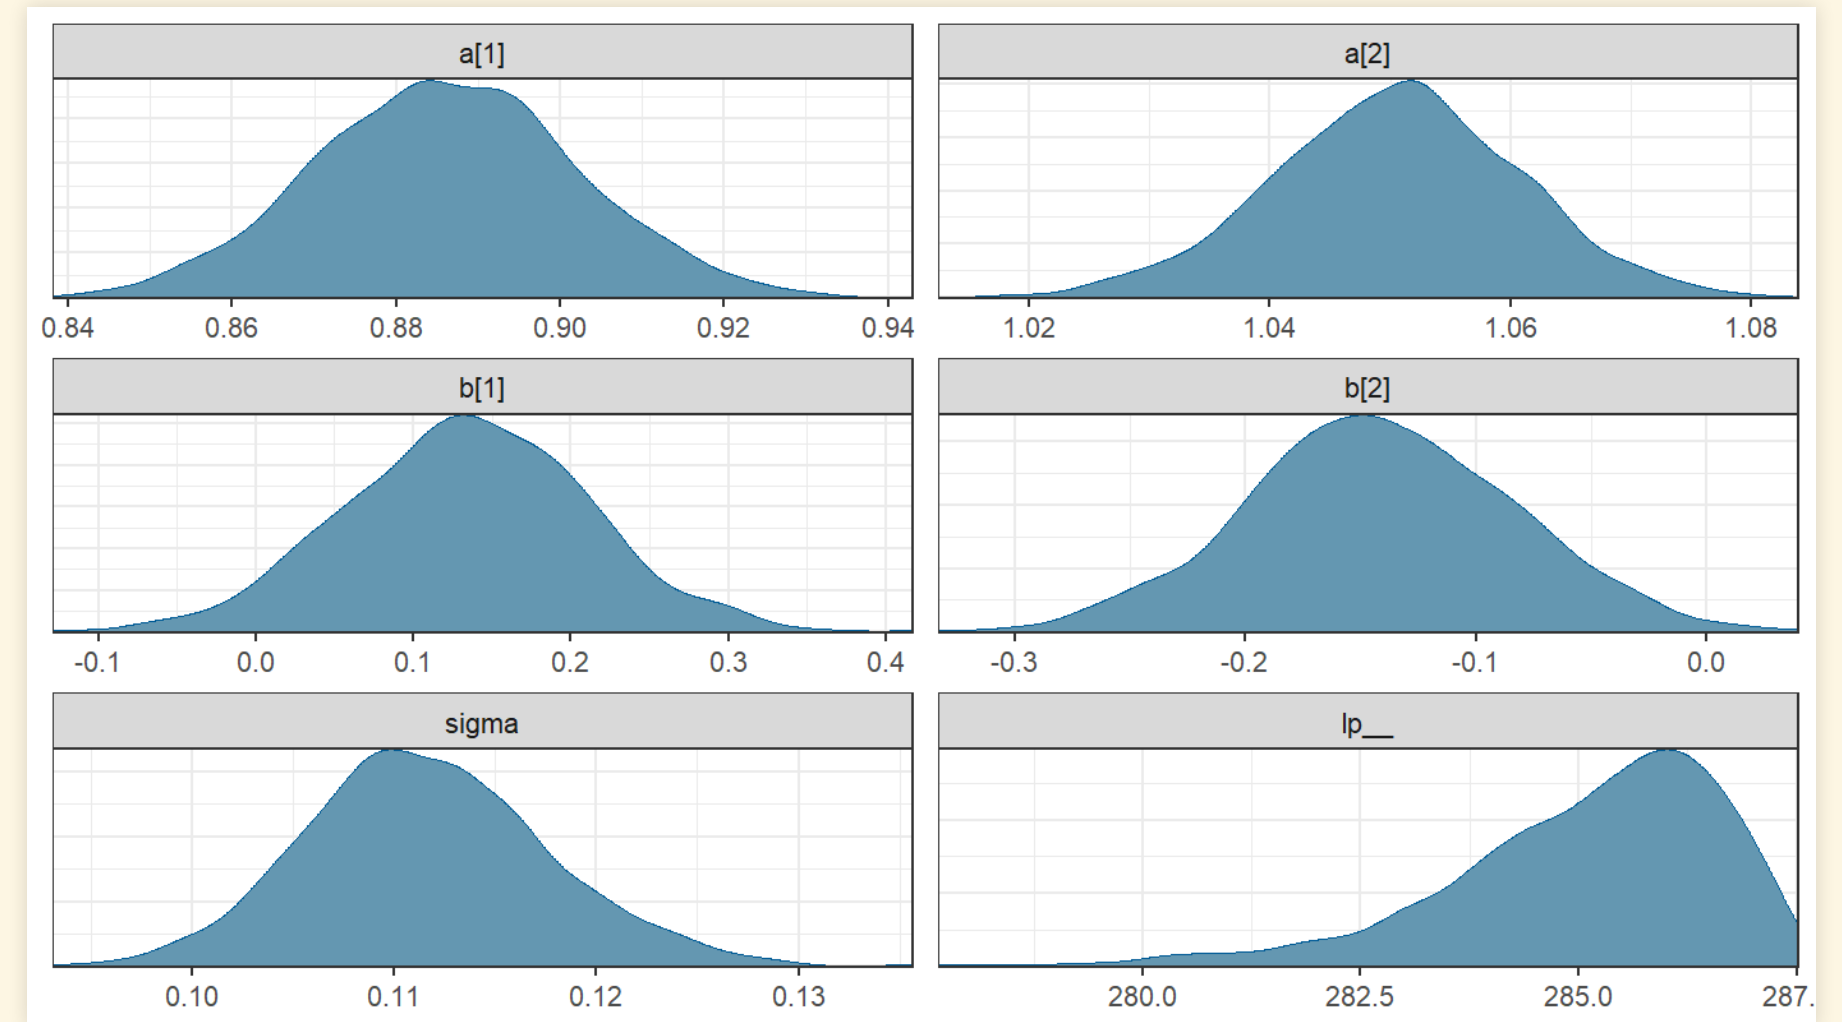
Plotting Model Results

```
library(bayesplot)

mdl_fit <- mdl_rug_ulam@stanfit

mcmc_dens(mdl_fit)
```

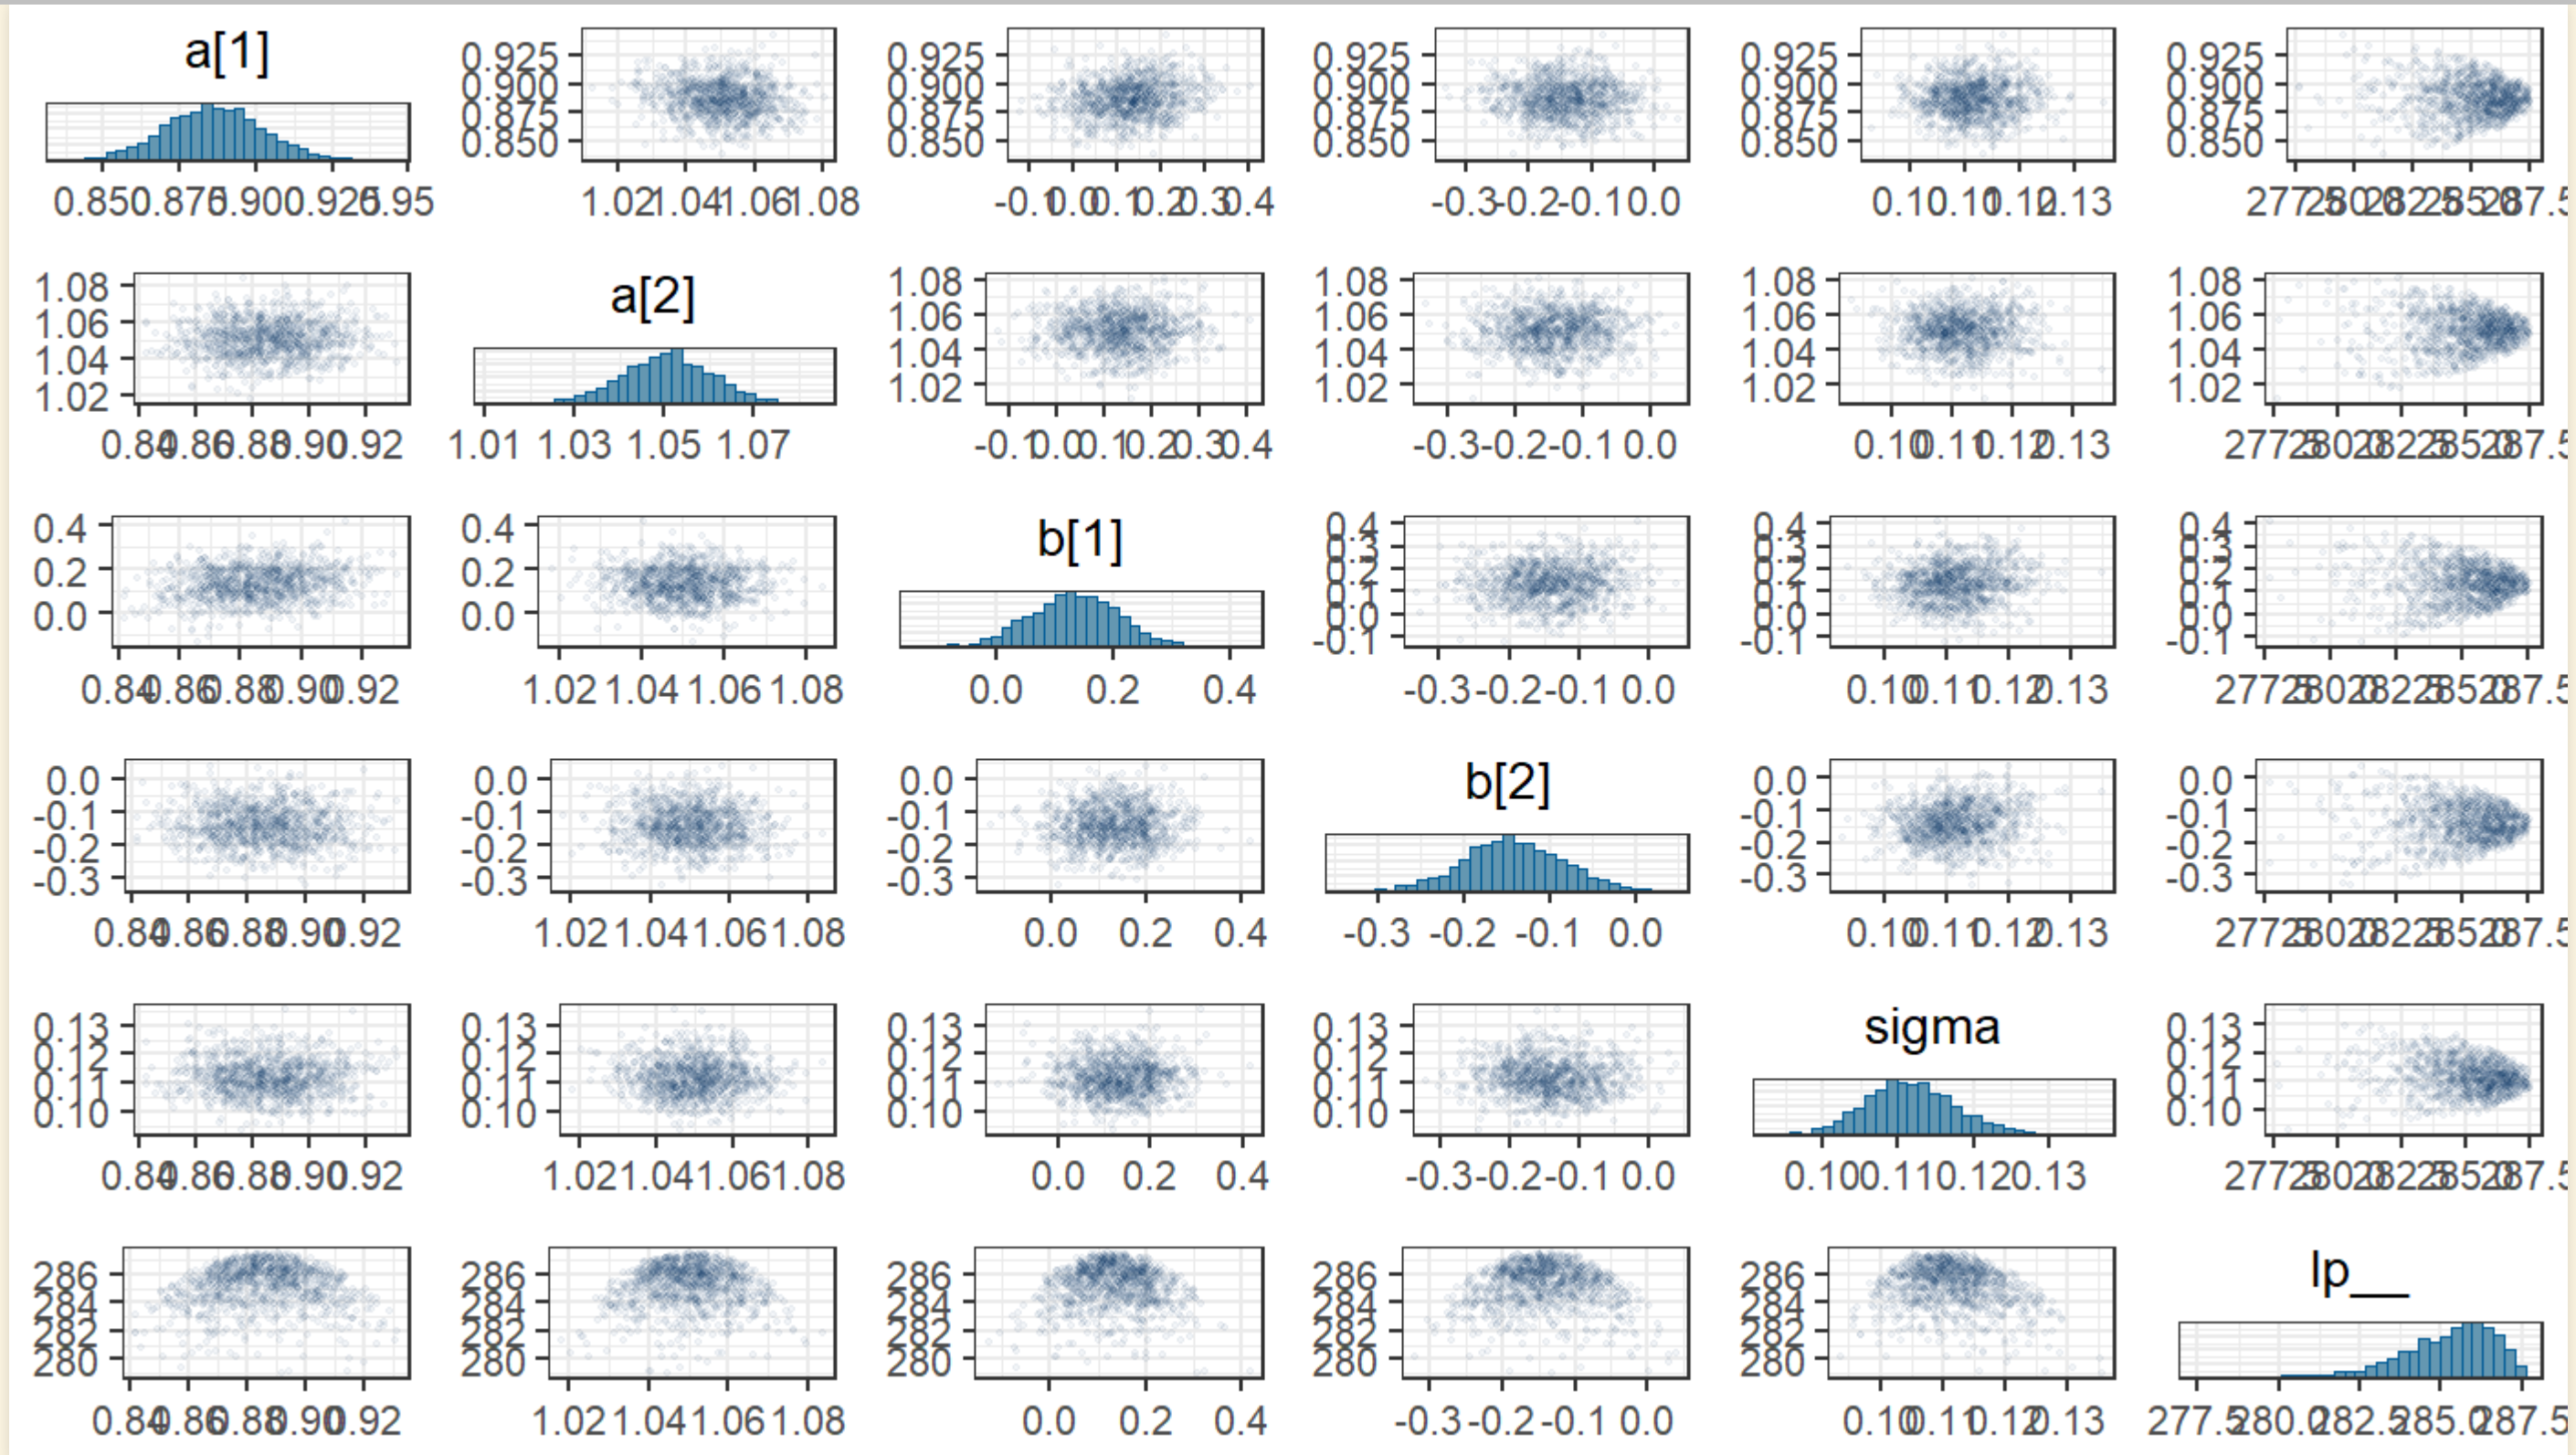
- The graphs show the probability density of each parameter.
- `lp__` is the log of the posterior probability density.
 - This graph shows how the MCMC samples are mostly drawn from higher probability regions of the posterior.



Model Diagnostics

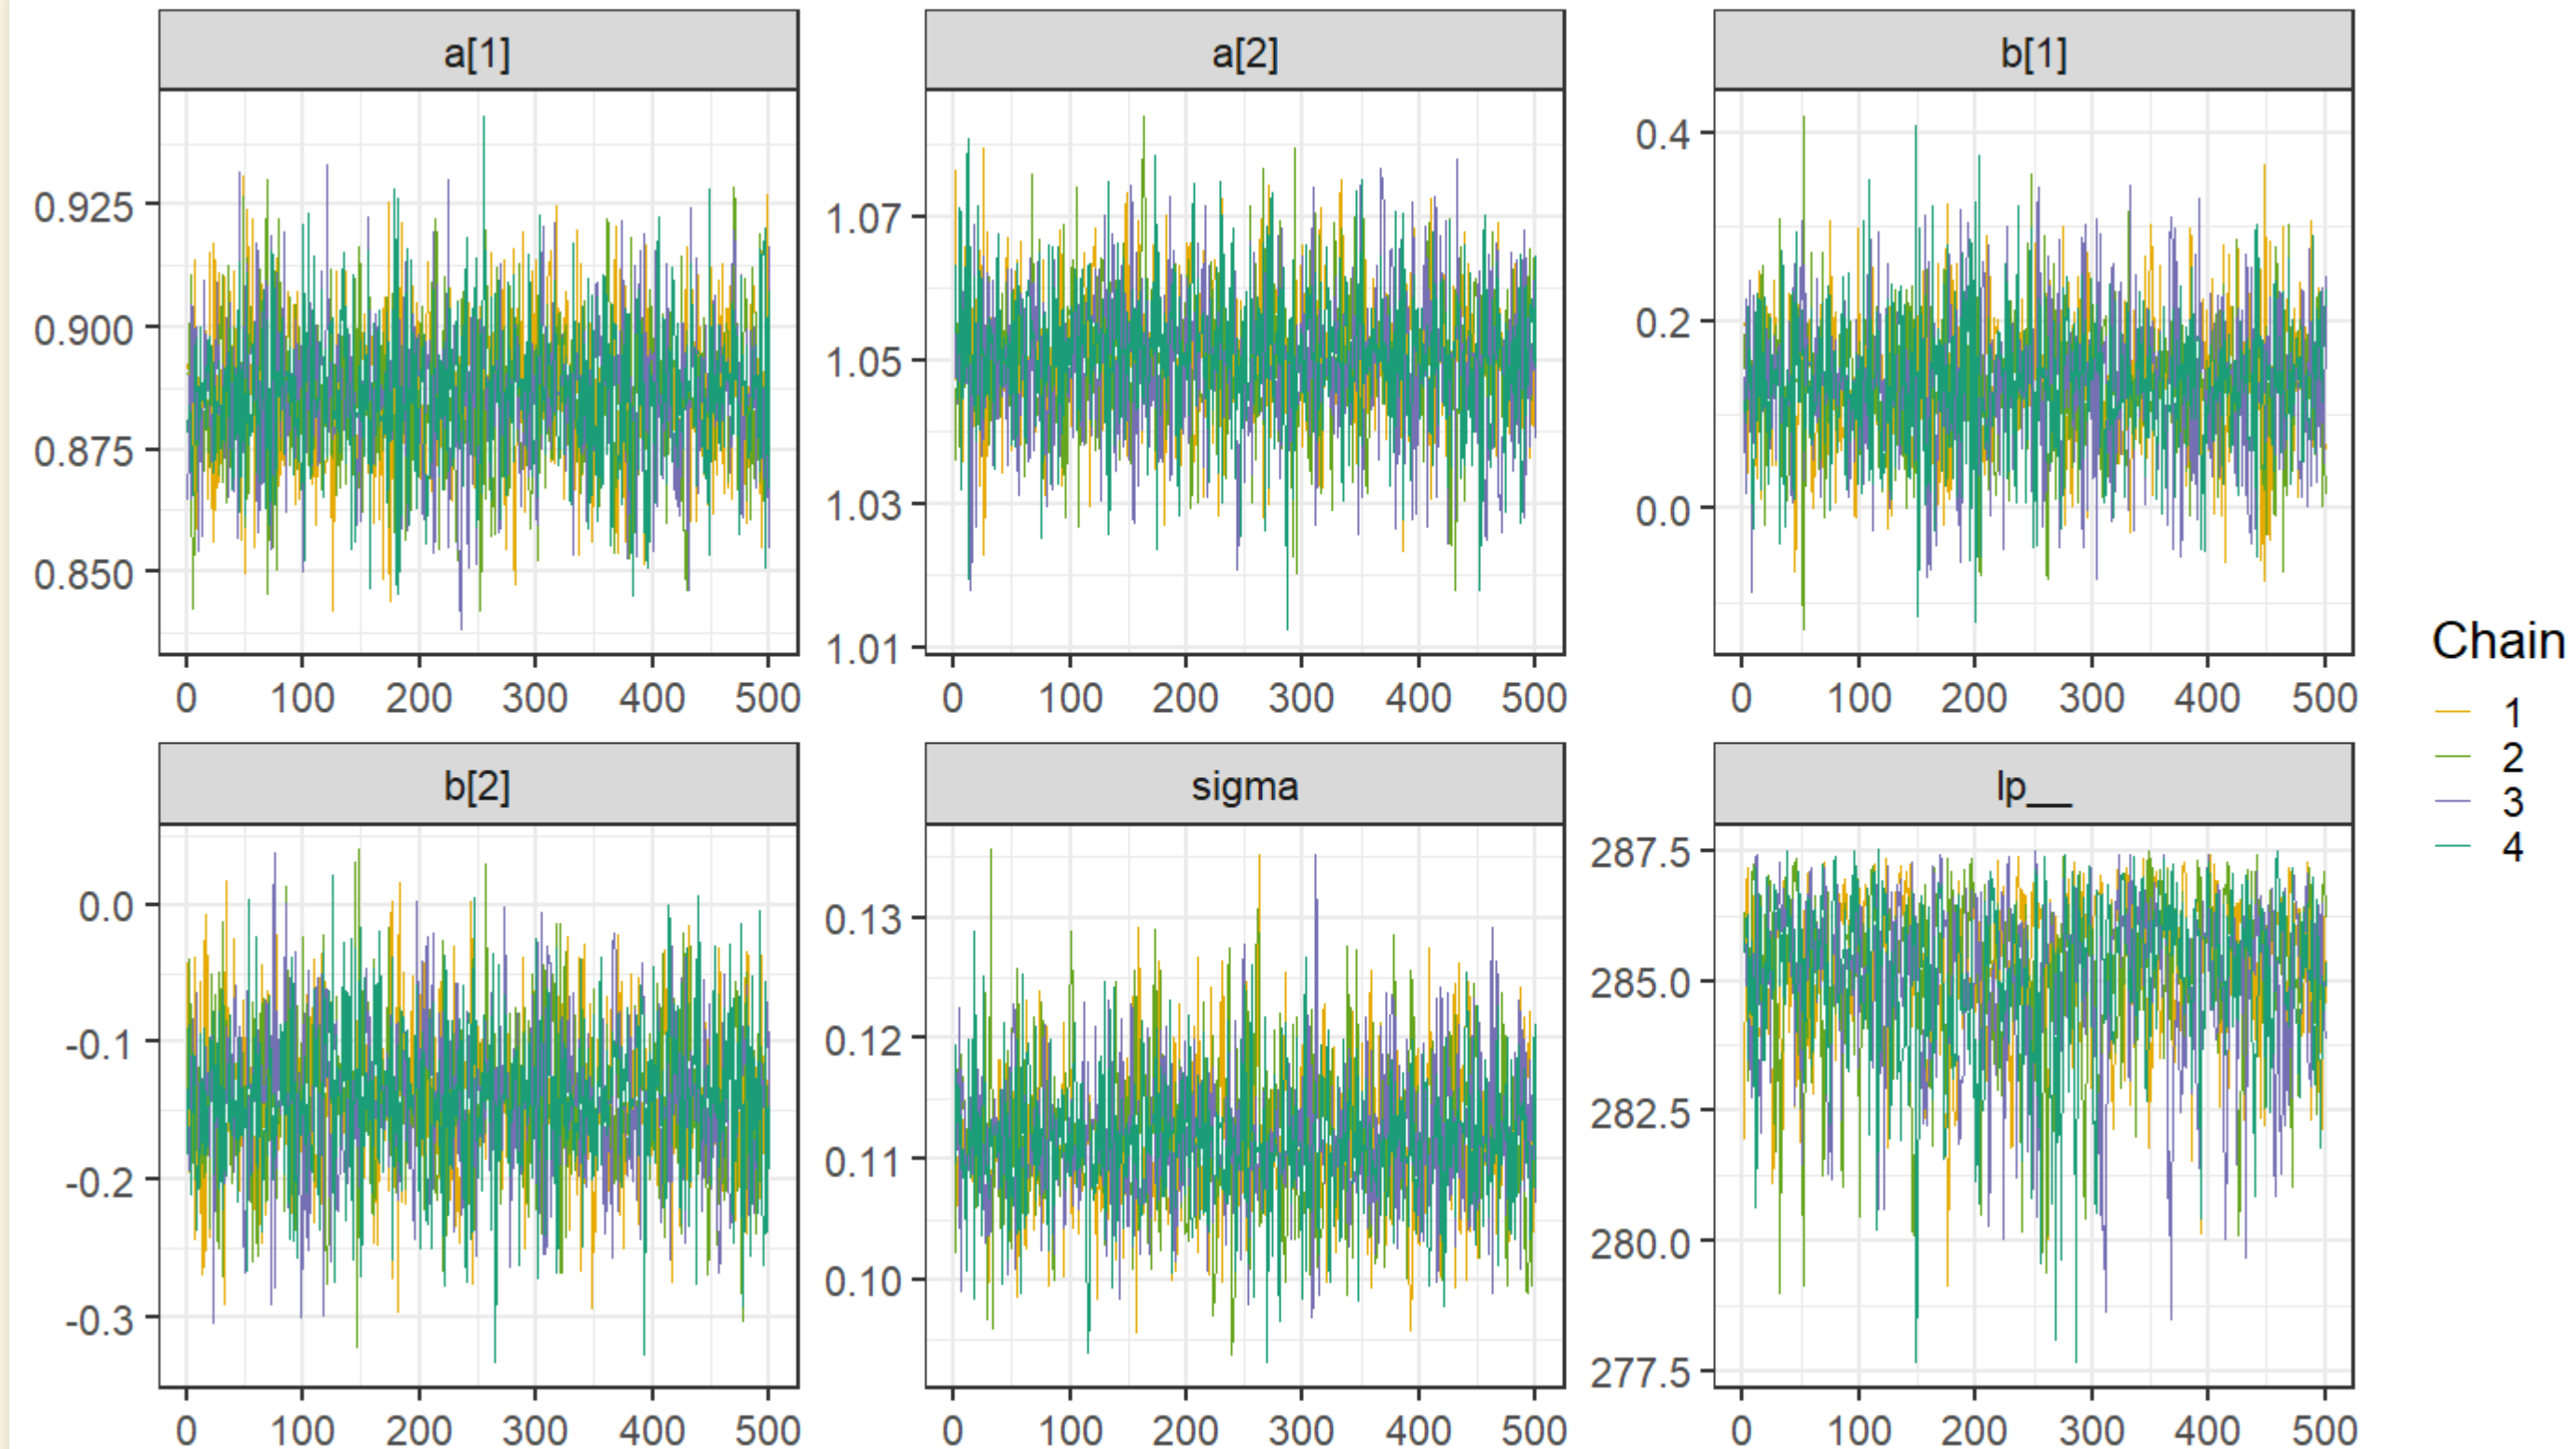
Pair-Correlation Plots

```
color_scheme_set("blue")
mcmc_pairs(mdl_fit, off_diag_args = list(size = 1, alpha = 0.05))
```



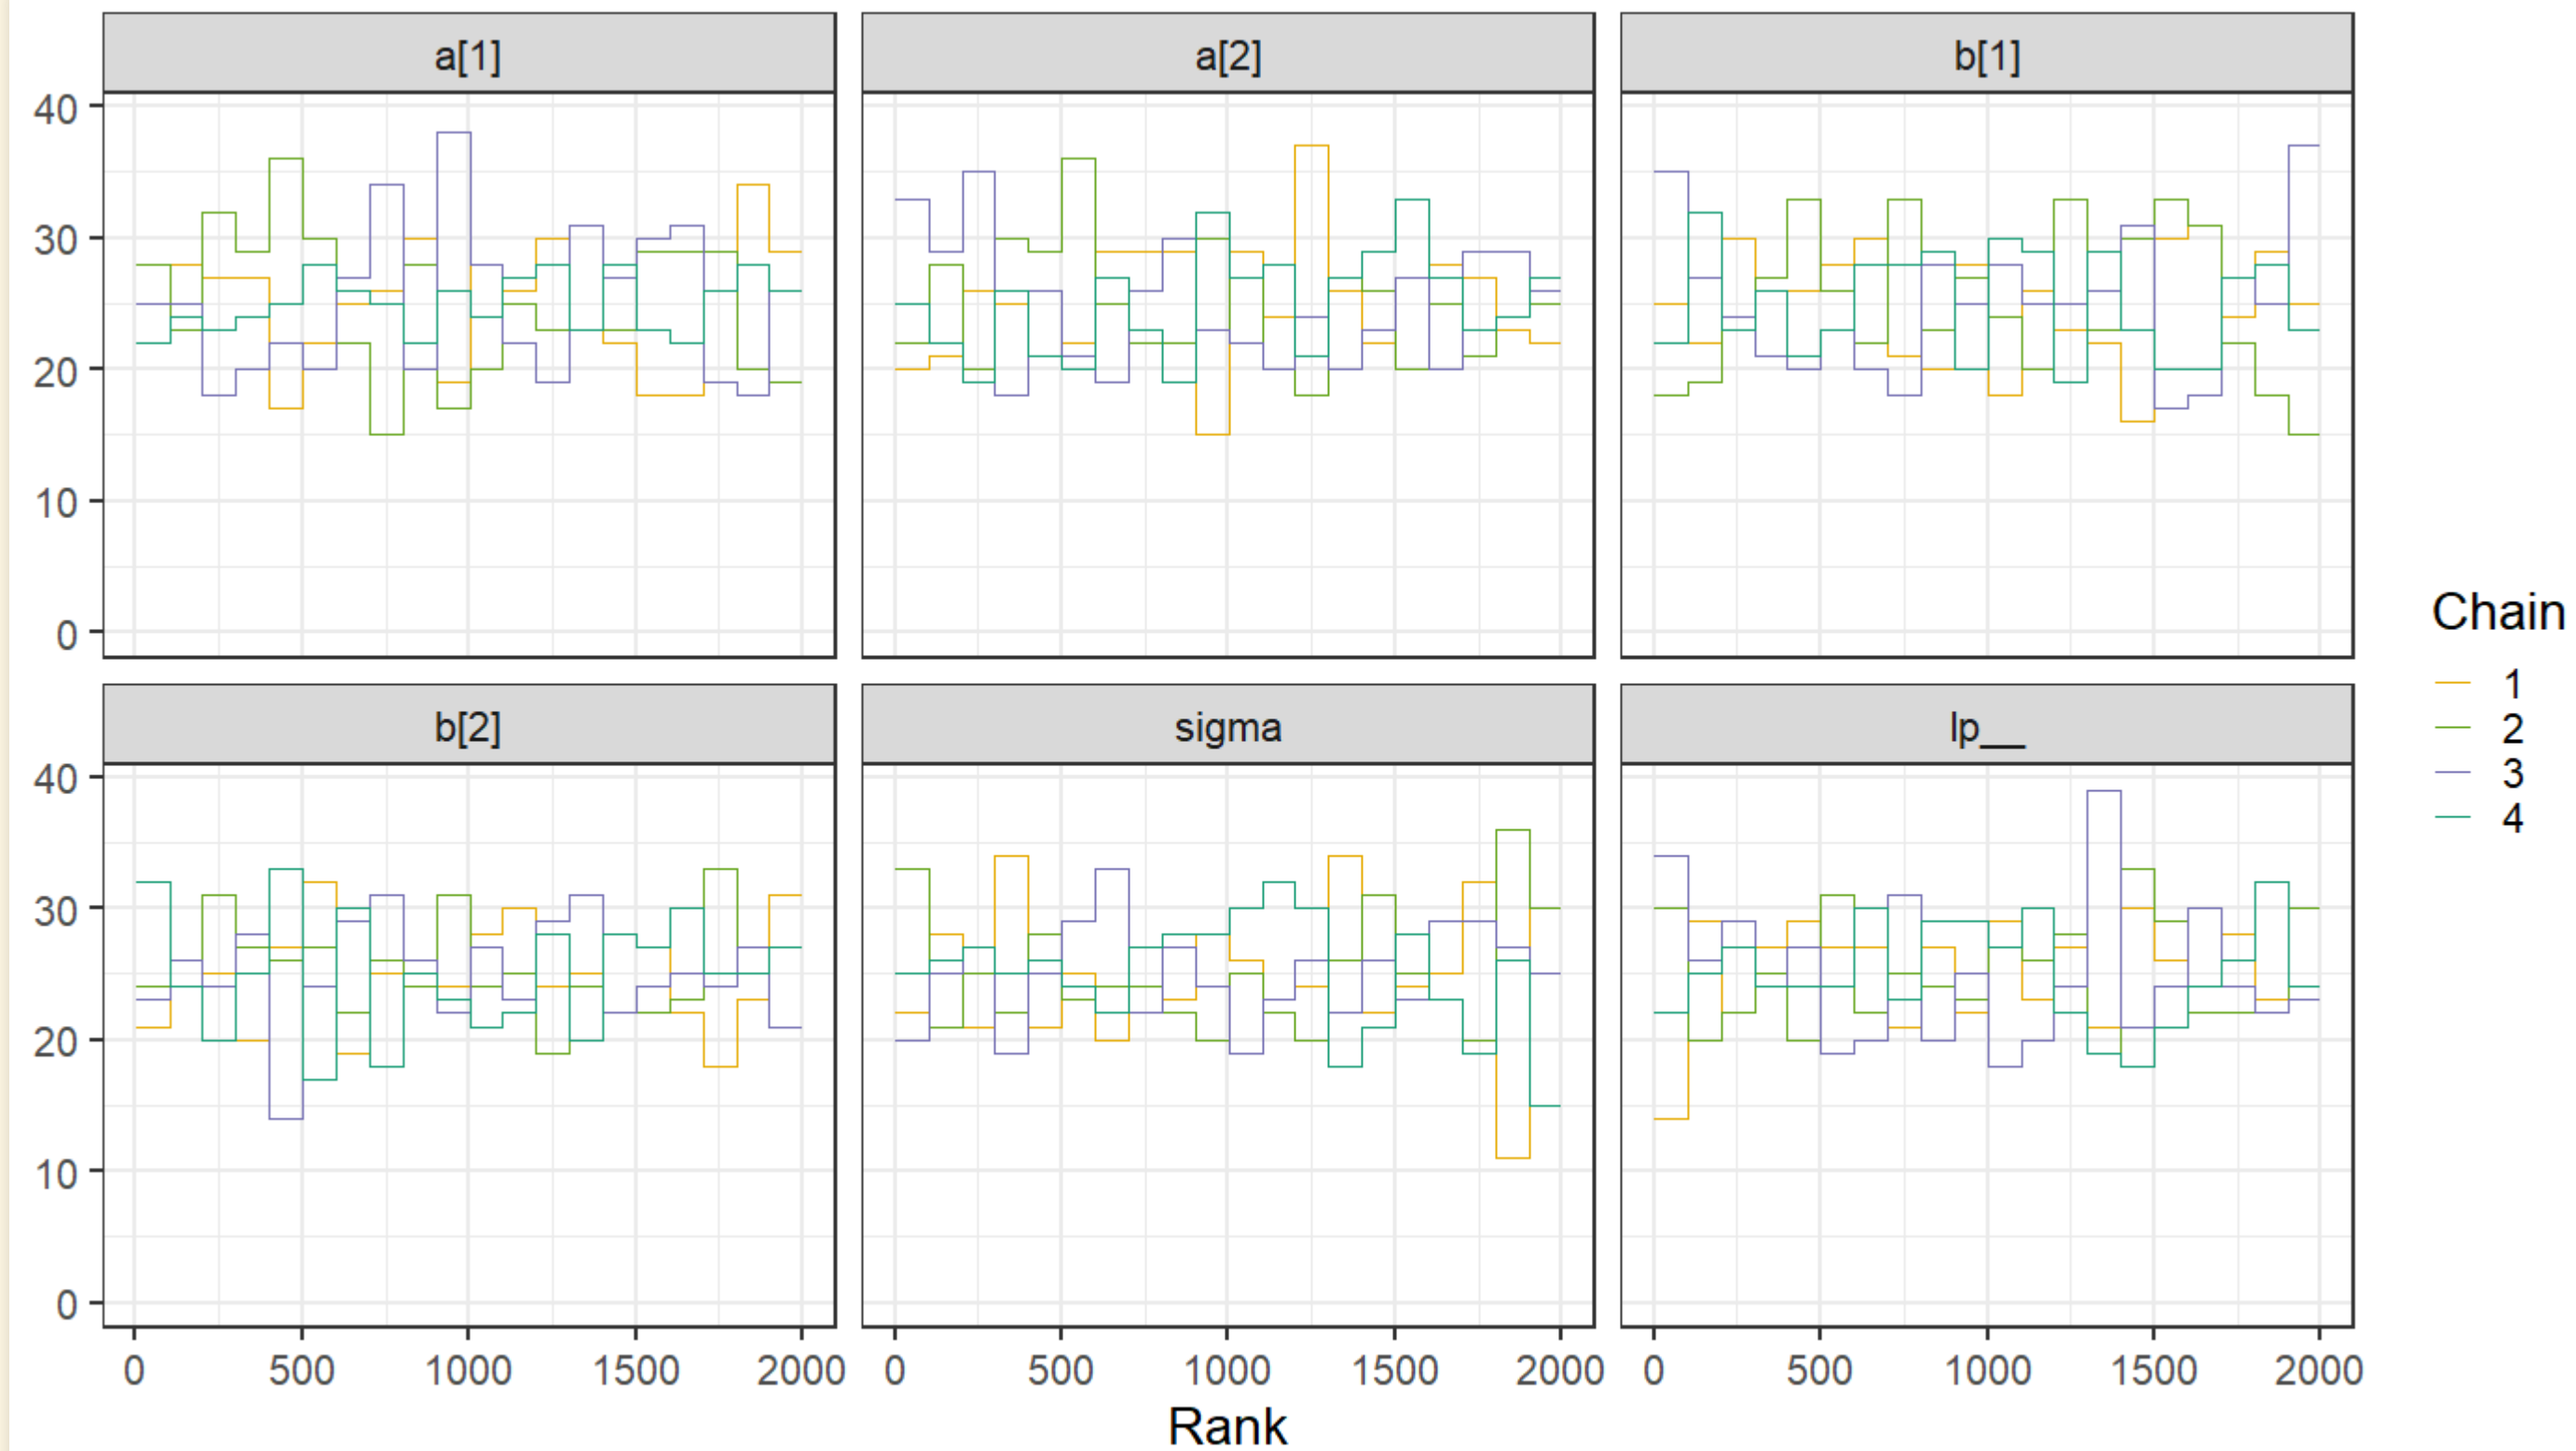
MCMC Trace Plots

```
color_scheme_set("brewer-Dark2")  
mcmc_trace(mdl_fit)
```



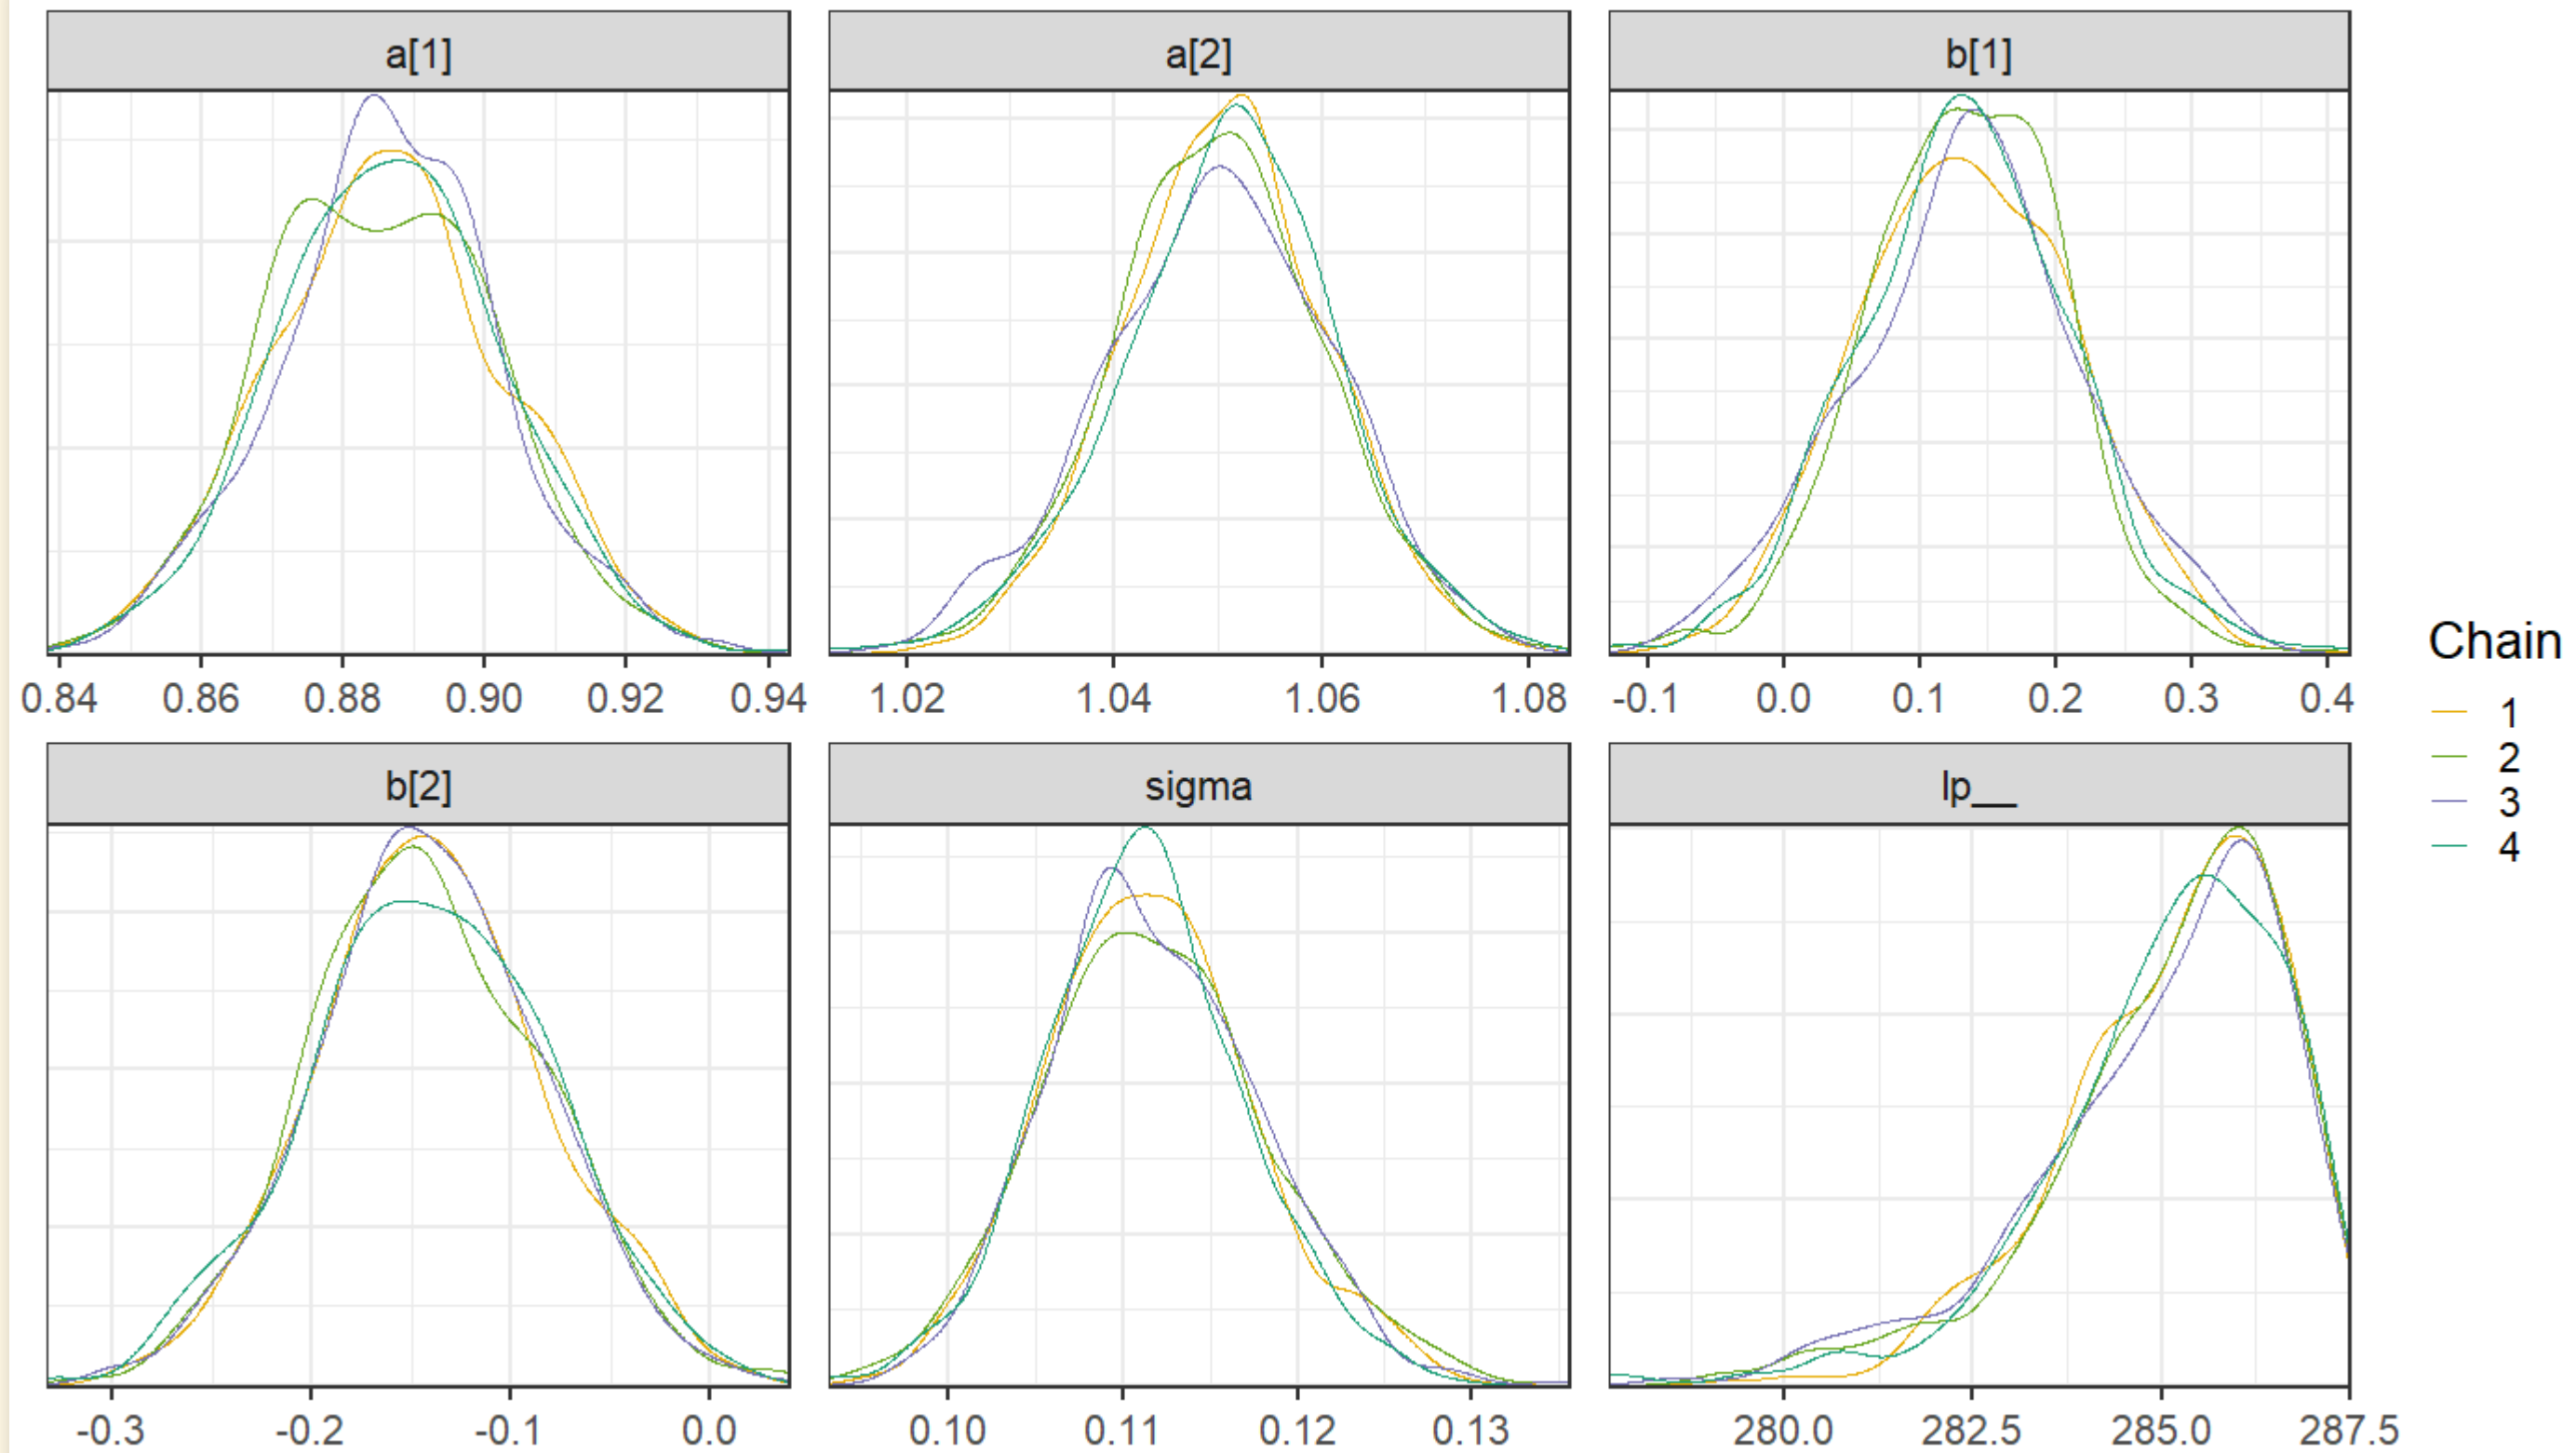
MCMC Trace-Rank Plots

```
mcmc_rank_overlay(mdl_fit)
```



Posterior Density Plots

```
mcmc_dens_overlay(mdl_fit)
```



MCMC Problems

Poorly Fit Model

```
y <- c(-1, 1)
set.seed(11)
mdl_bad <- ulam(
  alist(
    y ~ dnorm(mu, sigma),
    mu <- alpha,
    alpha ~ dnorm(0, 1000),
    sigma ~ dexp(0.0001)
  ), data = list(y = y), chains = 4, cores = 4 )
```

- 2 data observations
- Very uninformative priors
- Gives warning:

```
:: {bare .mtop-1 .max-listing .seventy}
```

```
Warning: 87 of 2000 (4.0%)
transitions ended with a
divergence.
See https://mc-
stan.org/misc/warnings for
details.
```

```
::
```

```
show(mdl_bad)
```

```
## Hamiltonian Monte Carlo approximation
## 2000 samples from 4 chains
##
## Sampling durations (seconds):
##           warmup sample total
## chain:1    0.07    0.04    0.11
## chain:2    0.07    0.06    0.13
## chain:3    0.15    0.06    0.20
## chain:4    0.04    0.04    0.08
##
## Formula:
## y ~ dnorm(mu, sigma)
## mu <- alpha
## alpha ~ dnorm(0, 1000)
## sigma ~ dexp(1e-04)
```


```
precis_show(precis(mdl_bad, digits=2))
```

```
##           mean      sd    5.5%   94.5% n_eff Rhat4
## alpha   16.35   334.62 -354.99  481.62   194   1.02
## sigma  492.14 1303.25    4.84 2083.80   156   1.04
```

- **n_eff**: <200 effective samples out of 2000.
- **sd**: Huge standard deviations in posteriors.

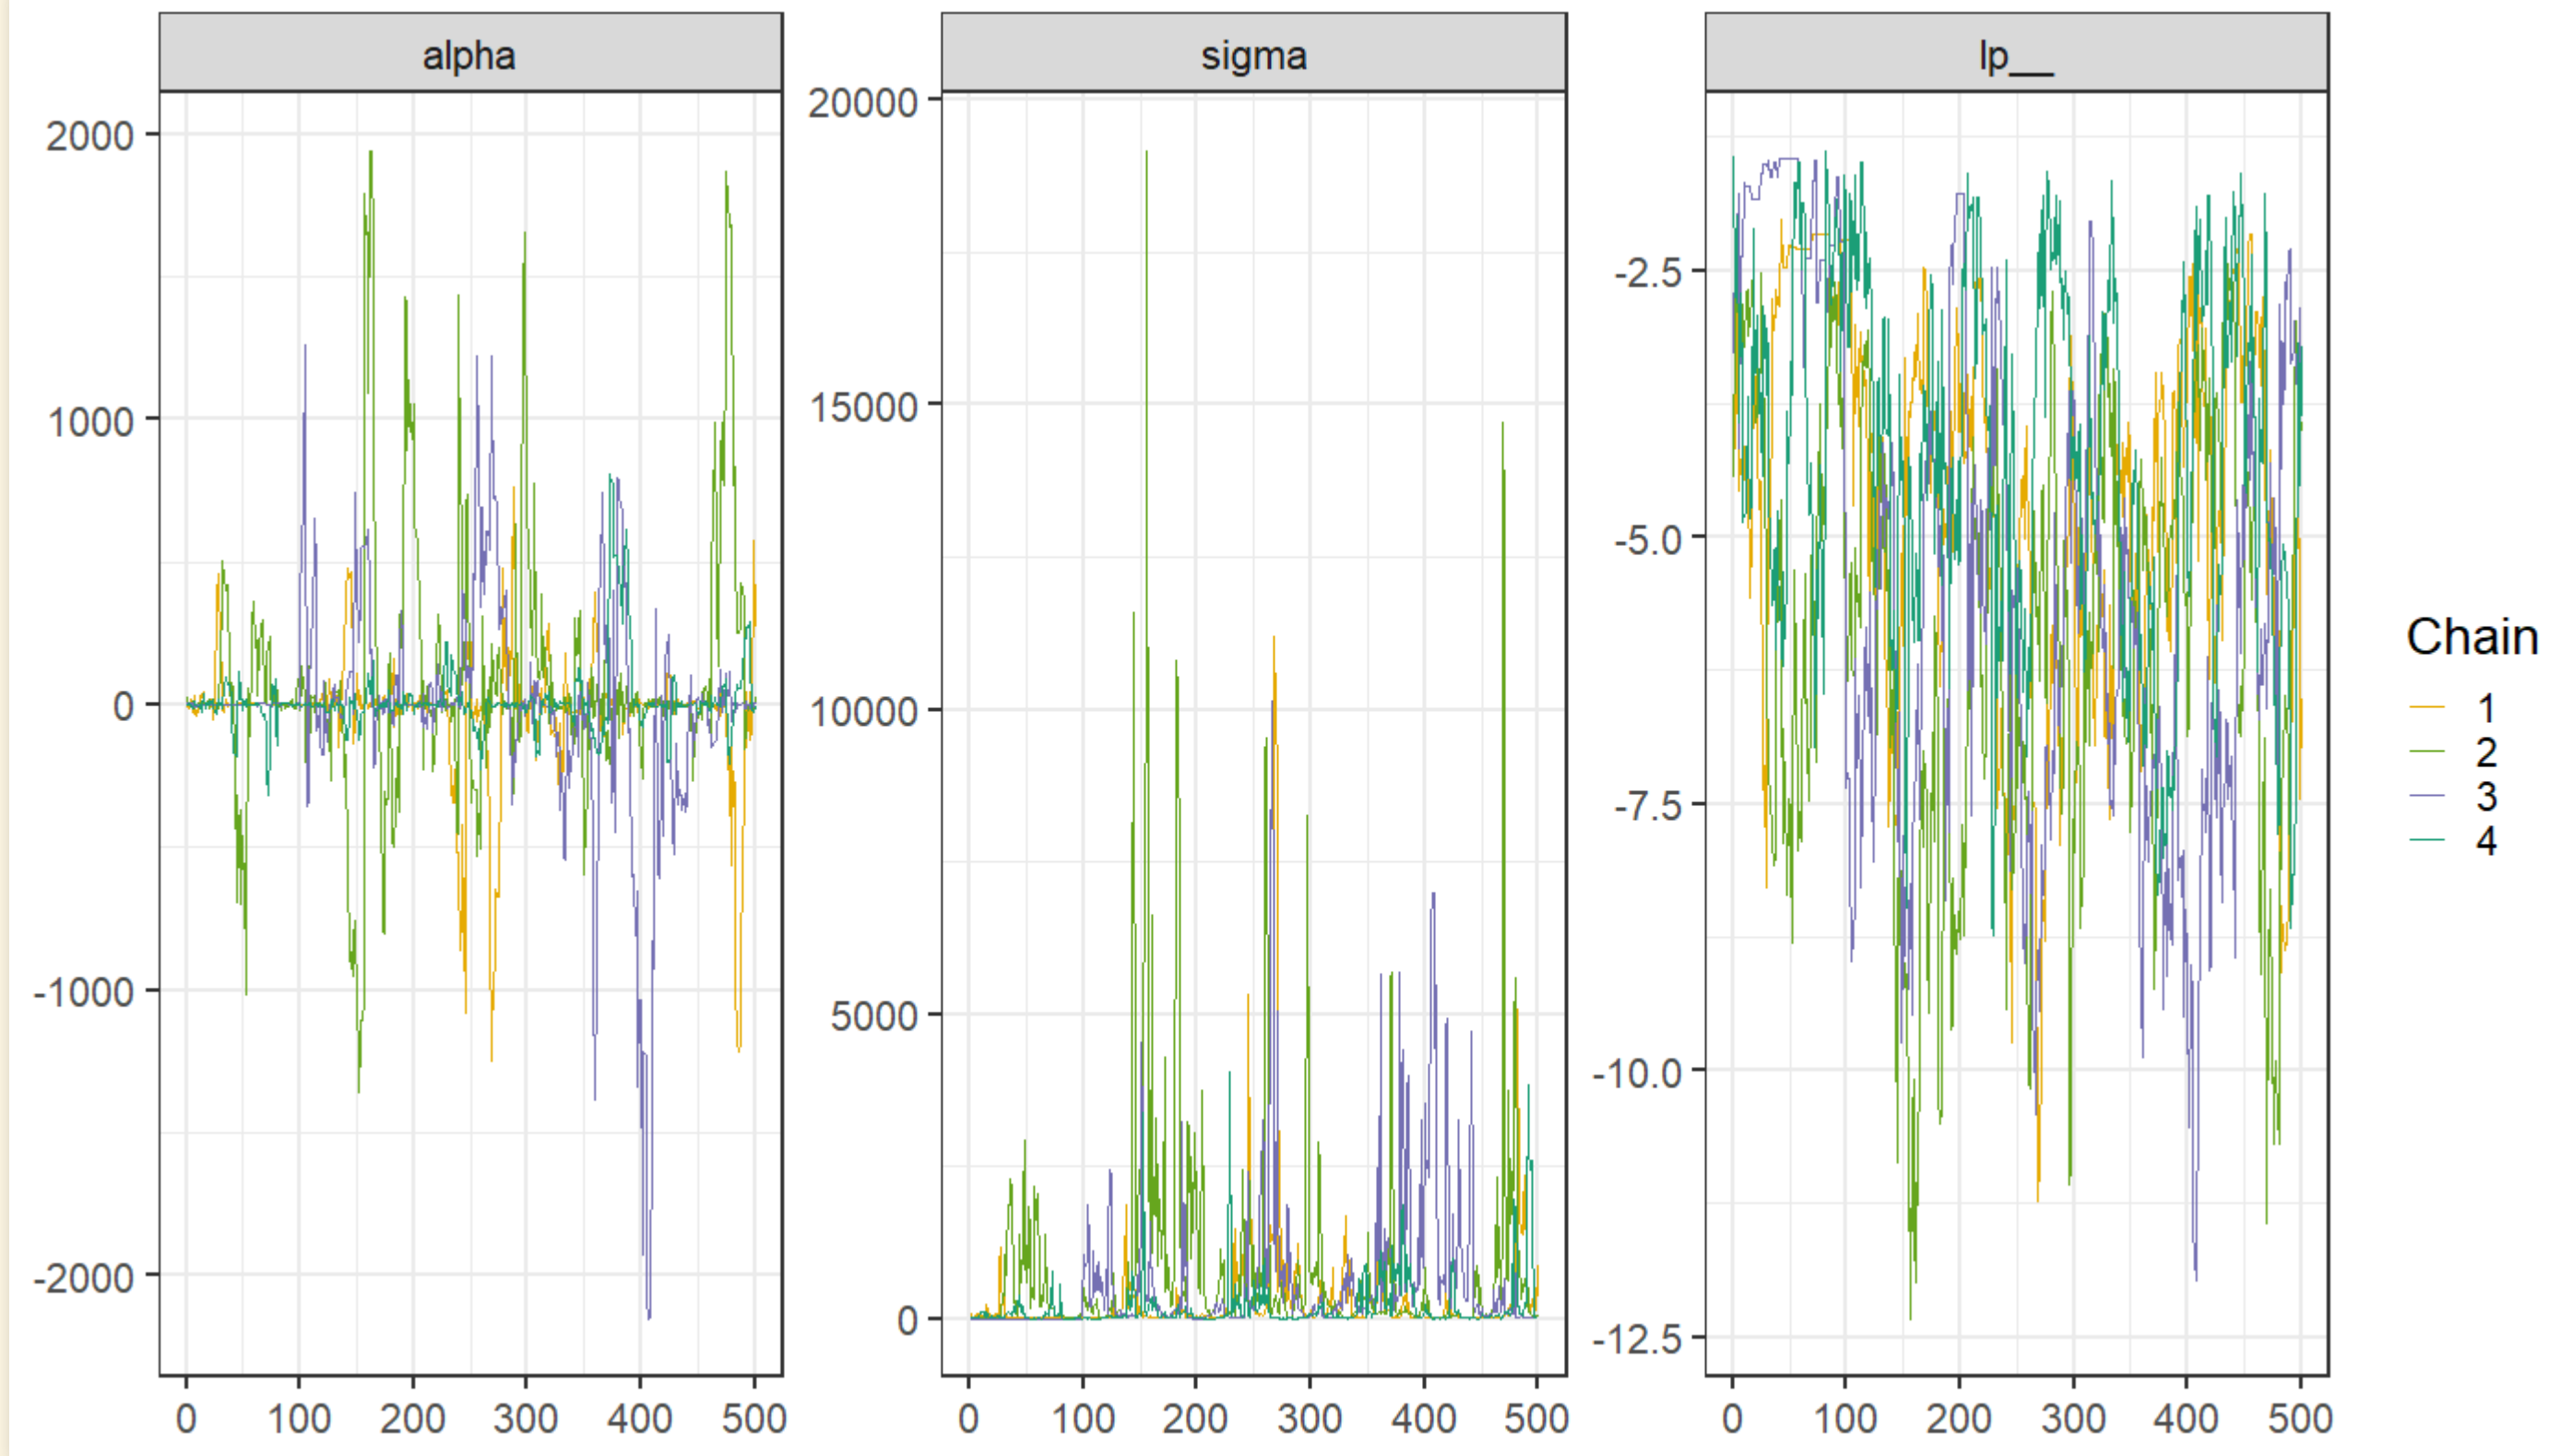
Show Posteriors

```
bad_md1_fit <- mdl_bad@stanfit  
color_scheme_set("brewer-Dark2")  
mcmc_dens_overlay(bad_md1_fit)
```




MCMC Trace Plots

```
mcmc_trace(bad_md1_fit)
```



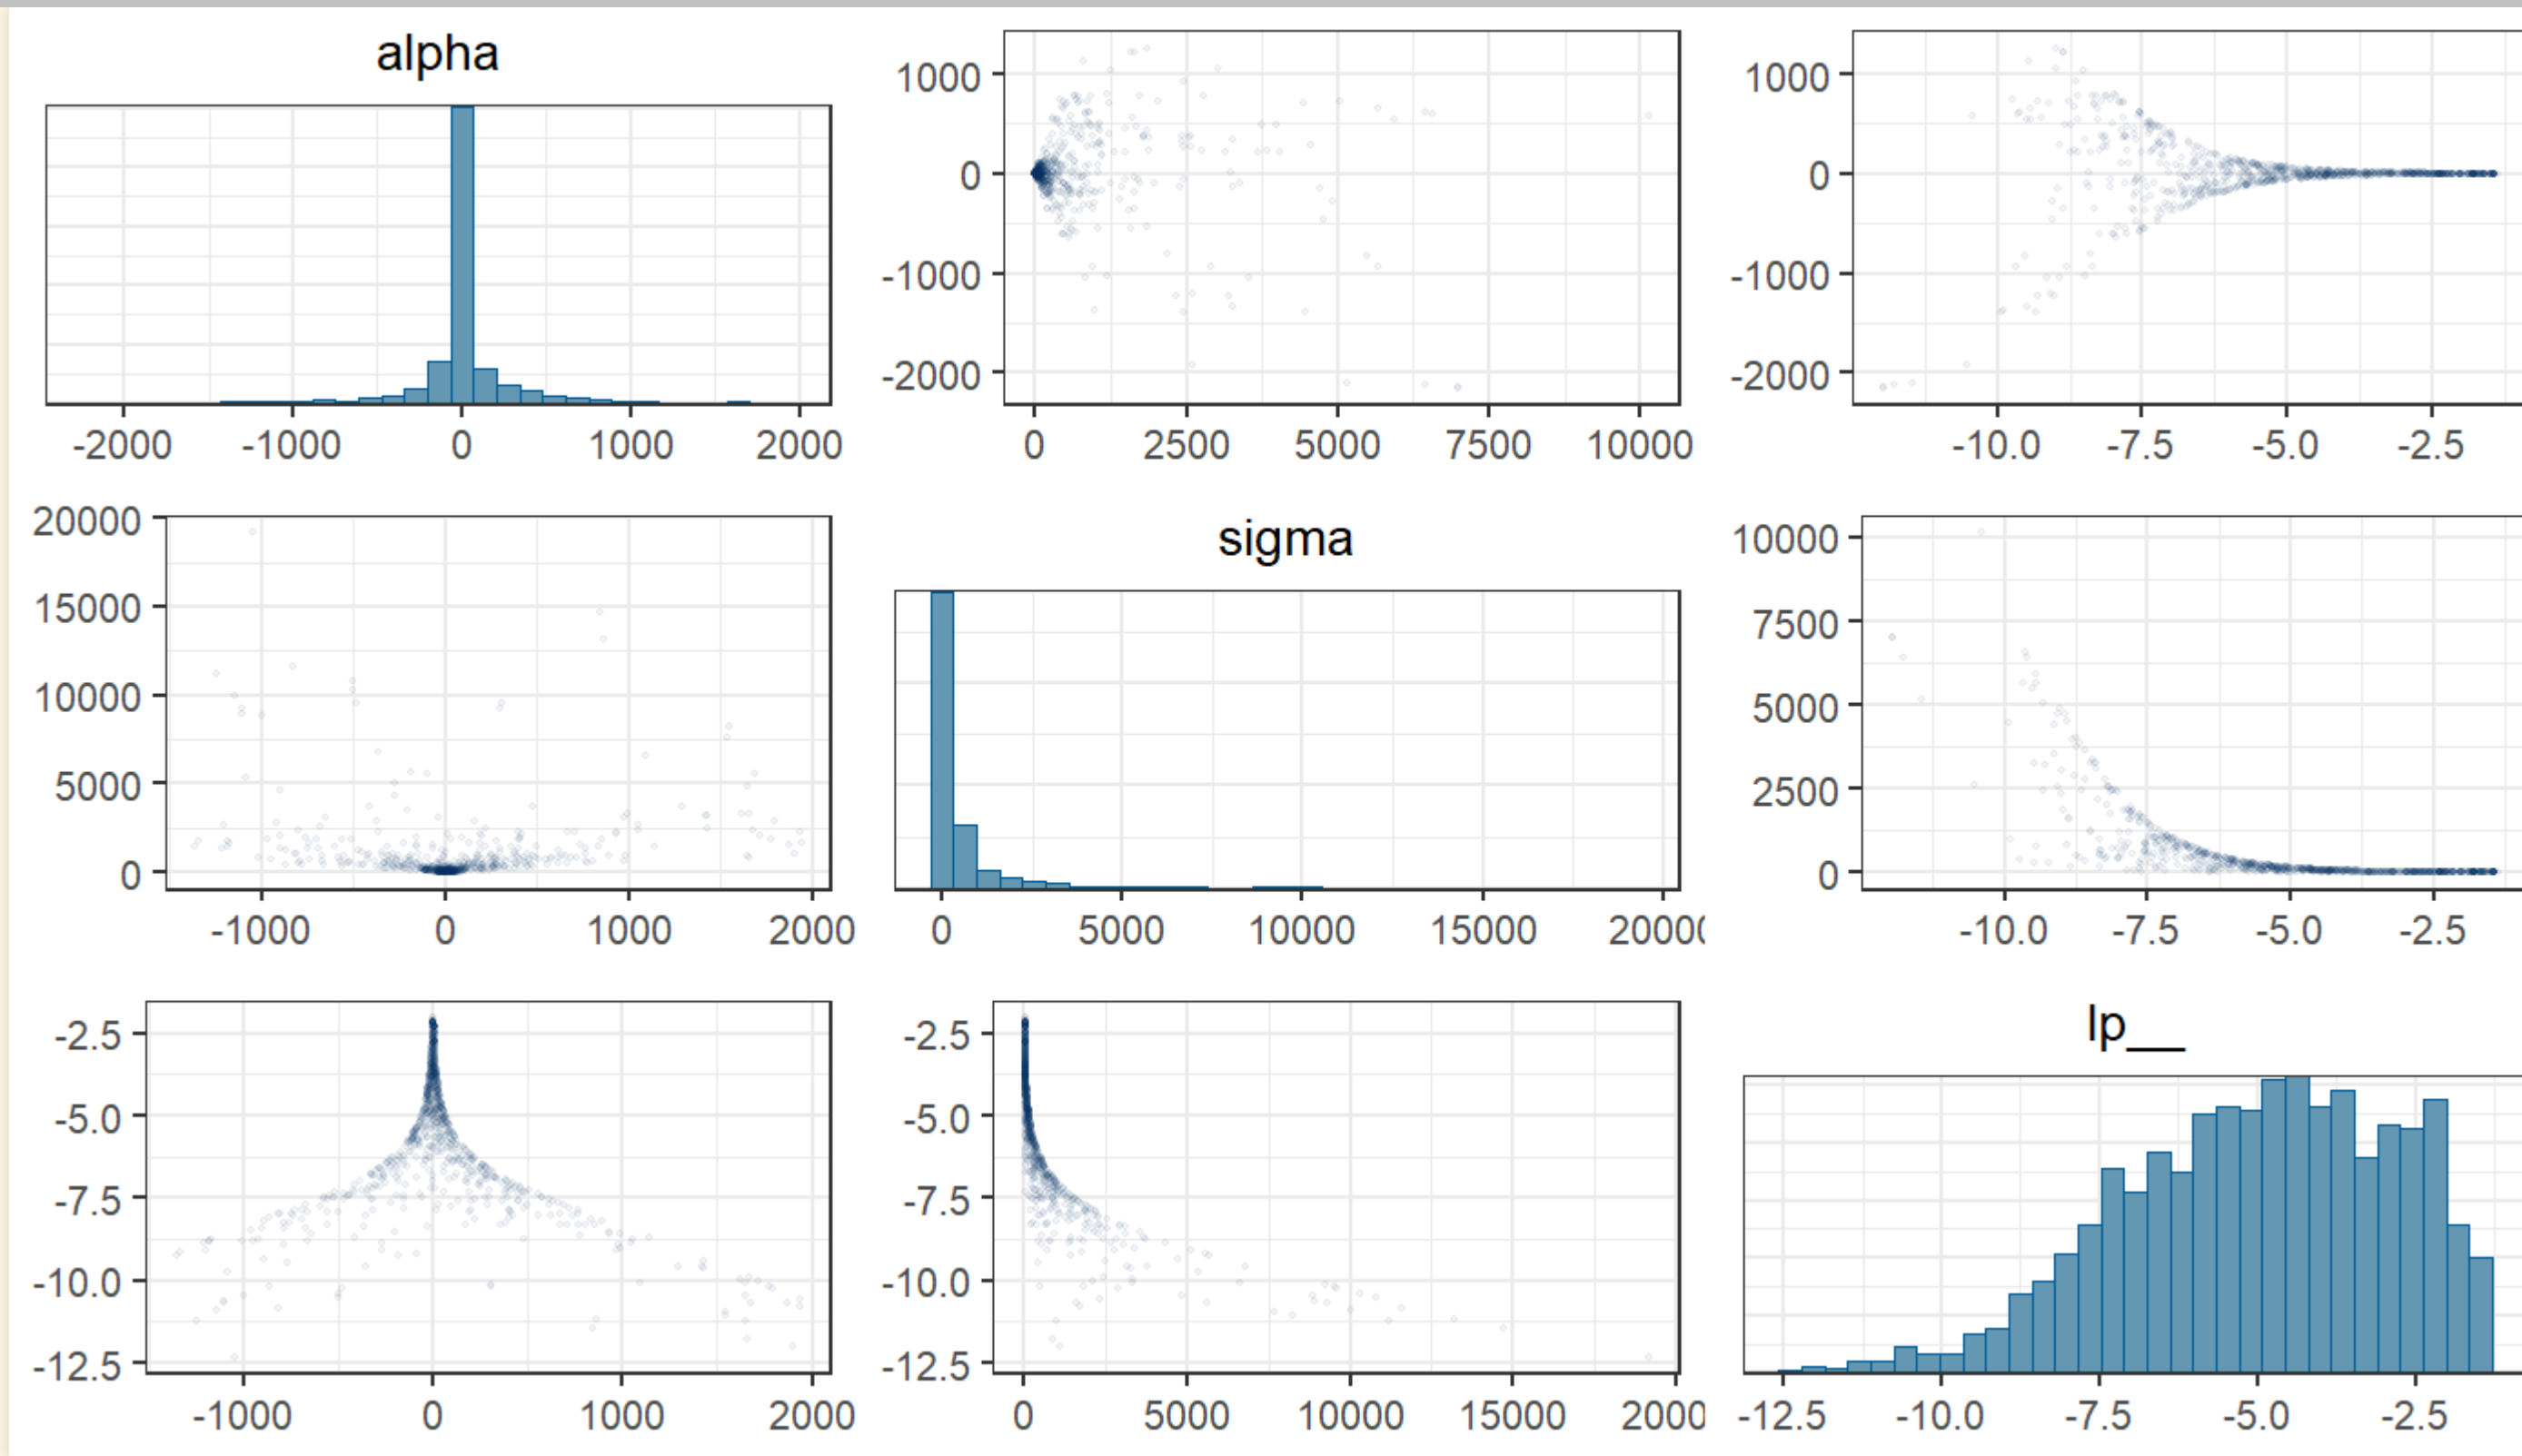
MCMC Trace-Rank Plots

```
mcmc_rank_overlay(bad_md1_fit)
```



Pair-Correlation Plots: Funnels

```
color_scheme_set("blue")
mcmc_pairs(bad_md1_fit, off_diag_args = list(size = 1, alpha = 0.05))
```



Better Model

```
set.seed(11)
mdl_better <- ulam(
  alist(
    y ~ dnorm(mu, sigma),
    mu <- alpha,
    alpha ~ dnorm(0, 10),
    sigma ~ dexp(1)
  ), data = list(y = y), chains = 4, cores = 4 )
```

- Replace uninformative priors with weakly informative priors. Weakly informative priors help a lot.

```
show(mdl_better)
```

```
## Hamiltonian Monte Carlo approximation
## 2000 samples from 4 chains
##
## Sampling durations (seconds):
##           warmup sample total
## chain:1      0.03      0.02   0.05
## chain:2      0.02      0.02   0.04
## chain:3      0.03      0.02   0.04
## chain:4      0.03      0.03   0.05
##
## Formula:
## y ~ dnorm(mu, sigma)
## mu <- alpha
## alpha ~ dnorm(0, 10)
## sigma ~ dexp(1)
```


```
precis_show(precis(mdl_bad, digits=2))
```

```
##           mean      sd    5.5%    94.5% n_eff Rhat4
## alpha  16.35  334.62 -354.99  481.62   194  1.02
## sigma 492.14 1303.25    4.84 2083.80   156  1.04
```

- Even weakly informative priors help a lot!
- `n_eff`: >400 effective samples.
- `sd`: Much smaller.


Show Posteriors

```
better_md1_fit <- mdl_better@stanfit  
mcmc_dens_overlay(better_md1_fit)
```




MCMC Trace Plots

```
color_scheme_set("brewer-Dark2")  
mcmc_trace(better_mdl_fit)
```



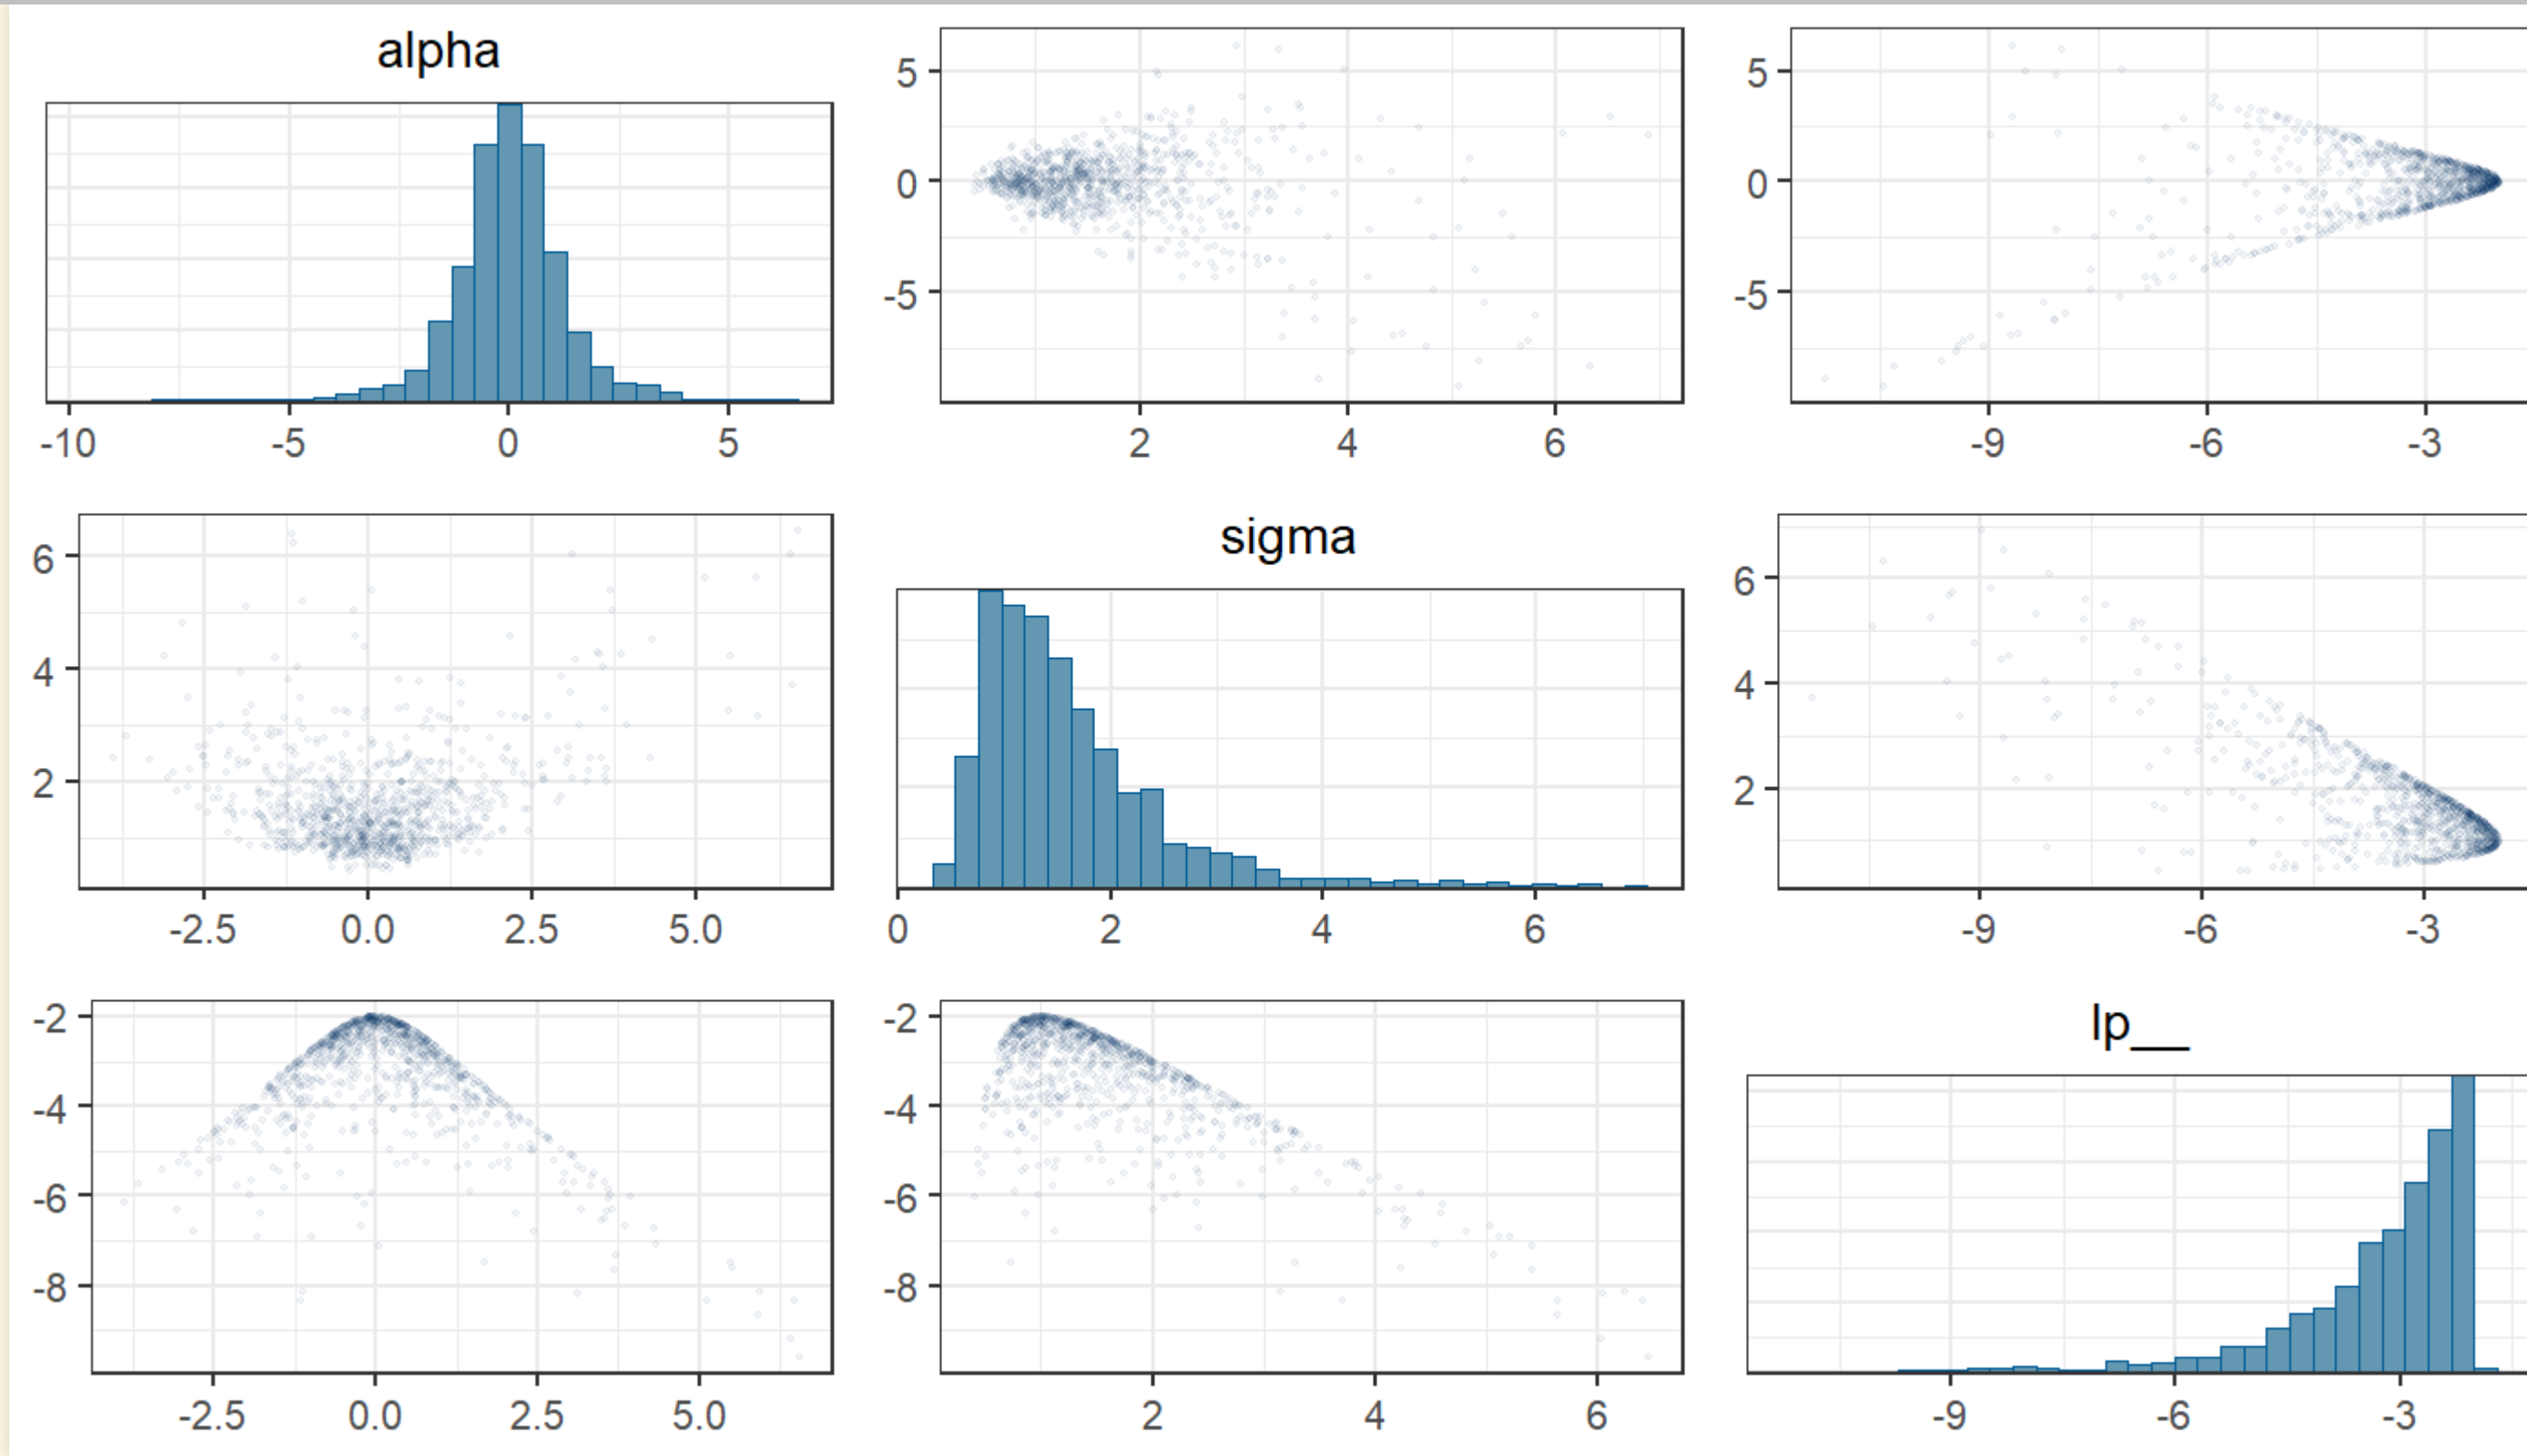
MCMC Trace-Rank Plots

```
mcmc_rank_overlay(better_mdl_fit)
```



Pair-Correlation Plots

```
color_scheme_set("blue")  
mcmc_pairs(better_mdl_fit, off_diag_args = list(size = 1, alpha = 0.05))
```



Non-Identifiable Models

Non-Identifiable Model

```
set.seed(41)
y <- rnorm(100, mean = 0, sd = 1)
```

- Fit model
$$\begin{aligned} y &\sim \text{Normal}(\mu, \sigma) \\ \mu &= \alpha_1 + \alpha_2 \\ \alpha_1 &\sim \text{Normal}(0, 1000) \\ \alpha_2 &\sim \text{Normal}(0, 1000) \\ \sigma &\sim \text{Exponential}(1) \end{aligned}$$

```
set.seed(11)
mdl_non_id <- ulam(
  alist(
    y ~ dnorm(mu, sigma),
    mu <- a1 + a2,
    a1 ~ dnorm(0, 1000),
    a2 ~ dnorm(0, 1000),
    sigma ~ dexp(1)
  ), data = list(y = y), chains = 4, cores = 4 )
```

```
show(mdl_non_id)
```

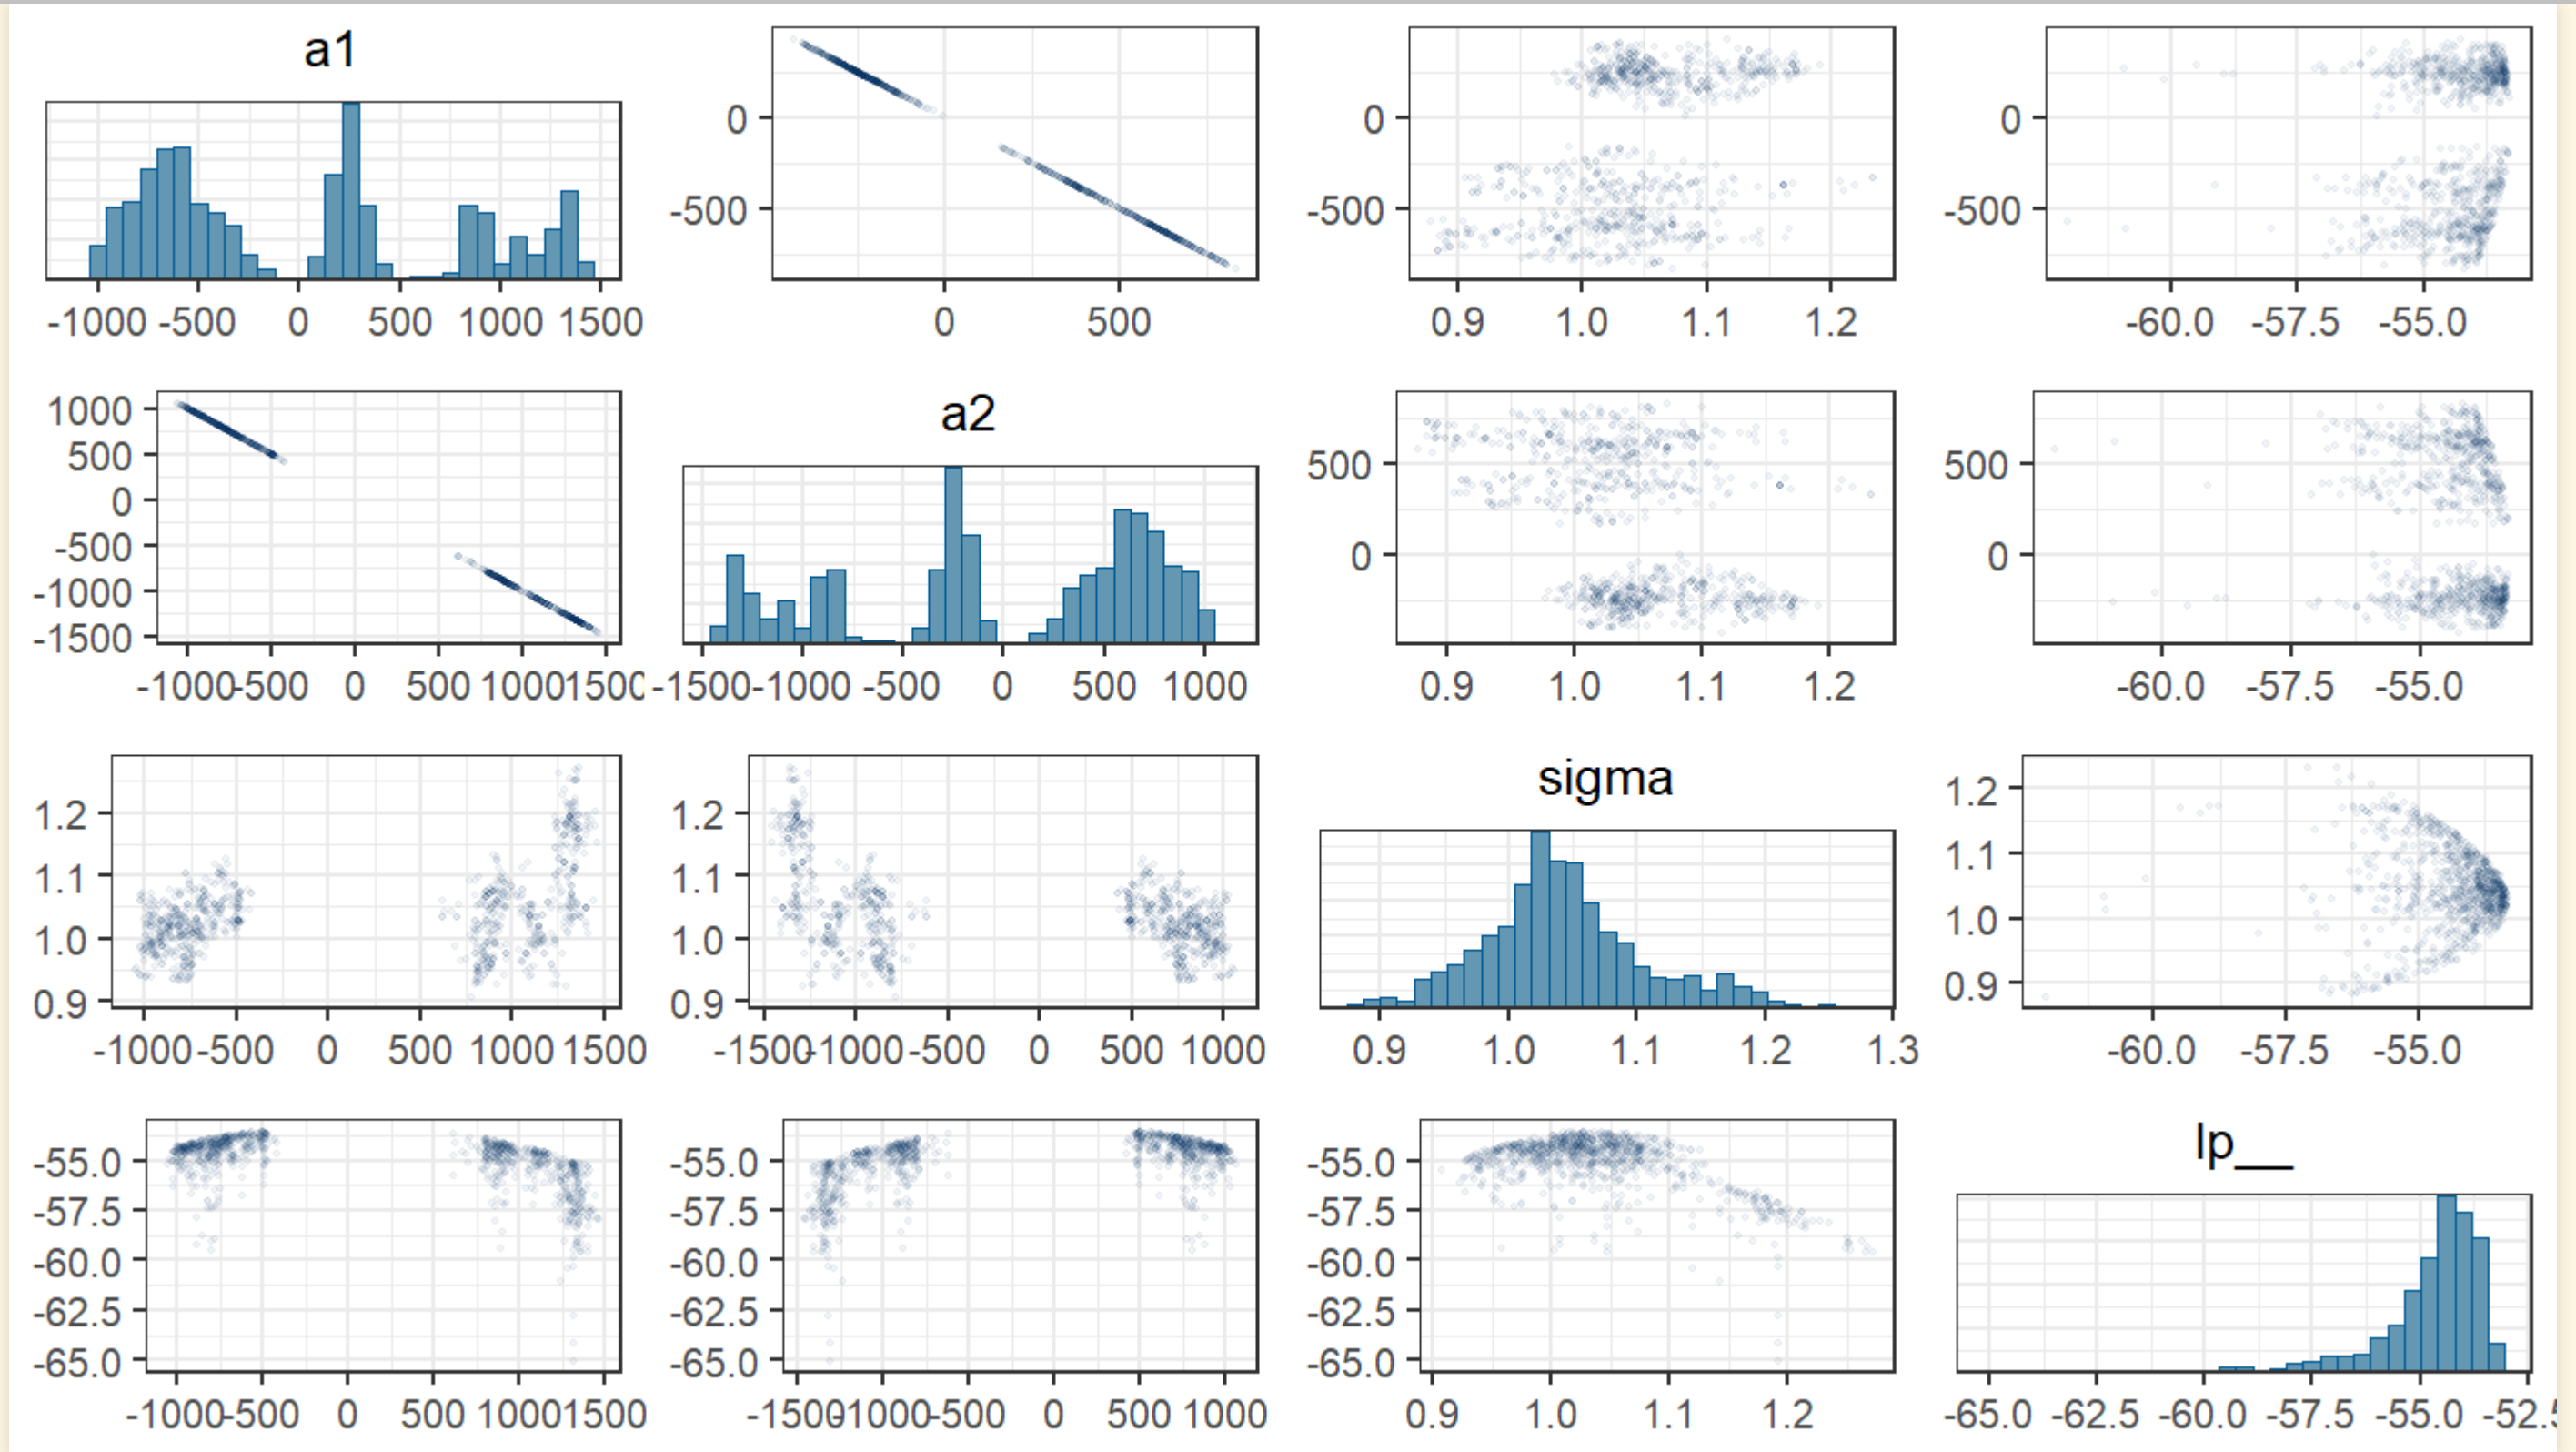
```
## Hamiltonian Monte Carlo approximation
## 2000 samples from 4 chains
##
## Sampling durations (seconds):
##           warmup sample total
## chain:1      1.20      1.30    2.50
## chain:2      1.08      1.35    2.43
## chain:3      1.12      1.33    2.45
## chain:4      1.19      1.25    2.45
##
## Formula:
## y ~ dnorm(mu, sigma)
## mu <- a1 + a2
## a1 ~ dnorm(0, 1000)
## a2 ~ dnorm(0, 1000)
## sigma ~ dexp(1)
```

```
precis_show(precis(mdl_non_id, digits=2))
```

	mean	sd	5.5%	94.5%	n_eff	Rhat4
a1	11.76	744.23	-907.39	1317.93	2	5.85
a2	-11.58	744.24	-1317.69	907.51	2	5.85
sigma	1.04	0.06	0.95	1.17	32	1.16


Pair-Correlation Plots

```
color_scheme_set("blue")
nid_md1_fit <- mdl_non_id@stanfit
mcmc_pairs(nid_md1_fit, off_diag_args = list(size = 1, alpha = 0.05))
```



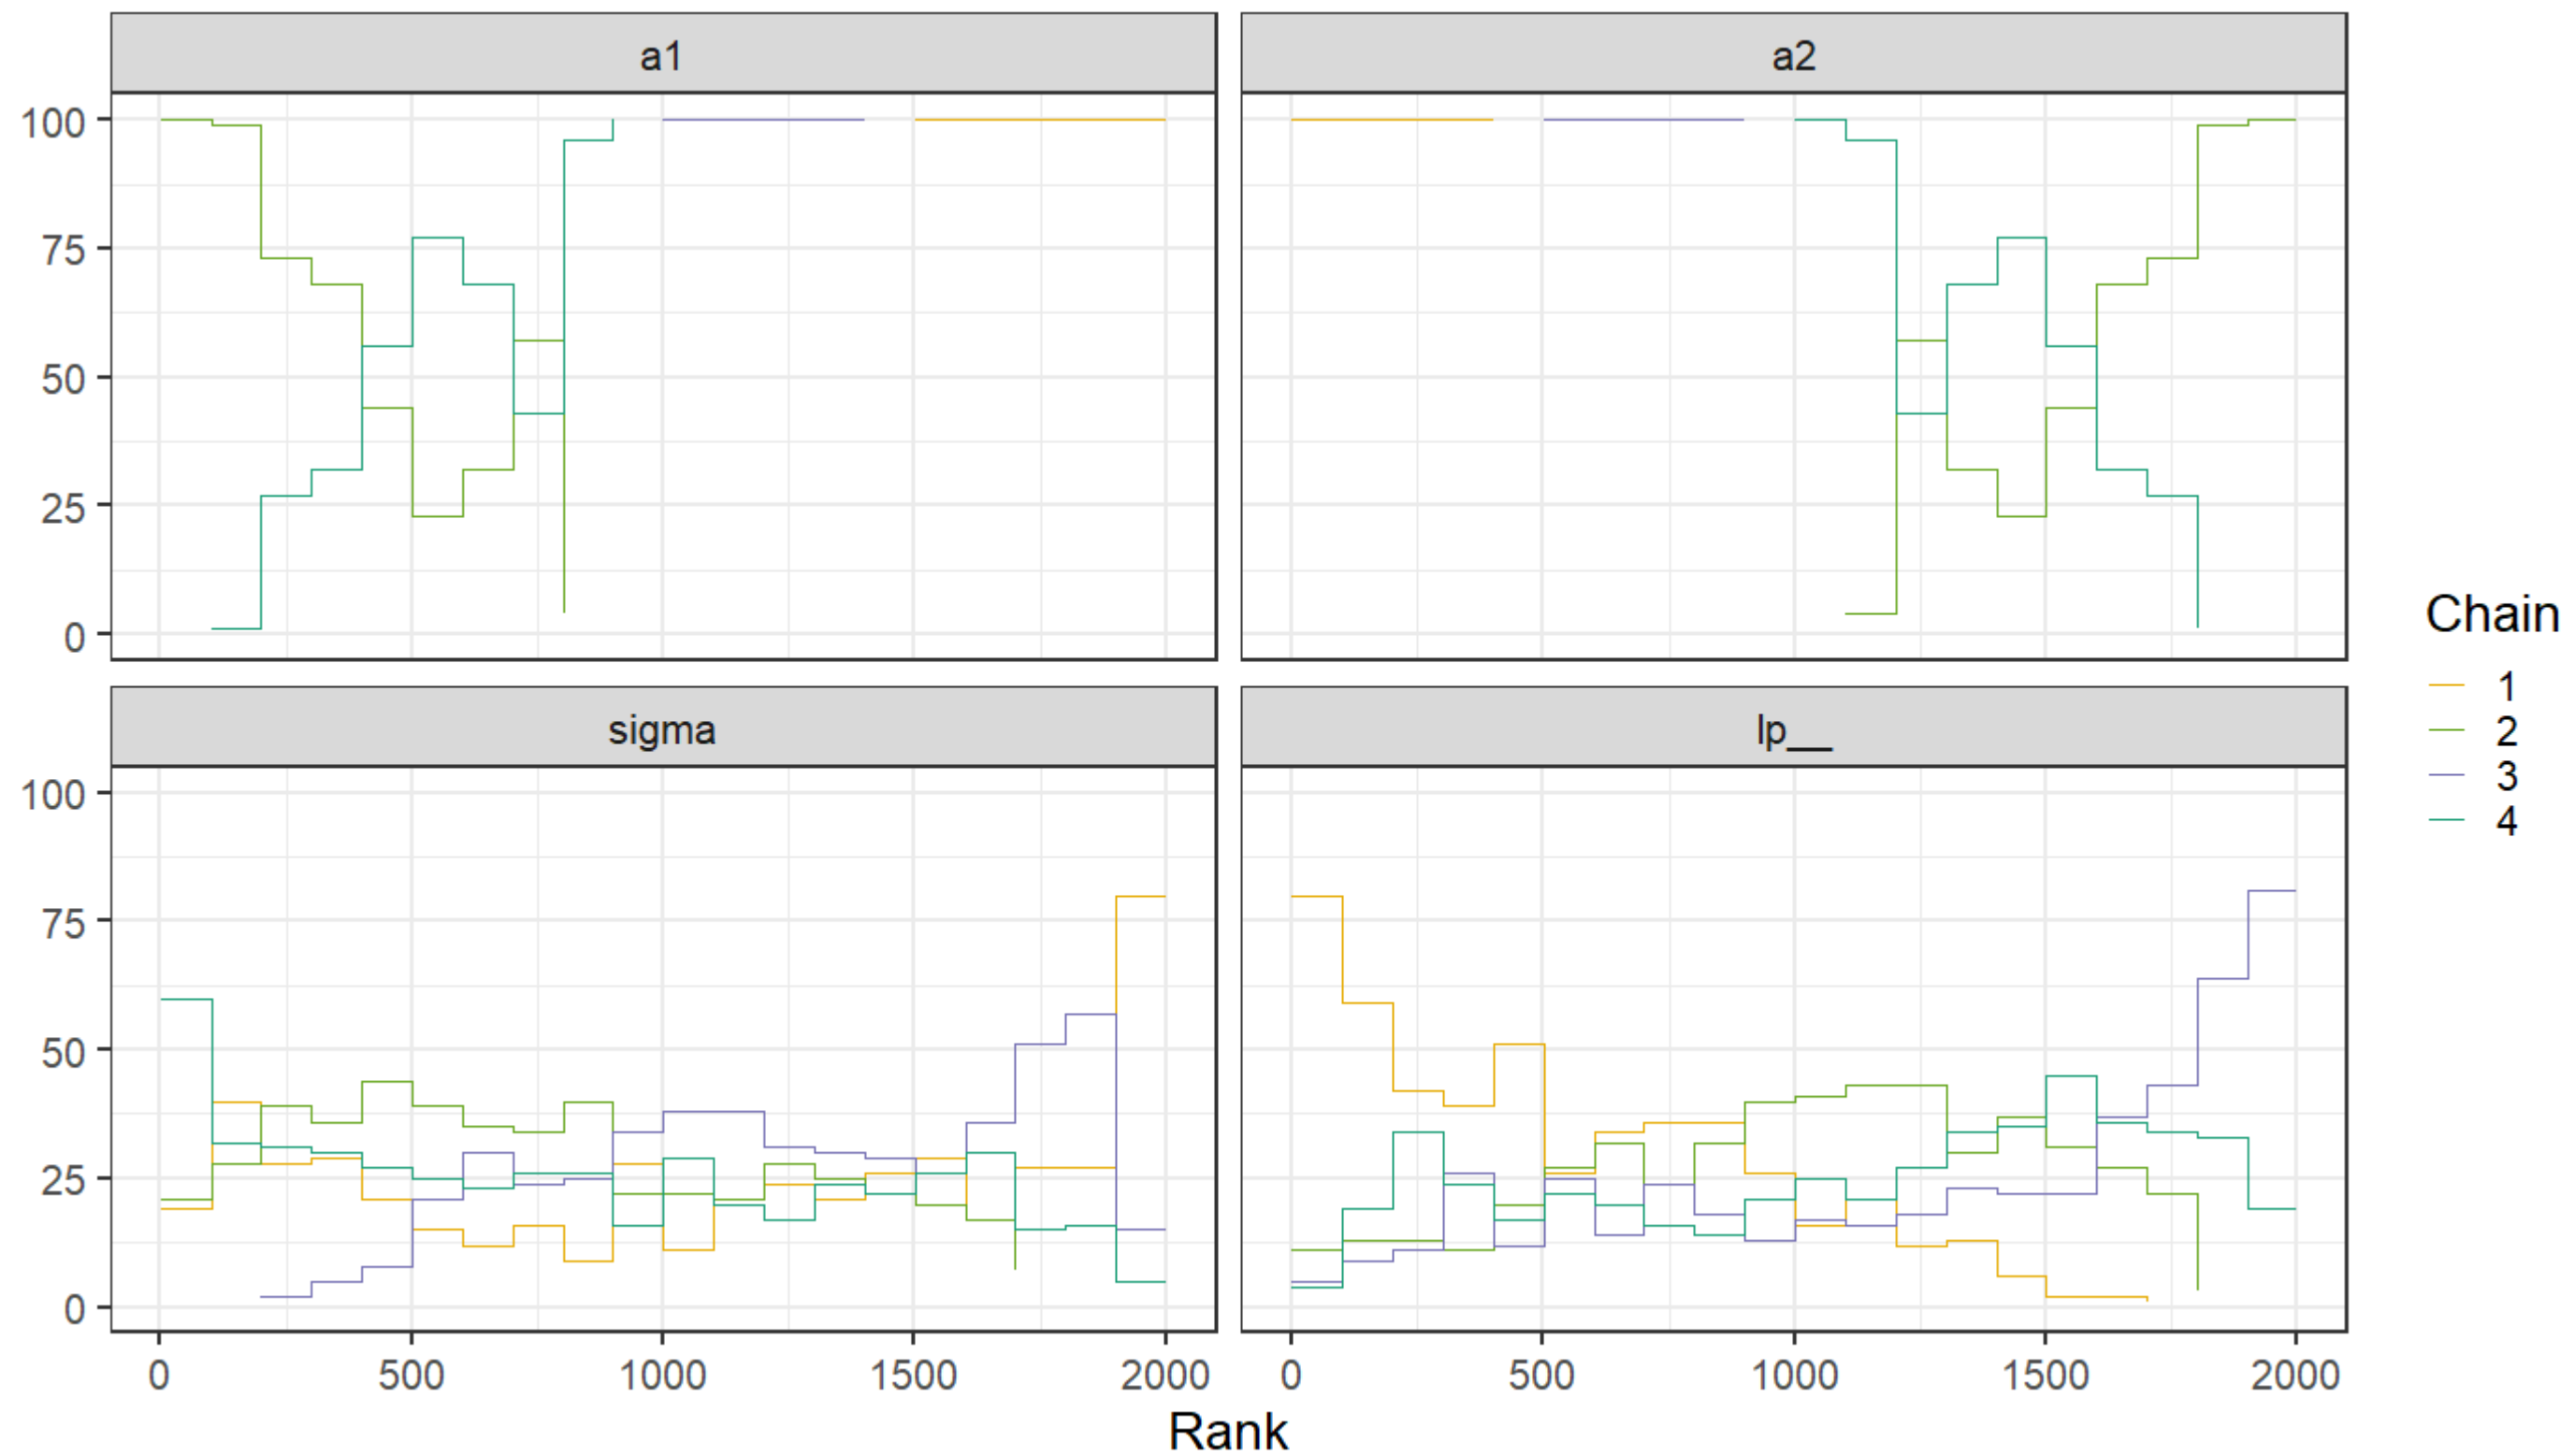
MCMC Trace Plots

```
color_scheme_set("brewer-Dark2")  
mcmc_trace(nid_md1_fit)
```



MCMC Trace-Rank Plots

```
mcmc_rank_overlay(nid_md1_fit)
```



Regularize Model

- Model with regularizing priors
$$\begin{aligned} y &\sim \text{Normal}(\mu, \sigma) \\ \mu &= \alpha_1 + \alpha_2 \\ \alpha_1 &\sim \text{Normal}(0, 10) \\ \alpha_2 &\sim \text{Normal}(0, 10) \\ \sigma &\sim \text{Exponential}(1) \end{aligned}$$

```
set.seed(11)
mdl_reg <- ulam(
  alist(
    y ~ dnorm(mu, sigma),
    mu <- a1 + a2,
    a1 ~ dnorm(0, 10),
    a2 ~ dnorm(0, 10),
    sigma ~ dexp(1)
  ), data = list(y = y), chains = 4, cores = 4 )
```

```
show(mdl_reg)
```

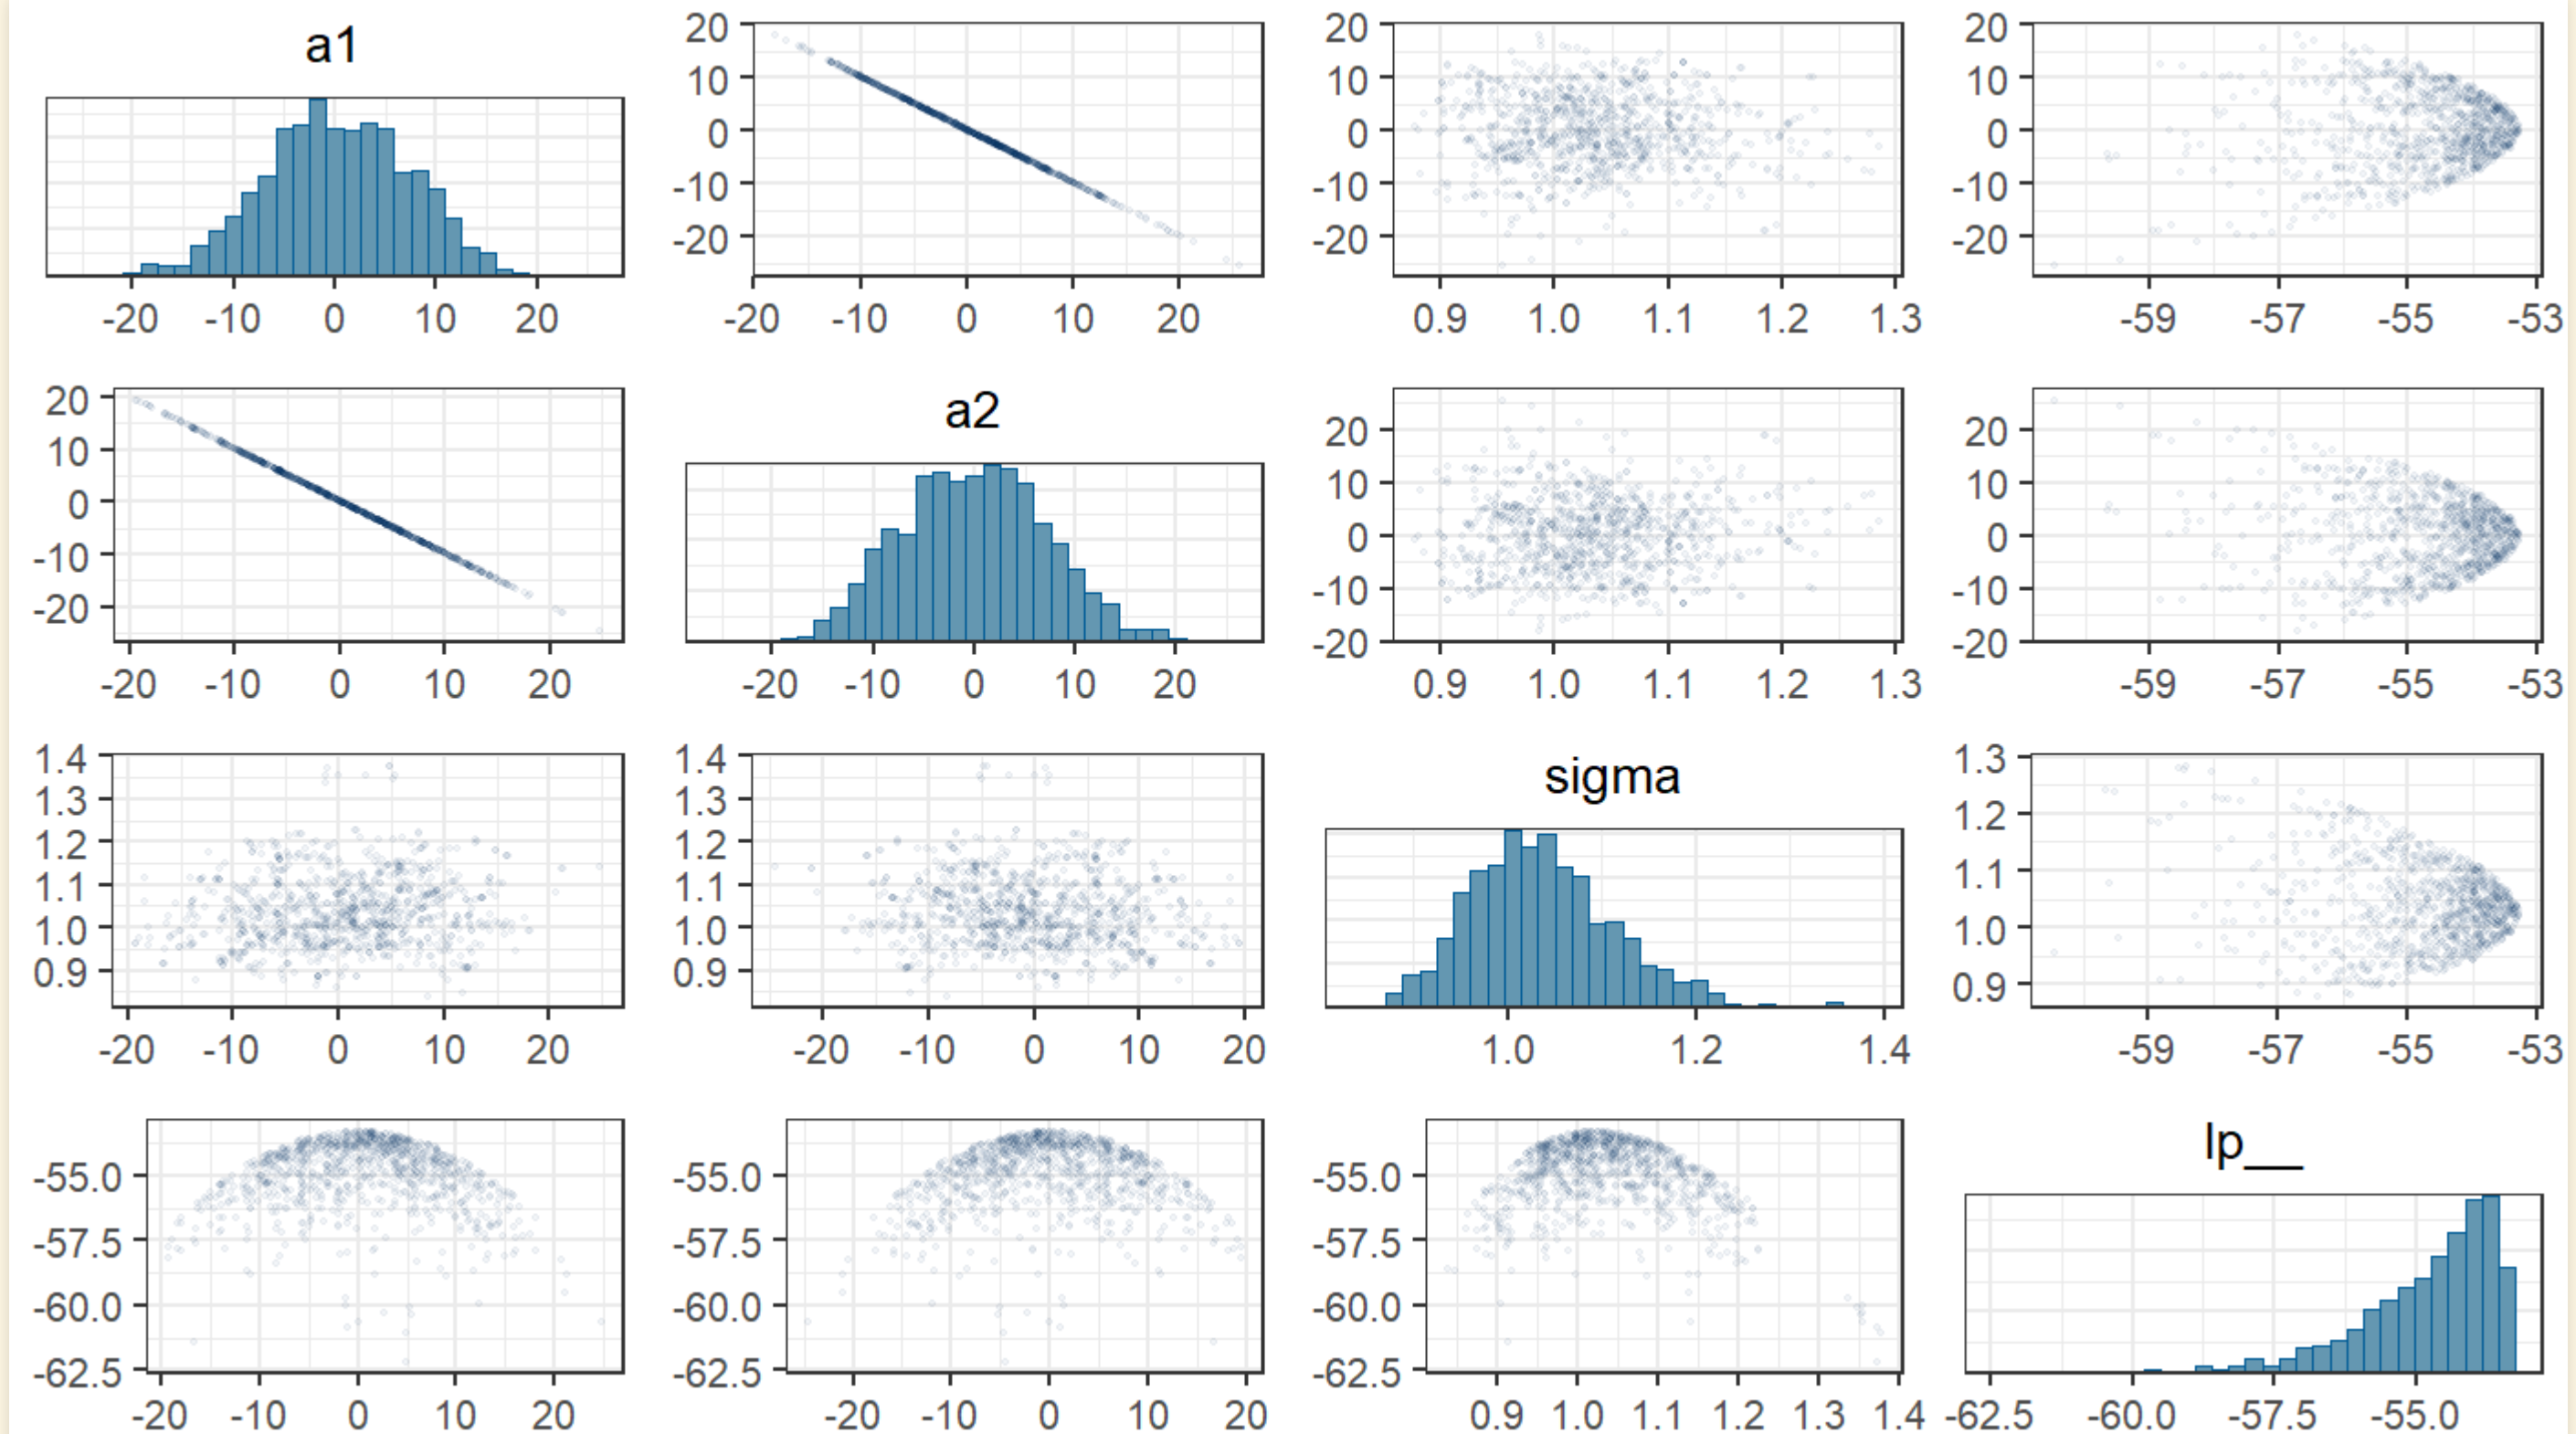
```
## Hamiltonian Monte Carlo approximation
## 2000 samples from 4 chains
##
## Sampling durations (seconds):
##           warmup sample total
## chain:1      0.46      0.45    0.91
## chain:2      0.43      0.49    0.92
## chain:3      0.39      0.42    0.81
## chain:4      0.47      0.44    0.92
##
## Formula:
## y ~ dnorm(mu, sigma)
## mu <- a1 + a2
## a1 ~ dnorm(0, 10)
## a2 ~ dnorm(0, 10)
## sigma ~ dexp(1)
```

```
precis_show(precis(mdl_reg, digits=2))
```

	mean	sd	5.5%	94.5%	n_eff	Rhat4
a1	0.25	7.19	-11.21	11.45	670	1.01
a2	-0.06	7.19	-11.20	11.40	670	1.01
sigma	1.04	0.08	0.93	1.17	616	1.00

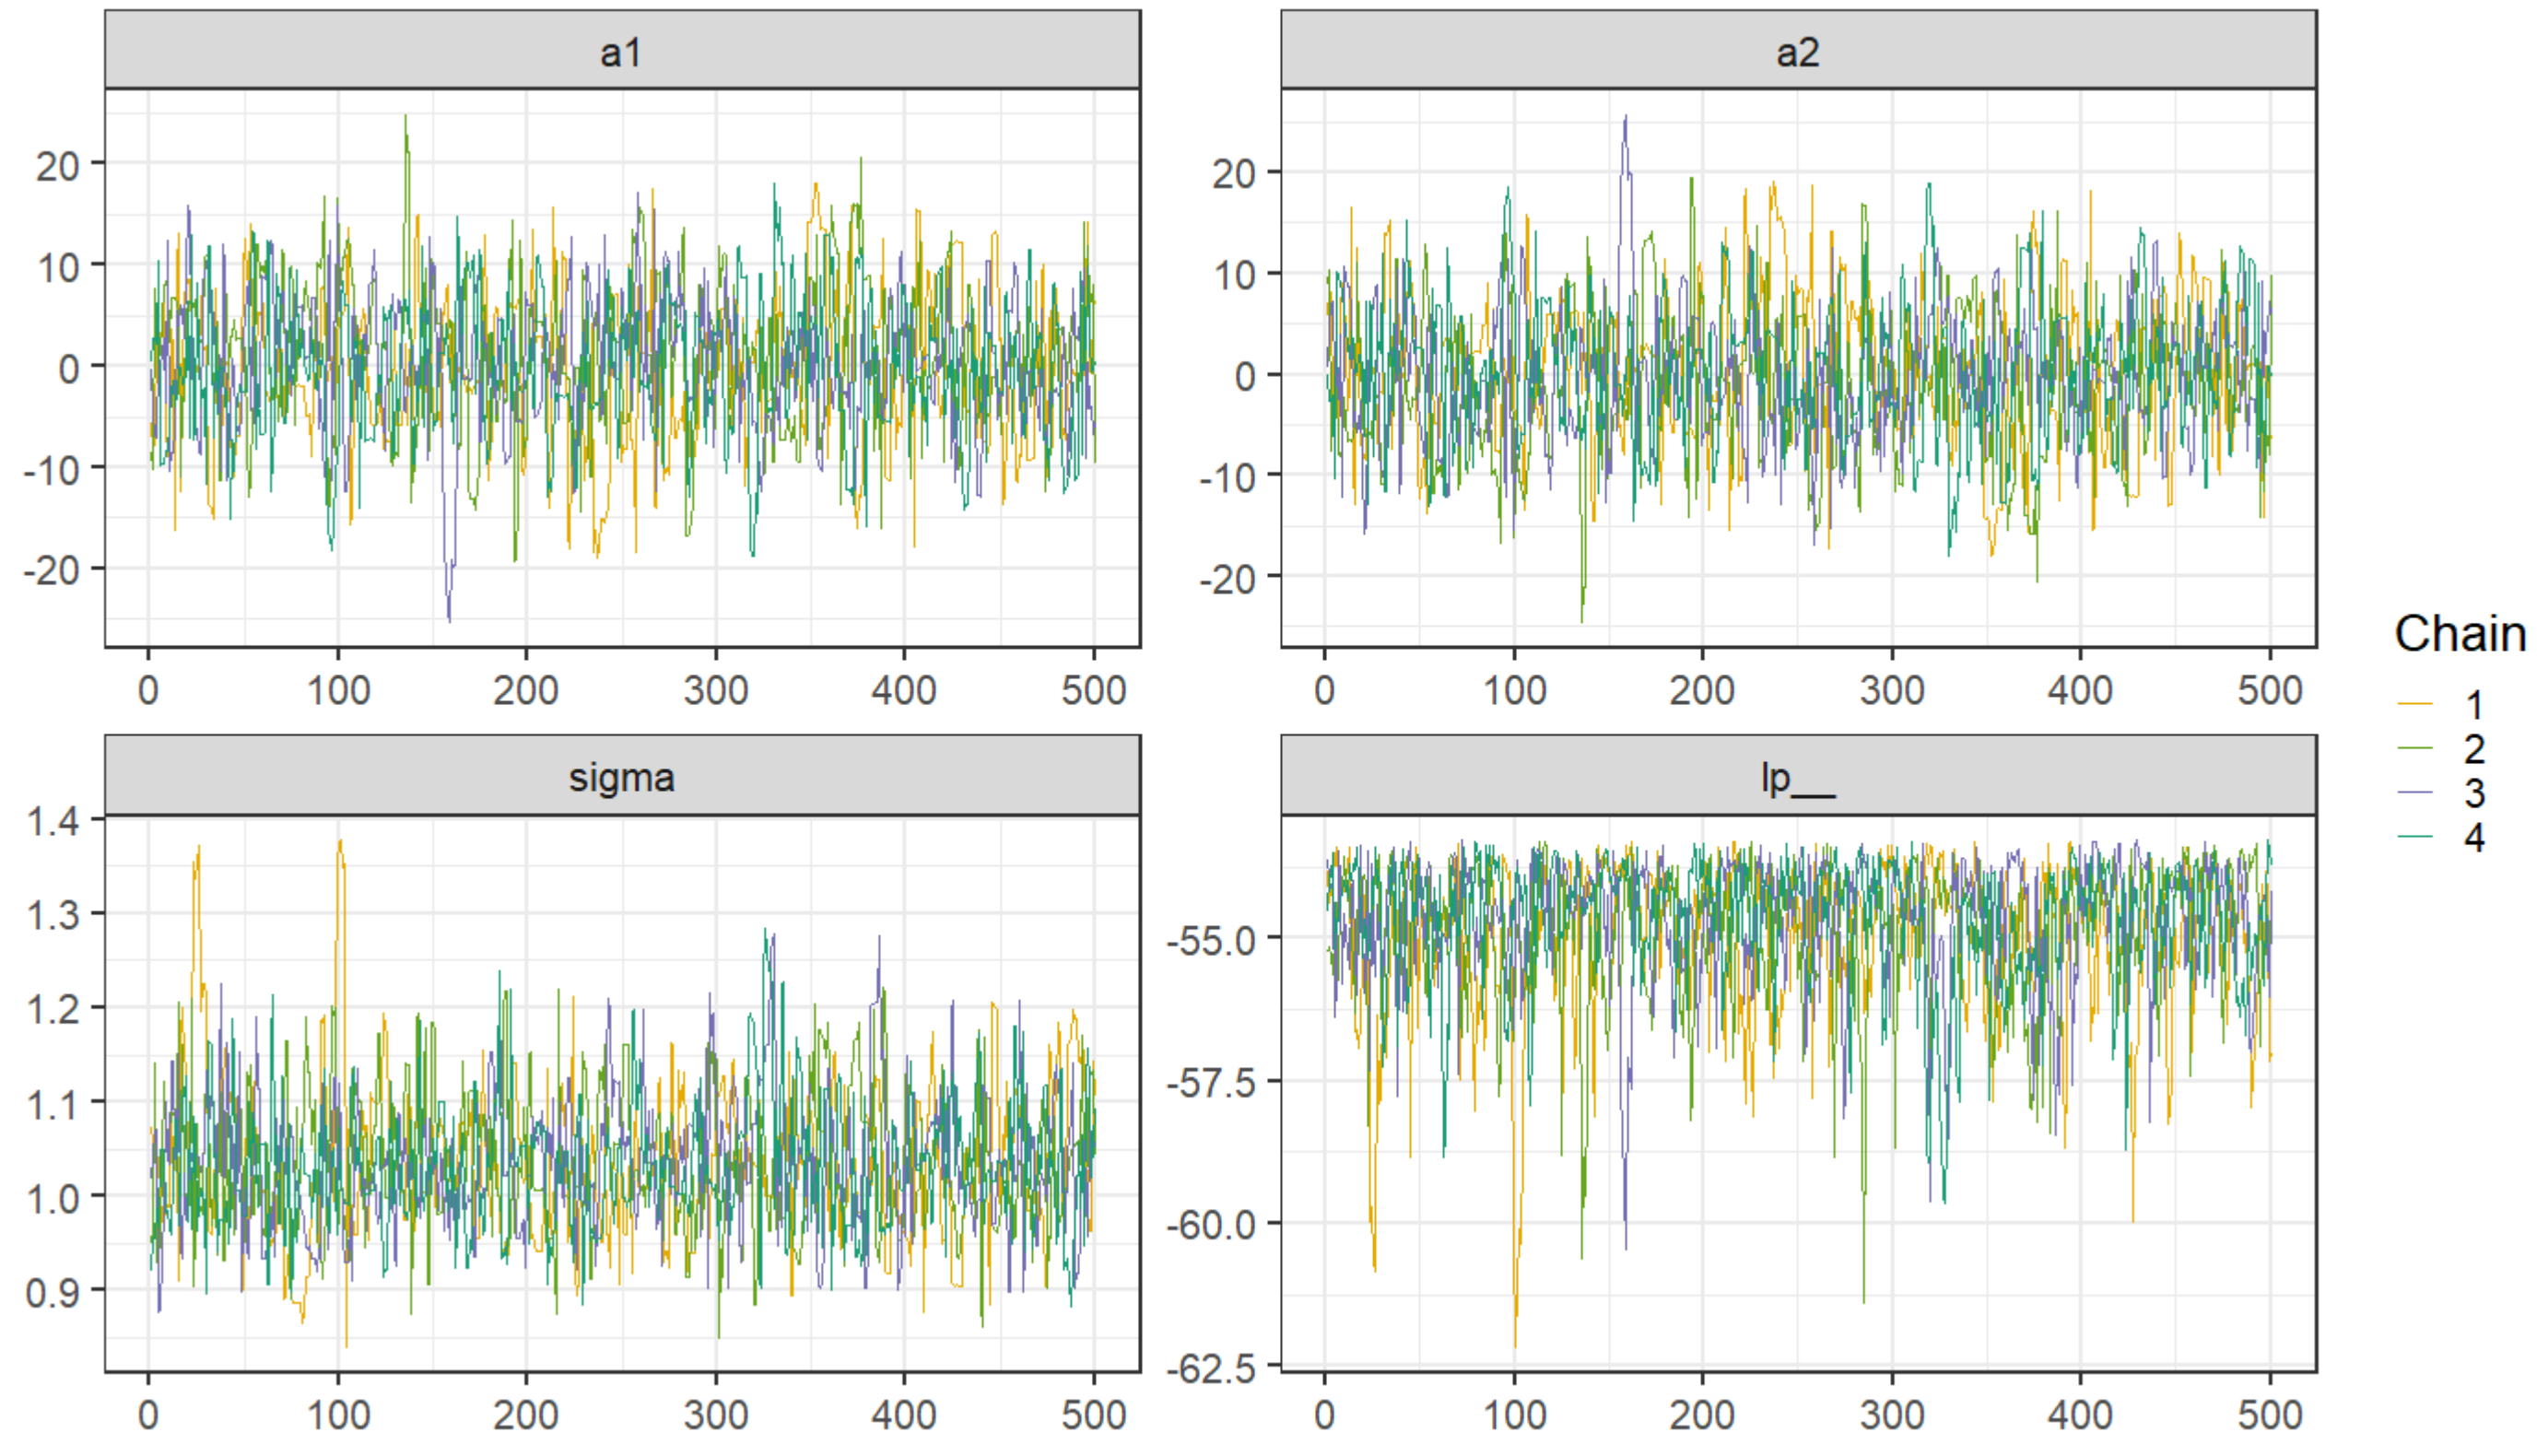
Pair-Correlation Plots

```
color_scheme_set("blue")
reg_md1_fit <- mdl_reg@stanfit
mcmc_pairs(reg_md1_fit, off_diag_args = list(size = 1, alpha = 0.05))
```



MCMC Trace Plots

```
color_scheme_set("brewer-Dark2")  
mcmc_trace(reg_mdl_fit)
```



MCMC Trace-Rank Plots

```
mcmc_rank_overlay(reg_md1_fit)
```

