

# Generalized Linear Madness

EES 5891-03

Bayesian Statistical Methods

Jonathan Gilligan

Class #24: Tuesday, November 29 2022

# Generalized Linear Models and Beyond

# Generalized Linear Models

- Linear Models:

$$y \sim \text{Normal}(\mu, \sigma)$$
$$\mu = \alpha + \sum_i \beta_i x_i$$

- Polynomial models become linear models

$$y \sim \text{Normal}(\mu, \sigma)$$
$$\mu = \alpha + \sum_i \beta_i x^i$$

- transform:

$$x_1 = x$$

$$x_2 = x^2$$

...

$$\mu = \alpha + \sum_i \beta_i x_i$$

- Generalized Linear Models (GLMs)

$$y \sim F(\mu)$$
$$\text{Link}(\mu) = \alpha + \sum_i \beta_i x_i$$

- $F$  could be Poisson, Binomial, etc.

- Link could be logit, log, etc.

- Generalized Additive Models (GAMs)

$$y \sim F(\mu)$$
$$\mu = \alpha + \sum_i f_i(x)$$

or


$$\mu = \alpha + \sum_i f_i(x_1, x_2, \dots)$$



# Beyond Linearity

- Linear models are widely applicable
  - You can use them out of the box
  - `quap` and `ulam` automatically take a linear model specification (`alist`) and turn it into R or Stan code
- Scientific models
  - Model equations from scientific knowledge
  - They may not be linear
  - Often require custom (“bespoke”) programming
    - Stan programming language

# Example: “Vitruvian Can”



# Vitruvian Can Model

- Height and Weight data set (from Chapter 4)

```
data(Howell11)
d <- Howell11

d <- d %>% mutate(h = height / mean(height),
                  w = weight / mean(weight))
```

- Model people as cylinders

$$V = \pi r^2 h$$


$$r = ph$$

$$V = \pi(ph)^2 h = \pi p^2 h^3$$

$$W = kV = k\pi p^2 h^3$$

- $V$  = volume
- $W$  = weight

- This is clearly *not* a linear model



# Statistical Model

$$W \sim \text{Log-Normal}(\mu, \sigma)$$

$$\exp(\mu) = \pi k p^2 h^3$$

$$k \sim \text{some prior}$$

$$p \sim \text{some prior}$$

$$\sigma \sim \text{Exponential}(1)$$

- Log-normal distribution ensures that  $W > 0$
- $\mu$  is the mean of  $\log(W)$
- Choosing priors
  - People are a lot less wide than tall, so
    - $r < h/2$
    - $p < 0.5$
    - Prior:  $p \sim \text{Beta}(2, 18)$ 
      - Mean =  $2/(2 + 18) = 0.1$
  - $k$  is density (weight / volume)
    - Prior:  $k \sim \text{Exponential}(0.5)$

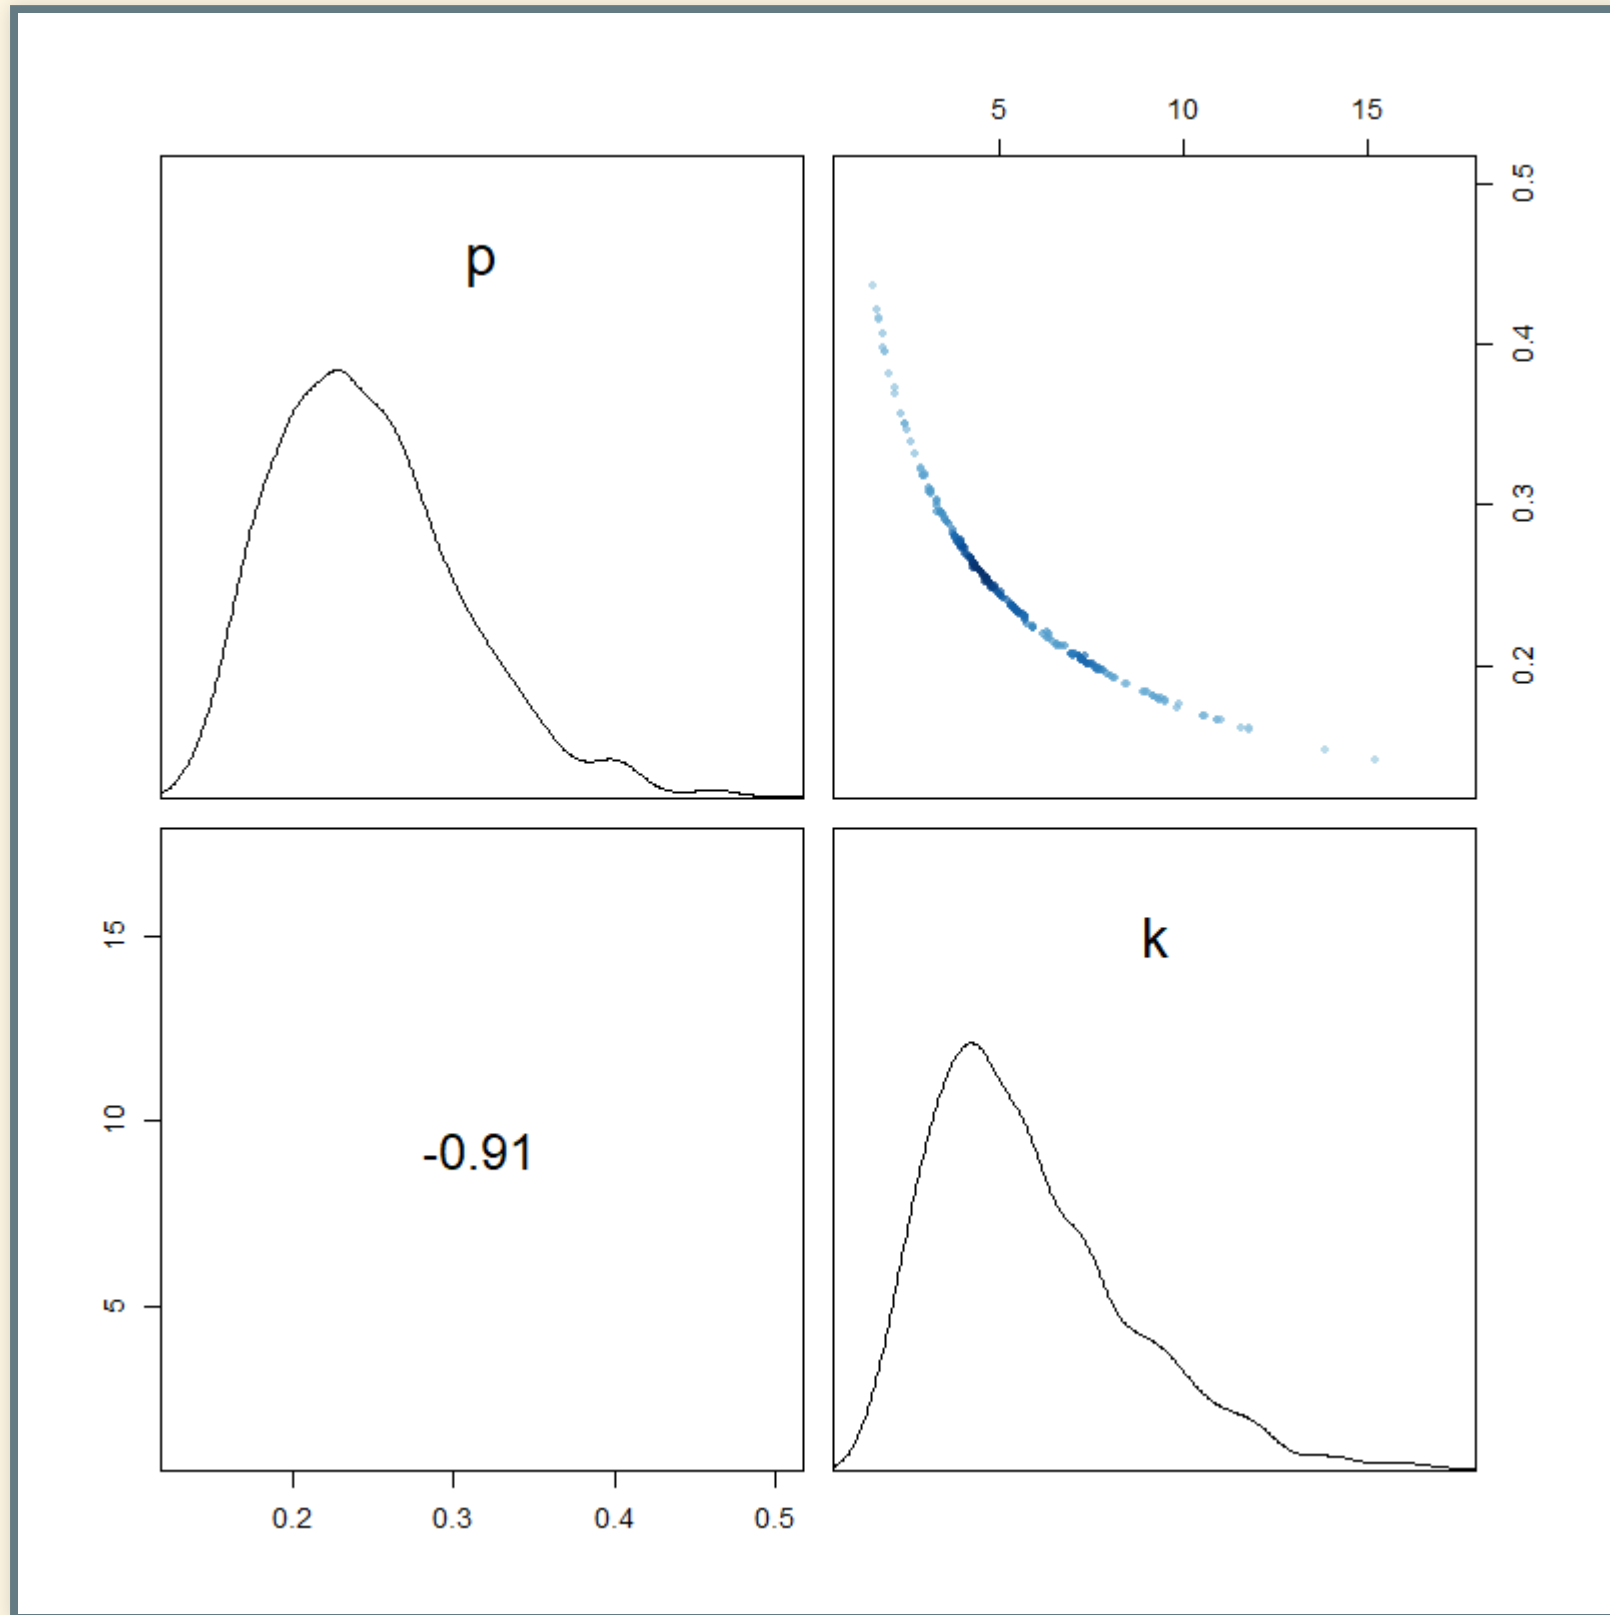
```
mdl <- ulam(  
  alist(  
    w ~ dlnorm(mu, sigma),  
    exp(mu) <- 3.14159265 * k * p^2 * h^3,  
    p ~ dbeta(2, 18),  
    k ~ dexp(0.5),  
    sigma ~ dexp(1)  
  ), data = d, chains = 4, cores = 4  
)
```

```
precis(mdl, digits = 2)
```


##		mean	sd	5.5%	94.5%	n_eff	Rhat4
##	p	0.25	0.06	0.17	0.35	499	1.01
##	k	5.81	2.70	2.47	10.92	570	1.00
##	sigma	0.21	0.01	0.20	0.22	770	1.01



# Model Results



```
pairs(mdl, pars = c("p", "k"))
```



```
new_data <- tibble(h = seq(0, max(d$h), length.out = 30))
w_sim <- predicted_draws(mdl, new_data, value = "w_pred") %>%
  group_by(.row) %>%
  summarize(h = head(h, 1), w = mean(w_pred),
            PI = list(set_names(PI(w_pred), c("min", "max")))) %>%
  unnest_wider(PI)
ggplot(w_sim, aes(x = h, y = w)) +
  geom_line(size = 1, color = "darkblue") +
  geom_ribbon(aes(ymin = min, ymax = max), fill = "darkblue",
            alpha = 0.3) +
  geom_point(data = d, size = 1, alpha = 0.3, color = "darkgreen")
  +
  labs(x = "height (scaled)", y = "weight (scaled)")
```

# Alternate Formulation

- Our model is

$$W \sim \text{Log-Normal}(\mu, \sigma)$$
$$\exp(\mu) = \pi k p^2 h^3$$

- But we could also write it


$$\log(W) \sim \text{Normal}(\mu, \sigma)$$
$$\mu = \log(\pi k p^2 h^3)$$
$$= \log(\pi) + \log(k) + 2 \log(p) + 3 \log(h)$$

which turns it into a linear regression.

# Latent-Variable Models

# Hidden Minds & Observed Behavior

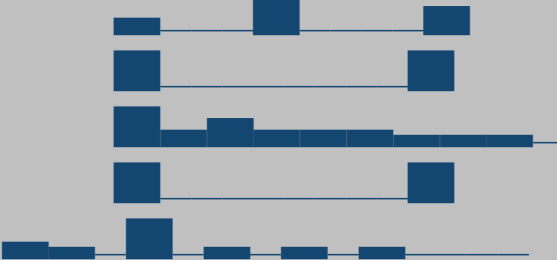
- Identifying children's decision-making strategies
  - "Choice box" has three colored tubes
    - When a ball is dropped into a tube, a toy comes out of the box.
  - Experiment: Groups of 5 children
    - Subject sees four other children use the choice box
    - 3 children choose the same color (**majority color**)
    - 1 child chooses another color (**minority color**)
    - The third color is unused
  - Which color does the subject choose?



# Data

```
data(Boxes)
precis(Boxes, digits = 2)
```

```
## 'data.frame': 629 obs. of 5 variables:
##      mean    sd 5.5% 94.5%      histogram
## y      2.12 0.73    1    3
## gender  1.51 0.50    1    2
## age     8.03 2.50    5   13
## majority_first 0.48 0.50    0    1
## culture  3.75 1.96    1    8
```



- **y**: Which color did they choose?
  1. Unchosen color
  2. Majority color
  3. Minority color
- **majority\_first**: Was the majority color chosen before the minority color?

```
table(Boxes$y) / length(Boxes$y)
```

```
##
##      1      2      3
## 0.2114467 0.4562798 0.3322734
```

- 45% of children chose the color that the majority of the four children chose.
- Does this mean that 45% of the children used a strategy of following the majority choice?
- Simulation:
  - Half of the children always choose the majority color (2).
  - The other half choose at random.

```
N <- 30
set.seed(163)
y1 <- sample(1:3, size = N/2, replace = TRUE)
y2 <- rep(2, N/2)
y <- sample(c(y1, y2))
table(y) / N
```

```
## y
##      1      2      3
## 0.1666667 0.6666667 0.1666667
```

- 2/3 of the children choose the majority color, but only half of them are actually following the majority.

# Making a Model

- We want to identify the strategies children are using,
  - We don't know what's going on in their heads.
  - We only know which colors they choose, not why.
- Possible strategies:
  1. Follow the majority
  2. Follow the minority
  3. Maverick (choose the unused color)
  4. Random choice
  5. Follow the first color chosen
- Each strategy has a *prior* for the choices (1, 2, or 3)
- Develop a posterior probability for each strategy

- The probability of choosing color  $y$  is

$$\Pr(y) = \sum_{s=1}^5 p_s \Pr(y|s),$$

where  $s$  indicates the strategy,  $p_s$  is the probability of using strategy  $s$ .

- **Simplex:** The probabilities of the 5 strategies form a vector  $(p_1, p_2, p_3, p_4, p_5)$ , where they add up to 1.
- **Dirichlet prior:** for  $(p_1, p_2, p_3, p_4, p_5)$ .

$$p \sim \text{Dirichlet}([4, 4, 4, 4, 4])$$

The  $p$ 's don't have to be the same, but each has the same probability to be larger or smaller

- Larger numbers instead of 4 would raise the probability for all the  $p$ 's being equal.

# Coding the Model

$y \sim \text{Categorical}(\theta)$

$\theta_j = \sum_{s=1}^5 p_s \Pr(j|s) \quad \text{for } j = 1 \dots 3$

$p \sim \text{Dirichlet}([4, 4, 4, 4, 4])$

- This is a complicated model
  - Can't use `ulam` or `quap`
  - Code it directly in Stan

```
data("Boxes_model")
cat(Boxes_model)
```

```
##
## data{
##   int N;
##   int y[N];
##   int majority_first[N];
## }
## parameters{
##   simplex[5] p;
## }
## model{
##   vector[5] phi;
##
##   // prior
##   p ~ dirichlet( rep_vector(4,5) );
##
##   // probability of data
##   for ( i in 1:N ) {
##     if ( y[i]==2 ) phi[1]=1; else phi[1]=0; // majority
##     if ( y[i]==3 ) phi[2]=1; else phi[2]=0; // minority
##     if ( y[i]==1 ) phi[3]=1; else phi[3]=0; // maverick
##     phi[4]=1.0/3.0; // random
##     if ( majority_first[i]==1 ) // follow first
##       if ( y[i]==2 ) phi[5]=1; else phi[5]=0;
##     else
##       if ( y[i]==3 ) phi[5]=1; else phi[5]=0;
##
##     // compute log( p_s * Pr(y_i|s) )
##     for ( j in 1:5 ) phi[j] = log(p[j]) + log(phi[j]);
##     // compute average log-probability of y_i
##     target += log_sum_exp( phi );
##   }
## }
```

# Running the Model

```
dat_list <- list(  
  N = nrow(Boxes),  
  y = Boxes$y,  
  majority_first = Boxes$majority_first  
)
```

```
mdl_box <- stan(model_code = Boxes_model, data = dat_list,  
  chains = 4, cores = 4)
```

```
mdl_box
```

```
## Inference for Stan model: rt_cmdstanr_82e68937d1e417e265d6739ef917f4e6-  
202211290539-1-4b9fa0.  
## 4 chains, each with iter=1000; warmup=500; thin=1;  
## post-warmup draws per chain=500, total post-warmup draws=2000.  
##  
##           mean se_mean   sd   2.5%   25%   50%   75%   97.5% n_eff Rhat  
## p[1]      0.26     0.00 0.04   0.18   0.23   0.26   0.28   0.33   555 1.01  
## p[2]      0.14     0.00 0.03   0.07   0.12   0.14   0.16   0.20   543 1.00  
## p[3]      0.15     0.00 0.03   0.08   0.13   0.15   0.17   0.20   562 1.00  
## p[4]      0.20     0.00 0.08   0.06   0.14   0.19   0.25   0.36   433 1.00  
## p[5]      0.26     0.00 0.03   0.19   0.24   0.26   0.28   0.32  1554 1.00  
## lp__ -667.21     0.06 1.54 -671.17 -667.98 -666.85 -666.06 -665.34   602 1.00  
##  
## Samples were drawn using NUTS(diag_e) at Tue Nov 29 5:39:48 AM 2022.  
## For each parameter, n_eff is a crude measure of effective sample size,  
## and Rhat is the potential scale reduction factor on split chains (at  
## convergence, Rhat=1).
```

```
p_labels <- c("1 Majority", "2 Minority", "3 Maverick", "4  
  Random", "5 Follow First")  
plot(precis(mdl_box, 2), labels = p_labels)
```

